



ECE408 / CS483 / CSE408  
Summer 2025

Applied Parallel Programming

Lecture 5: Locality and  
Tiled Matrix Multiplication



# What Will You Learn Today?

- to evaluate the performance implications of global memory accesses
- prepare for Lab 3: tiled matrix multiplication
- to assess the benefits of tiling

# Kernel Invocation (Host-side Code) vk

```
// Setup the execution configuration
// BLOCK_WIDTH is a #define constant
dim3 dimGrid(ceil((1.0*Width)/BLOCK_WIDTH),
    ceil((1.0*Width)/BLOCK_WIDTH), 1);

dim3 dimBlock(BLOCK_WIDTH, BLOCK_WIDTH, 1);

// Launch the device computation threads!
MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd, Pd, Width);
```

# The Problem: Accesses to Global Memory

```
__global__ void MatrixMulKernel(float* d_M, float* d_N, float* d_P, int Width)
{
    // Calculate the row index of the d_P element and d_M
    int Row = blockIdx.y*blockDim.y+threadIdx.y;
    // Calculate the column idenx of d_P and d_N
    int Col = blockIdx.x*blockDim.x+threadIdx.x;

    if ((Row < Width) && (Col < Width)) {
        float Pvalue = 0;
        // each thread computes one element of the block sub-matrix
        for (int k = 0; k < Width; ++k)
            Pvalue += d_M[Row*Width+k] *
                      d_N[k*Width+Col];
        d_P[Row*Width+Col] = Pvalue;
    }
}
```

A diagram illustrating the memory access flow. It shows a yellow box labeled "accesses to global memory" containing the expression `d_M[Row*Width+k]`. An arrow points from this box to a yellow box labeled "d\_P[Row\*Width+Col] = Pvalue". Another arrow points from the expression `d_N[k*Width+Col];` to the same yellow box.

## Review: 4B of Data per FLOP

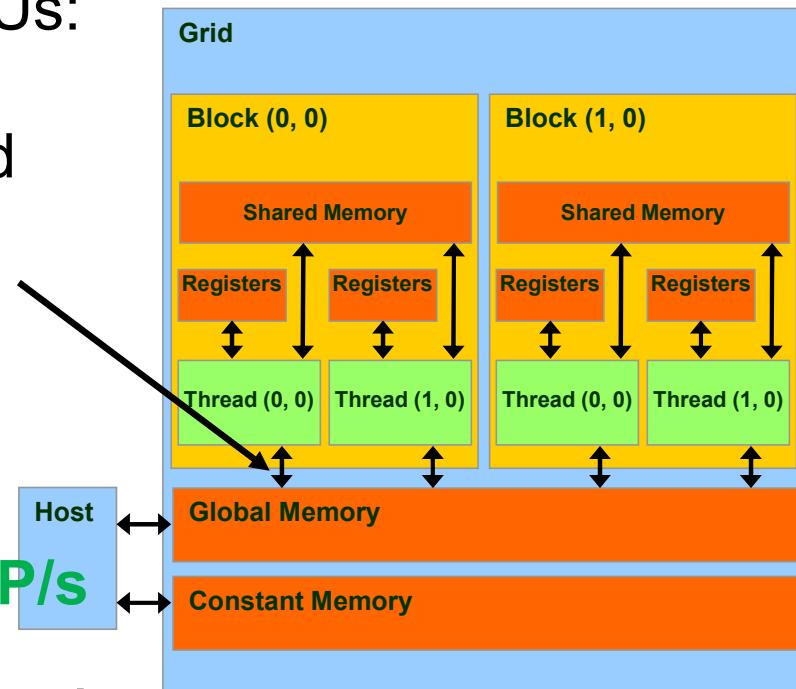
- Each threads access global memory
  - for elements of **M** and **N**:
    - 4B each**, or **8B per pair**.
    - (And once TOTAL to **P** per thread—ignore it.)
- With each pair of elements,
  - a thread does a single multiply-add,
    - 2 FLOP**—floating-point operations.
- So for every FLOP,
  - a thread needs** 4B from memory:
    - 4B / FLOP**.

## Review: Extremely Poor Performance

- One generation of GPUs:
  - **1,000 GFLOP/s** of compute power, and
  - **150 GB/s** of memory bandwidth.
- Dividing bandwidth by memory requirements:

$$\frac{150 \text{ GB/s}}{4 \text{ B/FLOP}} = 37.5 \text{ GFLOP/s}$$

which **limits computation!**



# The Solution? Reuse Memory Accesses!

But **37.5 GFLOPs is a limit.**

In an **actual execution**,

- memory is not busy all the time, and
- the code **runs at** about **25 GFLOPs**.

To get closer to 1,000 GFLOPs

- we **need to** drastically **cut down**
- **accesses to global memory.**

But ... how?



# Tiled Matrix-Matrix Multiplication using Shared Memory

# A Common Programming Strategy

- The dilemma:
  - Matrices **M** and **N** are large.
  - They **fit** easily **in global memory**, but that's **slow**.
  - **Shared memory is fast, but M and N don't fit.**
- The solution:
  - **Break M and N into tiles**
  - (called blocks in the much older CPU literature).
  - **Read a tile** into shared memory.
  - **Use the tile** from shared memory.
  - **Repeat** until done.

# A Common Programming Strategy

- In a GPU, **only threads in a block can use shared memory.**
- Thus, each **block** operates on **separate tiles**:
  - **Read tile(s)** into shared memory **using multiple threads** to exploit memory-level parallelism.
  - **Compute** based on shared memory tiles.
  - **Repeat.**
  - **Write results back to global memory.**

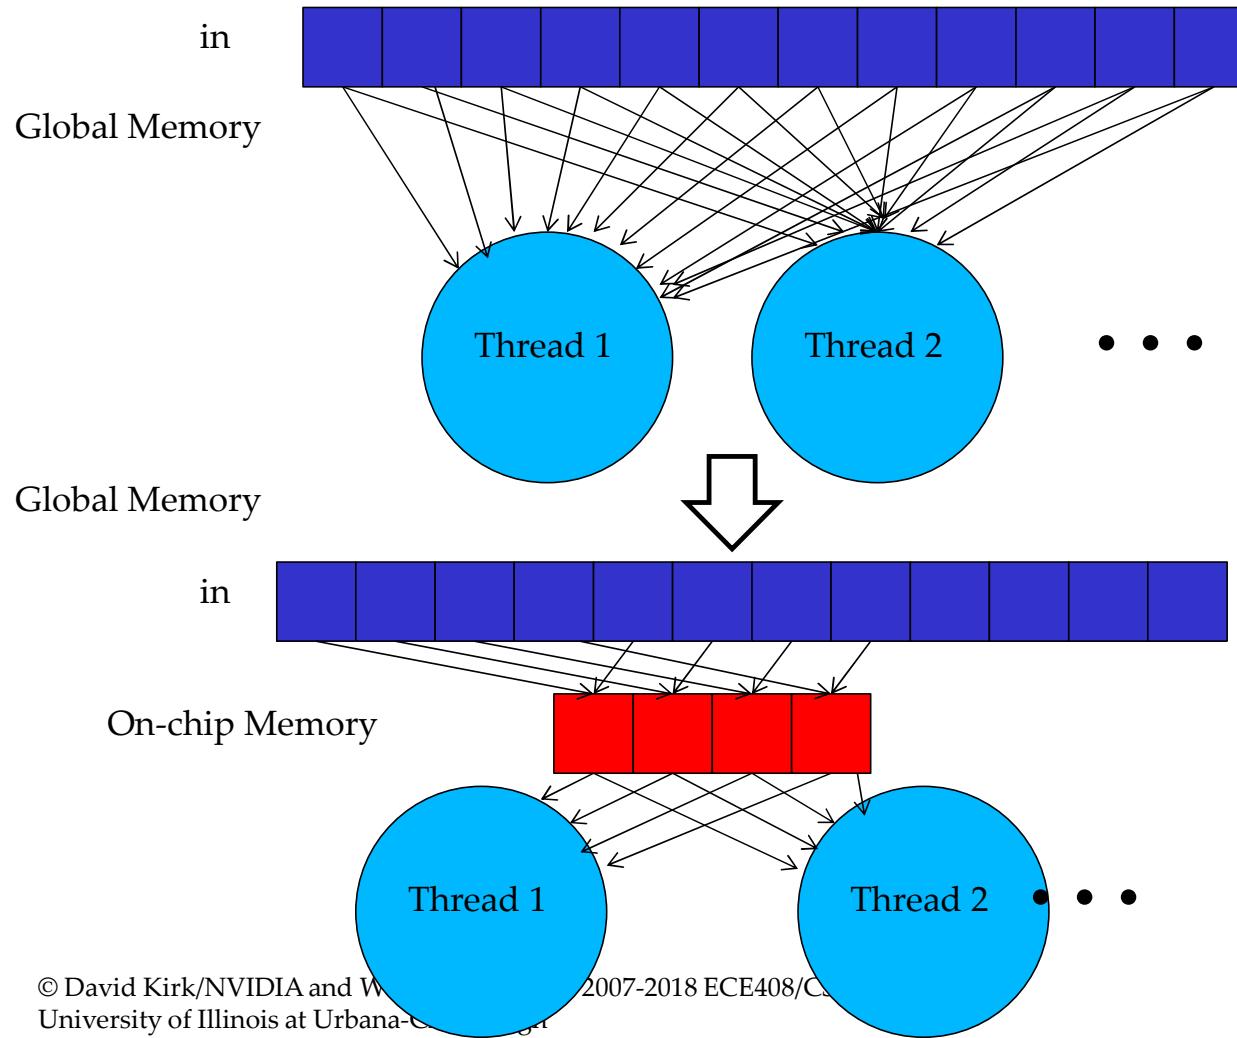
# Declaring Shared Memory Arrays

```
__global  
void MatrixMulKernel(float* M, float* N, float* P, int Width)  
{  
    __shared__ float subTileM[TILE_WIDTH][TILE_WIDTH];  
    __shared__ float subTileN[TILE_WIDTH][TILE_WIDTH];
```



Common across  
all threads in a  
block

# Shared Memory Tiling Basic Idea

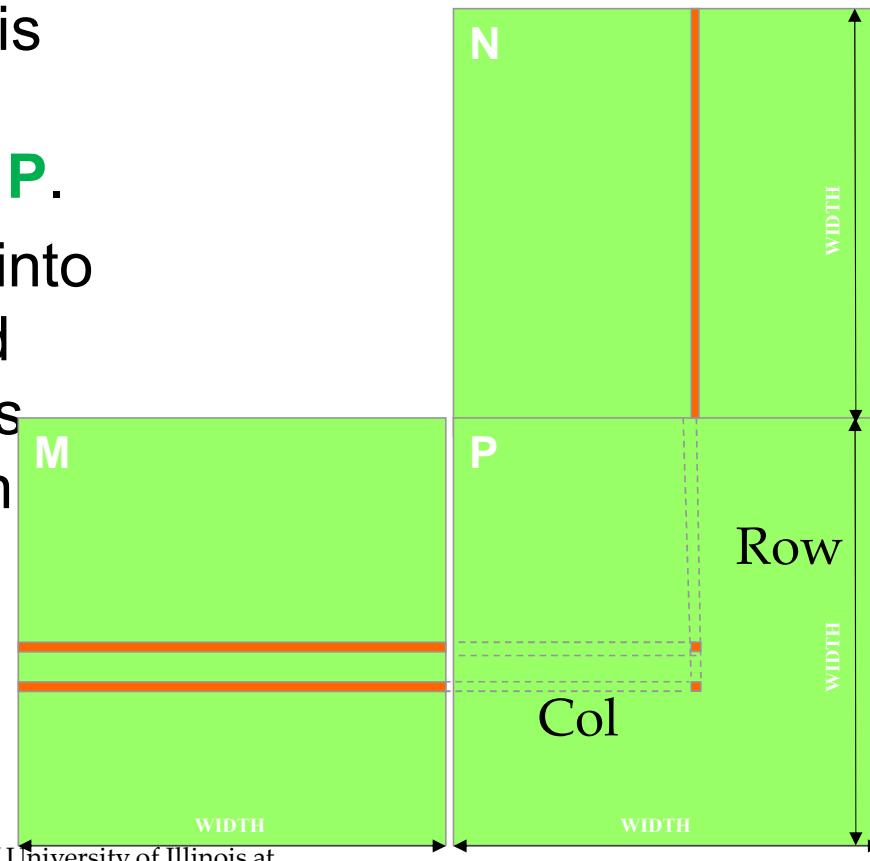


# Transform Global Memory Accesses into On-Chip Accesses

- **Identify a tile of global data**
  - with each datum accessed by multiple threads
  - and/or accessed repeatedly.
- **Change access pattern in kernel**
  - Load the tile from global memory into on-chip memory
  - Have the threads access the data from on-chip memory
  - Move on to the next tile

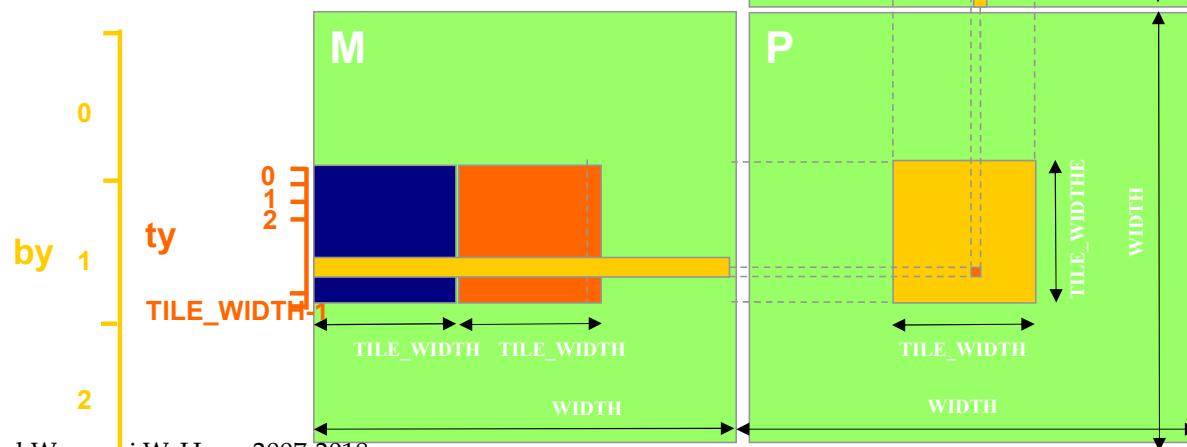
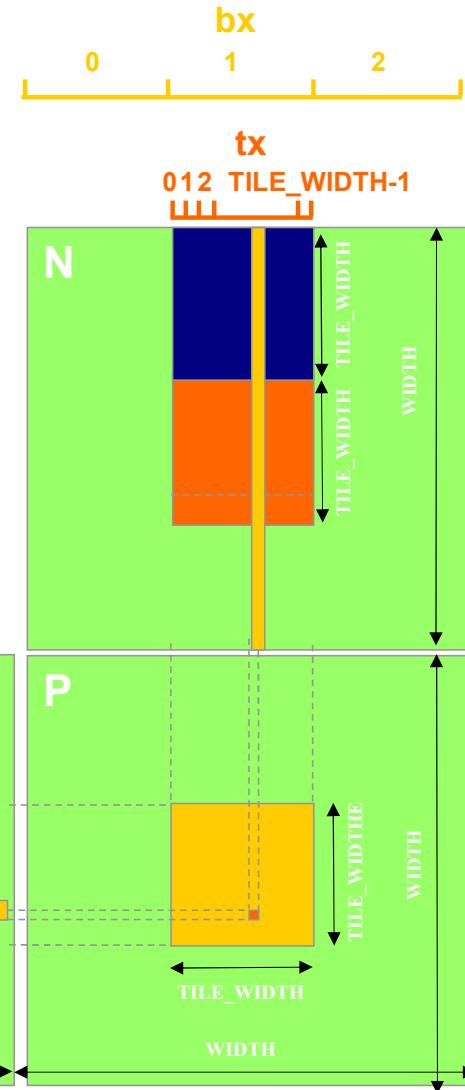
# Idea: Place global memory data into Shared Memory for reuse

- Each input element is used to calculate **WIDTH** elements of **P**.
- Load each element into Shared Memory and have several threads use the local version to reduce memory bandwidth.



# Tiled Multiply

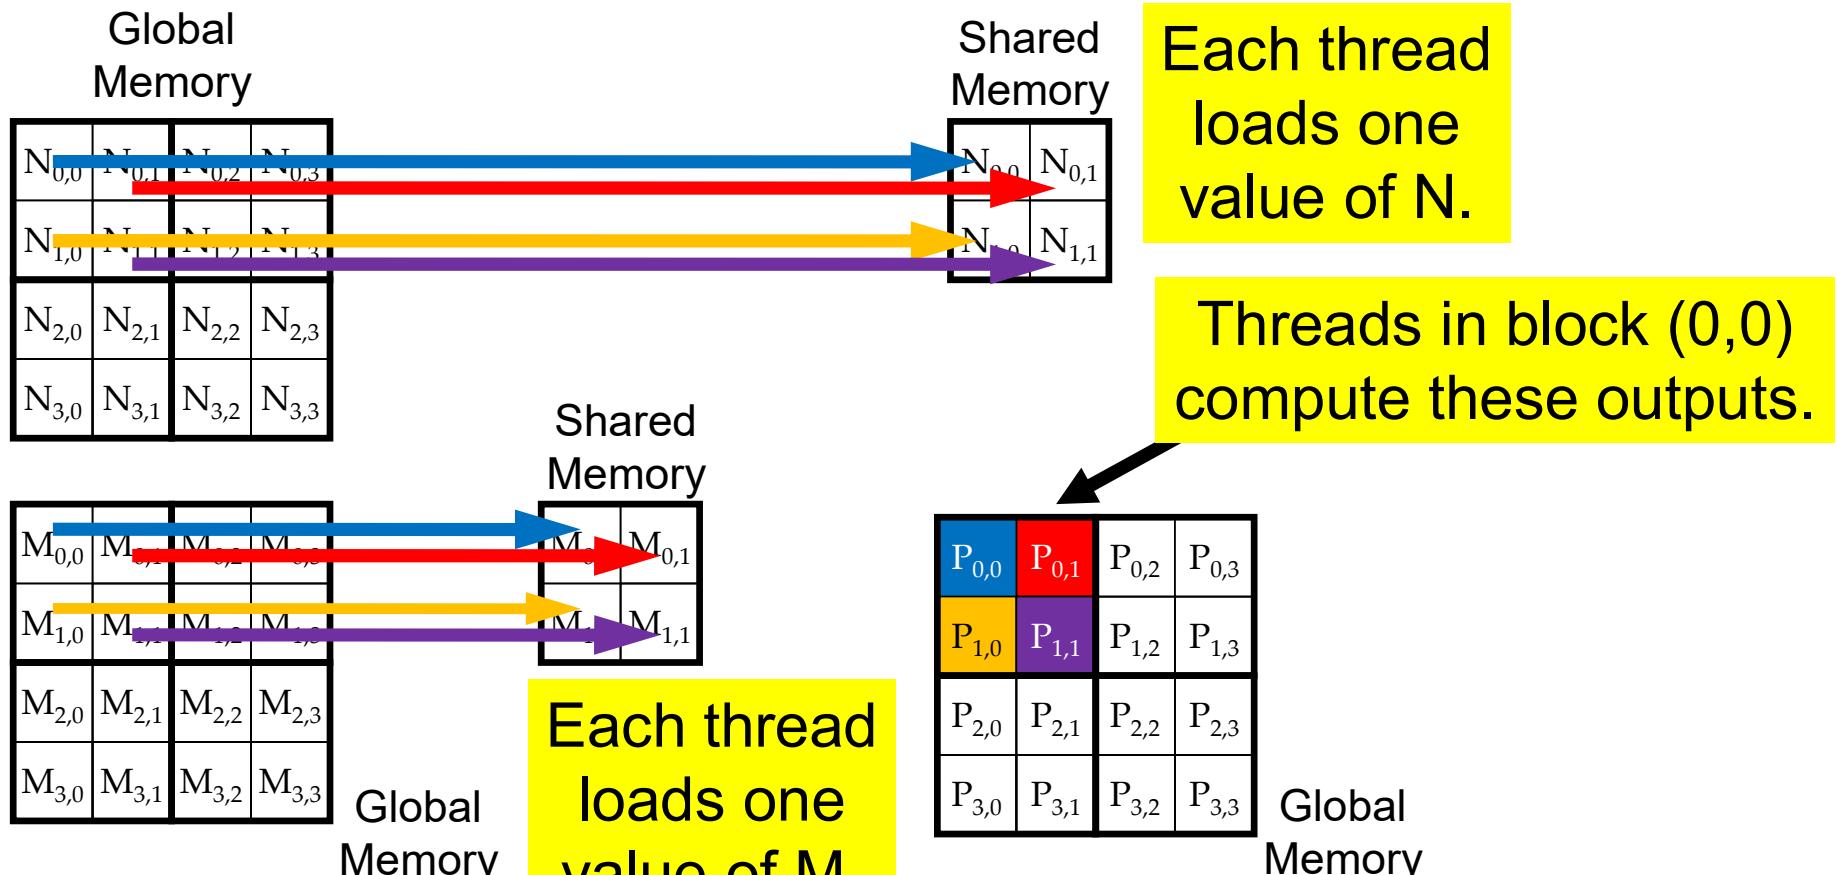
- Break up the execution of the kernel into phases so that the data accesses in each phase are focused on one subset (tile) of **M** and **N**



# Loading a Tile

- All threads in a block participate
  - Each thread loads
    - one **M** element and
    - one **N** element
  - in basic tiling code.
- Assign the loaded element to each thread such that the accesses within each warp is coalesced (more later).

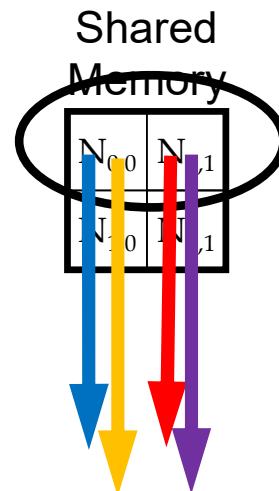
# Thread Block (0,0) Starts by Loading Two Tiles



# Thread Block (0,0) Computes on Shared Tiles (Iter 0)

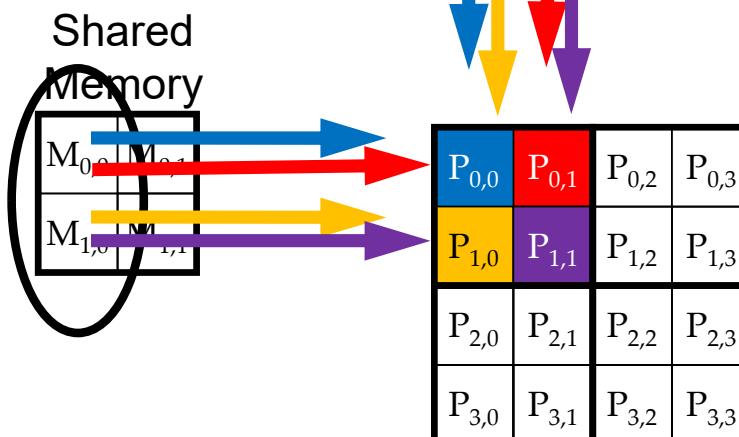
Global Memory			
N <sub>0,0</sub>	N <sub>0,1</sub>	N <sub>0,2</sub>	N <sub>0,3</sub>
N <sub>1,0</sub>	N <sub>1,1</sub>	N <sub>1,2</sub>	N <sub>1,3</sub>
N <sub>2,0</sub>	N <sub>2,1</sub>	N <sub>2,2</sub>	N <sub>2,3</sub>
N <sub>3,0</sub>	N <sub>3,1</sub>	N <sub>3,2</sub>	N <sub>3,3</sub>

Each thread loads one value of N (from tile).



Threads use shared data in compute loop iteration 0.

Global Memory			
M <sub>0,0</sub>	M <sub>0,1</sub>	M <sub>0,2</sub>	M <sub>0,3</sub>
M <sub>1,0</sub>	M <sub>1,1</sub>	M <sub>1,2</sub>	M <sub>1,3</sub>
M <sub>2,0</sub>	M <sub>2,1</sub>	M <sub>2,2</sub>	M <sub>2,3</sub>
M <sub>3,0</sub>	M <sub>3,1</sub>	M <sub>3,2</sub>	M <sub>3,3</sub>



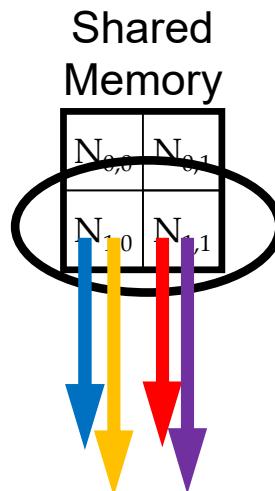
Each thread loads one value of M (from tile).

Global  
Memory

# Thread Block (0,0) Computes on Shared Tiles (Iter 1)

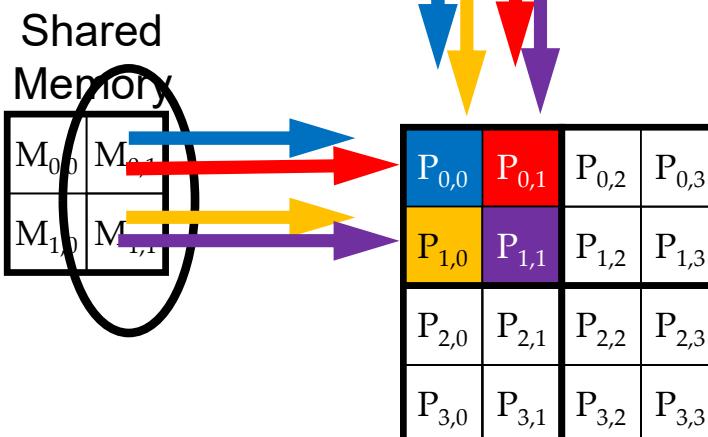
Global Memory			
N <sub>0,0</sub>	N <sub>0,1</sub>	N <sub>0,2</sub>	N <sub>0,3</sub>
N <sub>1,0</sub>	N <sub>1,1</sub>	N <sub>1,2</sub>	N <sub>1,3</sub>
N <sub>2,0</sub>	N <sub>2,1</sub>	N <sub>2,2</sub>	N <sub>2,3</sub>
N <sub>3,0</sub>	N <sub>3,1</sub>	N <sub>3,2</sub>	N <sub>3,3</sub>

Each thread loads one value of N (from tile).



Threads use shared data in compute loop iteration 1.

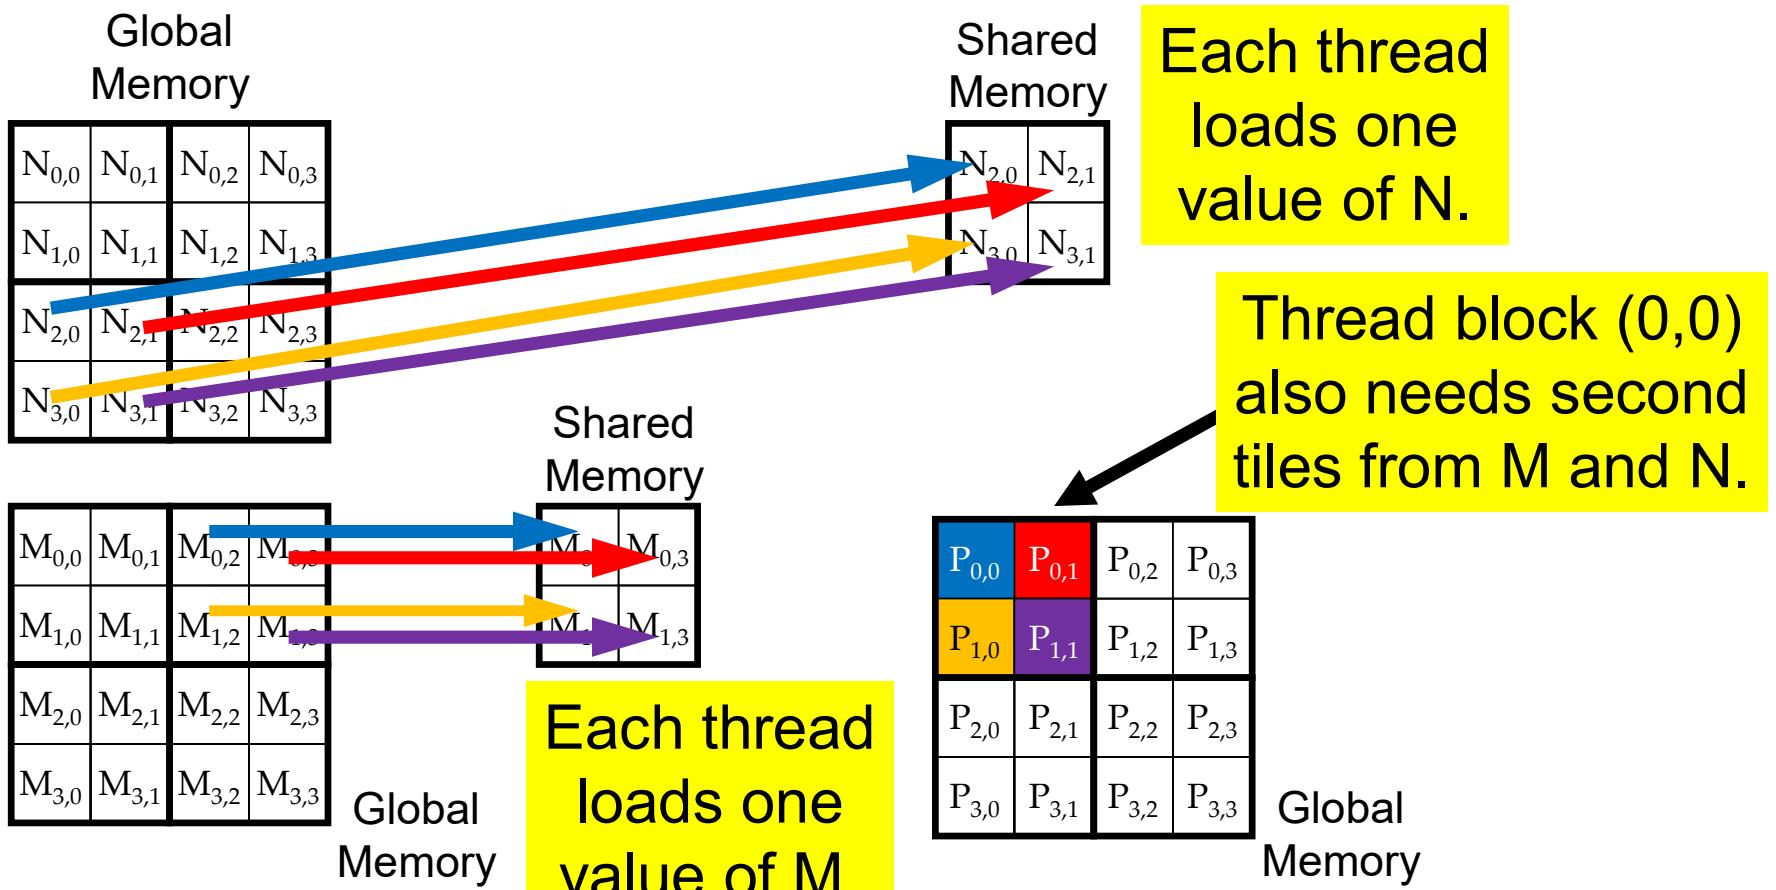
Global Memory			
M <sub>0,0</sub>	M <sub>0,1</sub>	M <sub>0,2</sub>	M <sub>0,3</sub>
M <sub>1,0</sub>	M <sub>1,1</sub>	M <sub>1,2</sub>	M <sub>1,3</sub>
M <sub>2,0</sub>	M <sub>2,1</sub>	M <sub>2,2</sub>	M <sub>2,3</sub>
M <sub>3,0</sub>	M <sub>3,1</sub>	M <sub>3,2</sub>	M <sub>3,3</sub>



Each thread loads one value of M (from tile).

Global Memory

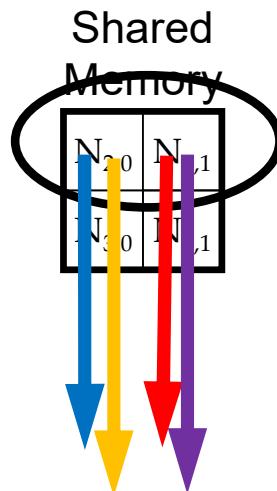
## Then Thread Block (0,0) Loads Two New Tiles



# Inner Loop Iteration is Now Identical (Iter 0)!

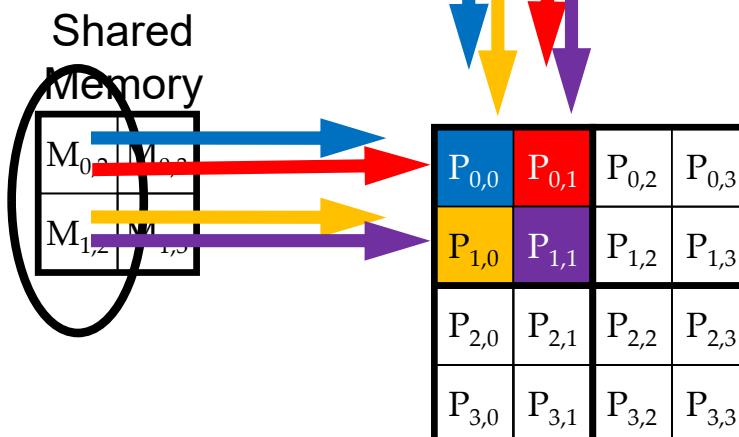
Global Memory			
N <sub>0,0</sub>	N <sub>0,1</sub>	N <sub>0,2</sub>	N <sub>0,3</sub>
N <sub>1,0</sub>	N <sub>1,1</sub>	N <sub>1,2</sub>	N <sub>1,3</sub>
N <sub>2,0</sub>	N <sub>2,1</sub>	N <sub>2,2</sub>	N <sub>2,3</sub>
N <sub>3,0</sub>	N <sub>3,1</sub>	N <sub>3,2</sub>	N <sub>3,3</sub>

Each thread loads one value of N (from tile).



Threads use shared data in compute loop iteration 0.

Global Memory			
M <sub>0,0</sub>	M <sub>0,1</sub>	M <sub>0,2</sub>	M <sub>0,3</sub>
M <sub>1,0</sub>	M <sub>1,1</sub>	M <sub>1,2</sub>	M <sub>1,3</sub>
M <sub>2,0</sub>	M <sub>2,1</sub>	M <sub>2,2</sub>	M <sub>2,3</sub>
M <sub>3,0</sub>	M <sub>3,1</sub>	M <sub>3,2</sub>	M <sub>3,3</sub>



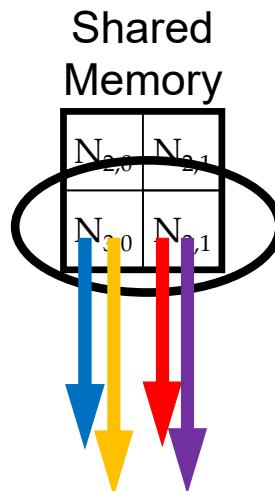
Each thread loads one value of M (from tile).

Global  
Memory

# Inner Loop Iteration is Now Identical (Iter 1)!

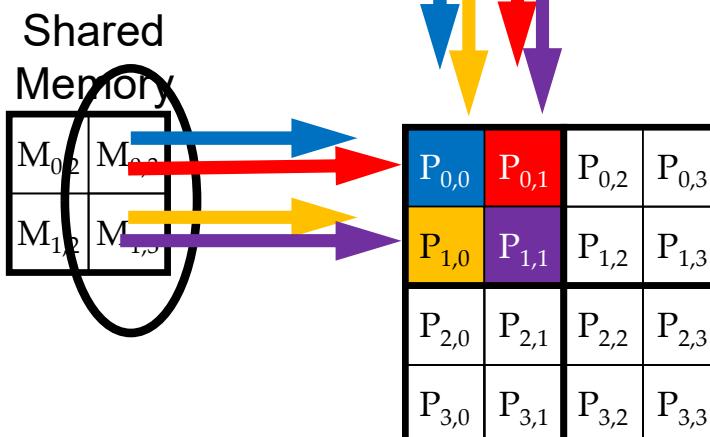
Global Memory			
N <sub>0,0</sub>	N <sub>0,1</sub>	N <sub>0,2</sub>	N <sub>0,3</sub>
N <sub>1,0</sub>	N <sub>1,1</sub>	N <sub>1,2</sub>	N <sub>1,3</sub>
N <sub>2,0</sub>	N <sub>2,1</sub>	N <sub>2,2</sub>	N <sub>2,3</sub>
N <sub>3,0</sub>	N <sub>3,1</sub>	N <sub>3,2</sub>	N <sub>3,3</sub>

Each thread loads one value of N (from tile).



Threads use shared data in compute loop iteration 1.

Global Memory			
M <sub>0,0</sub>	M <sub>0,1</sub>	M <sub>0,2</sub>	M <sub>0,3</sub>
M <sub>1,0</sub>	M <sub>1,1</sub>	M <sub>1,2</sub>	M <sub>1,3</sub>
M <sub>2,0</sub>	M <sub>2,1</sub>	M <sub>2,2</sub>	M <sub>2,3</sub>
M <sub>3,0</sub>	M <sub>3,1</sub>	M <sub>3,2</sub>	M <sub>3,3</sub>



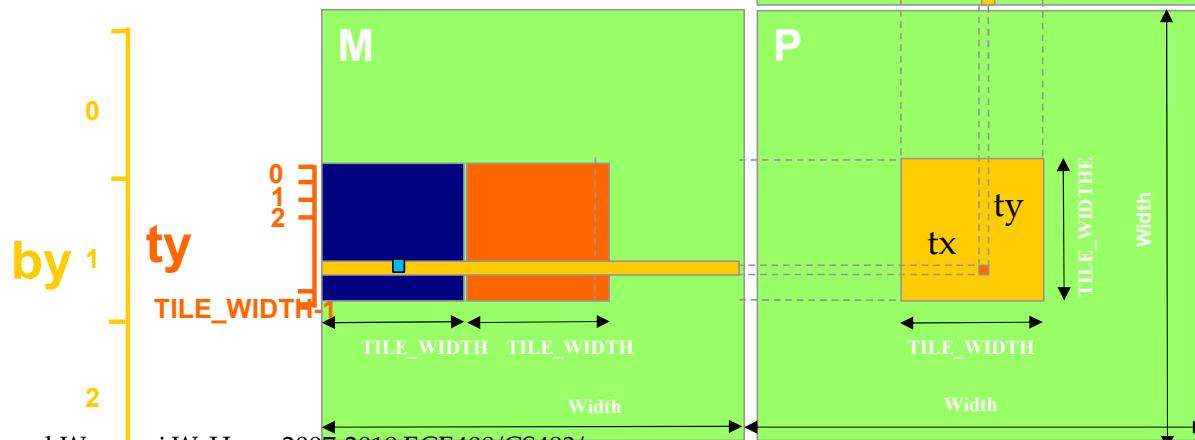
Each thread loads one value of M (from tile).

Global Memory

# Loading an Input Tile 0

Tile 0 2D indexing for each thread:

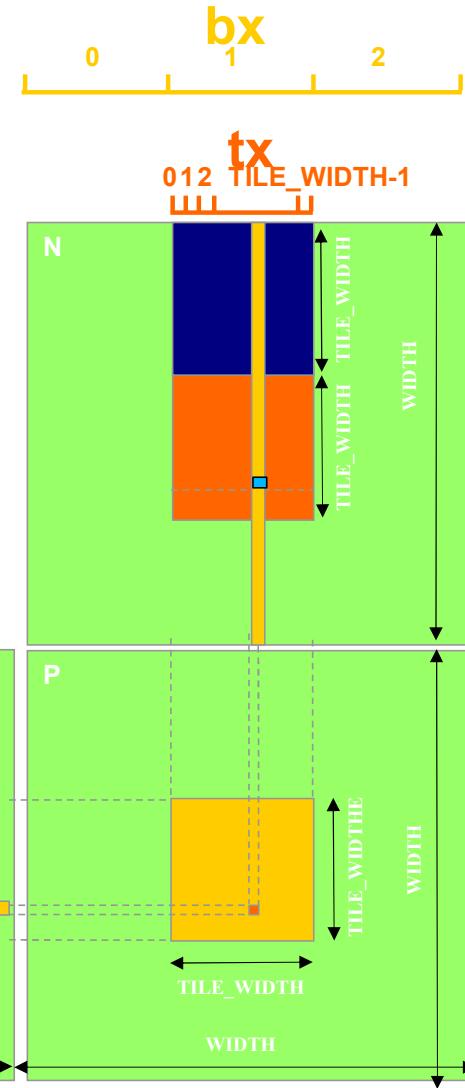
$M[\text{Row}] [\text{tx}]$   
 $N[\text{ty}] [\text{Col}]$



# Loading an Input Tile 1

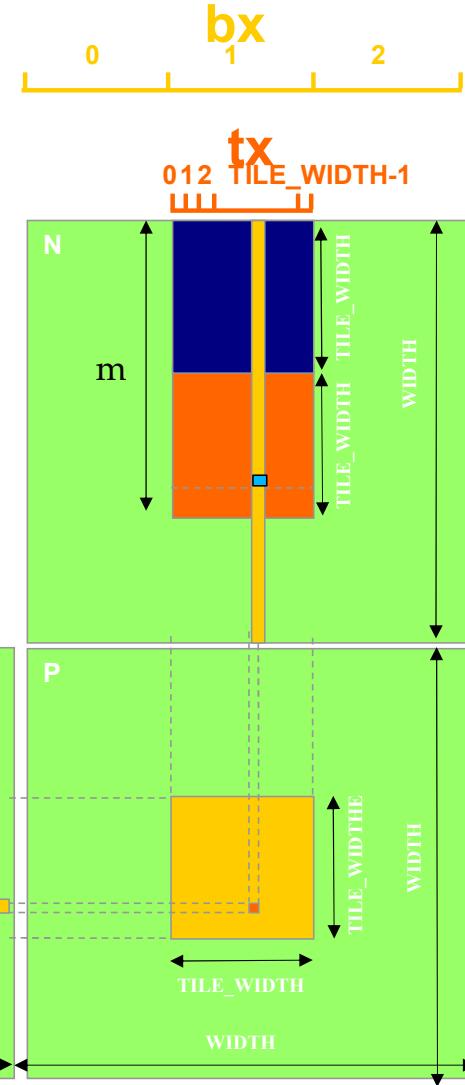
Accessing tile 1 in 2D indexing:

$M[\text{Row}] [1 * \text{TILE\_WIDTH} + \text{tx}]$   
 $N[1 * \text{TILE\_WIDTH} + \text{ty}] [\text{Col}]$



# Loading an Input Tile m

However, recall that M and N are dynamically allocated  
and can only use 1D indexing:

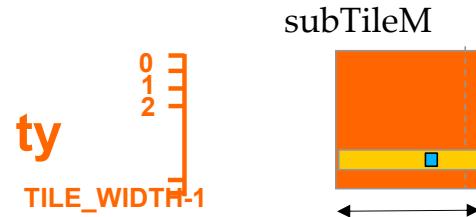
$$M[\text{Row}] [m * \text{TILE\_WIDTH} + tx]$$
$$M[\text{Row} * \text{Width} + m * \text{TILE\_WIDTH} + tx]$$
$$N[m * \text{TILE\_WIDTH} + ty] [\text{Col}]$$
$$N[(m * \text{TILE\_WIDTH} + ty) * \text{Width} + \text{Col}]$$


# Accessing a Tile

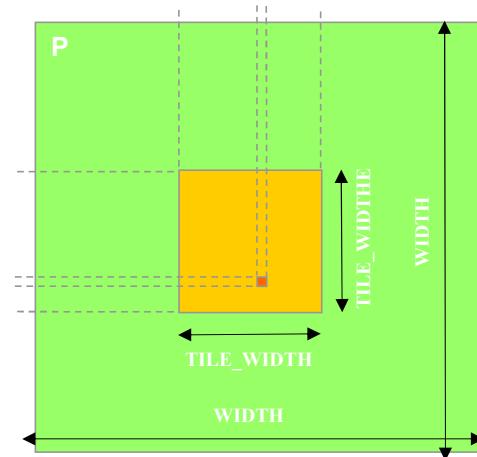
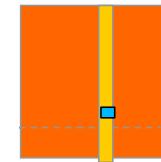
To perform the  $k^{\text{th}}$  step of the product within the tile:

$\text{subTileM}[ty][k]$

$\text{subTileN}[k][tx]$



$\text{subTileN}$





# We're Not There Yet!

- But ...
- How can a thread know ...
  - That another thread has finished its part of the tile?
  - Or that another thread has finished using the previous tile?

We need to synchronize!

# Leveraging Parallel Strategies

- **Bulk synchronous execution:**  
threads execute roughly in unison
  1. Do some work
  2. Wait for others to catch up
  3. Repeat
- **Much easier programming model**
  - Threads only parallel within a section
  - Debug lots of little programs
  - Instead of one large one.
- **Dominates high-performance applications**

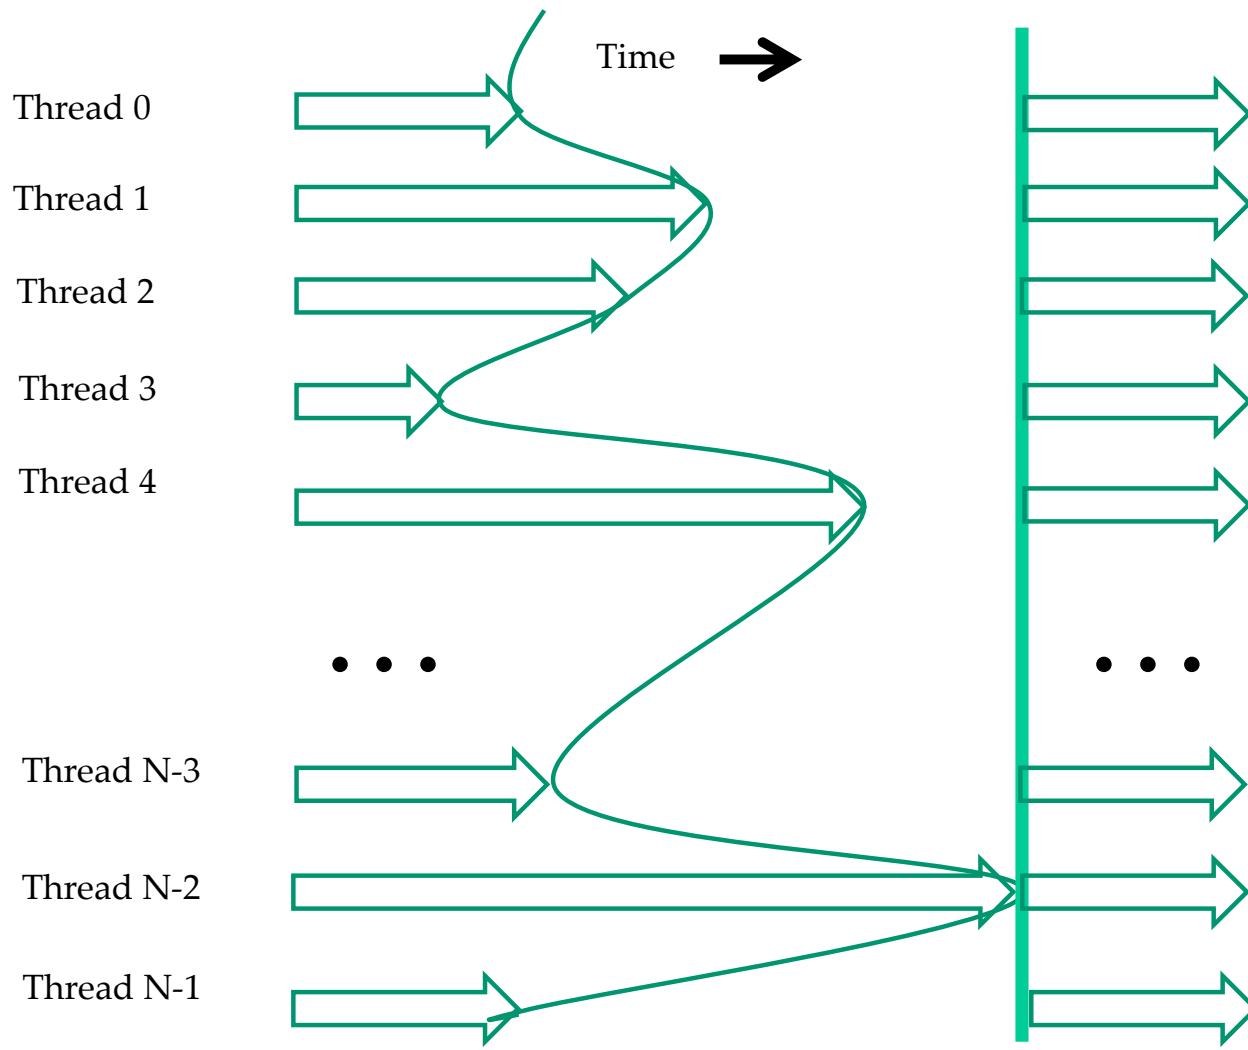
# Bulk Synchronous Steps Based on Barriers

- **How does it work?**

**Use a barrier** to wait for thread to ‘catch up.’

- A barrier is a synchronization point:

- **each thread calls a function** to enter barrier;
  - **threads block** (sleep) in barrier function  
**until all threads have called;**
  - **after last thread calls** function,  
**all threads continue** past the barrier.



# Use `__syncthreads` for CUDA Blocks

- How does it work in CUDA?  
Only **within thread blocks!**
- The function: `void __syncthreads(void);`
- N.B.
  - All threads in block must enter (no subsets).
  - All threads must enter the **SAME static call**  
(not the same as all threads calling function!).

# Barrier Trauma: What's Actually Done?

- **What exactly is guaranteed** to have finished?
  - Are **shared memory** operations before a barrier (e.g., stores) guaranteed to have completed?
  - What about **global memory** ops?
  - What about **atomic ops** with no return values?
  - What about I/O operations?
- CUDA manual: all global and shared memory ops (which presumably includes atomic variants) have completed.
- **Avoid assumptions about I/O** (such as printf).

# Tiled Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, int Width)
{
1.    __shared__ float subTileM[TILE_WIDTH][TILE_WIDTH];
2.    __shared__ float subTileN[TILE_WIDTH][TILE_WIDTH];

3.    int bx = blockIdx.x;    int by = blockIdx.y;
4.    int tx = threadIdx.x;  int ty = threadIdx.y;

        // Identify the row and column of the P element to work on
5.    int Row = by * TILE_WIDTH + ty; // note: blockDim.x == TILE_WIDTH
6.    int Col = bx * TILE_WIDTH + tx; //          blockDim.y == TILE_WIDTH
7.    float Pvalue = 0;

        // Loop over the M and N tiles required to compute the P element
        // The code assumes that the Width is a multiple of TILE_WIDTH!
8.    for (int m = 0; m < Width/TILE_WIDTH; ++m) {
            // Collaborative loading of M and N tiles into shared memory
9.        subTileM[ty][tx] = M[Row*Width + m*TILE_WIDTH+tx];
10.       subTileN[ty][tx] = N[(m*TILE_WIDTH+ty)*Width+Col];
11.       __syncthreads();
12.       for (int k = 0; k < TILE_WIDTH; ++k)
13.           Pvalue += subTileM[ty][k] * subTileN[k][tx];
14.       __syncthreads();
15.   }
16.   P[Row*Width+Col] = Pvalue;
}
```

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/  
University of Illinois at Urbana-Champaign

# Compare with Basic MM Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, int Width)
{
    // Calculate the row index of the P element and M
    int Row = blockIdx.y * blockDim.y + threadIdx.y;
    // Calculate the column index of P and N
    int Col = blockIdx.x * blockDim.x + threadIdx.x;

    if ((Row < Width) && (Col < Width)) {
        float Pvalue = 0;

        // each thread computes one element of the block sub-matrix
        for (int k = 0; k < Width; ++k)
            Pvalue += M[Row*Width+k] * N[k*Width+Col];

        P[Row*Width+Col] = Pvalue;
    }
}
```

# Use of Large Tiles Shifts Bottleneck

- Recall our example GPU: **1,000 GFLOP/s, 150 GB/s**
- **16x16 tiles** use each operand for 16 operations
  - **reduce global memory accesses by** a factor of **16**
  - **150GB/s** bandwidth supports  
 $(150/4)*16 = 600 \text{ GFLOPS!}$
- **32x32 tiles** use each operand for 32 operations
  - **reduce global memory accesses by** a factor of **32**
  - **150 GB/s** bandwidth supports  
 $(150/4)*32 = 1,200 \text{ GFLOPS!}$
  - **Memory bandwidth is no longer the bottleneck!**

## Also Need Parallel Accesses to Memory

- Shared memory size
  - implementation dependent
  - **64kB** per SM in Maxwell (48kB max per block)
- Given **TILE\_WIDTH of 16** (256 threads / block),
  - each thread block uses  
 $2 * 256 * 4B = 2kB$  of shared memory,
  - which limits active blocks to 32;
  - max. of 2048 threads per SM,
  - which limits blocks to 8.
  - Thus up to  $8 * 512 = \mathbf{4,096 \text{ pending loads}}$   
(2 per thread, 256 threads per block)

## Another Good Choice: 32x32 Tiles

- Given **TILE\_WIDTH of 32** (1,024 threads / block),
  - each thread block uses  $2 \times 1024 \times 4B = 8kB$  of shared memory,
  - which limits active blocks to 8;
  - max. of 2,048 threads per SM,
  - which limits blocks to 2.
  - Thus up to  $2 \times 2,048 = \mathbf{4,096 \text{ pending loads}}$   
(2 per thread, 1,024 threads per block)

**(same memory parallelism exposed)**

# Current GPU? Use Device Query

- Number of devices in the system

```
int dev_count;  
cudaGetDeviceCount( &dev_count);
```

- Capability of devices

```
cudaDeviceProp      dev_prop;  
for (i = 0; i < dev_count; i++) {  
    cudaGetDeviceProperties( &dev_prop, i);  
  
    // decide if device has sufficient resources and capabilities  
}
```

- `cudaDeviceProp` is a built-in C structure type
  - `dev_prop.dev_prop.maxThreadsPerBlock`
  - `Dev_prop.sharedMemoryPerBlock`
  - ...



**QUESTIONS?**

**READ CHAPTER 4!**

# Problem Solving

Barriers are needed to ensure that one thread's writes are visible to other threads—for communication between threads in a block.

Consider some code...

```
int32_t tx = threadIdx.x;  
int32_t ty = threadIdx.y;  
__shared__ float tile[TW][TW];  
tile[ty][tx] = myNumber;
```

Q: Do we need to add a call to `__syncthreads` here?

```
otherNumber = tile[ty][tx];
```

A: No `__syncthreads` call needed: each thread reads only the value that it wrote itself!

# Problem Solving

Here's a slightly different code...

```
int32_t tx = threadIdx.x;  
int32_t ty = threadIdx.y;  
__shared__ float tile[TW][TW];  
tile[ty][tx] = myNumber;
```

```
otherNumber = tile[tx][ty];
```

Q: Do we need to add a call to `__syncthreads` here?

A: Yes! Otherwise, the values copied into each thread's `otherNumber` variable may not be deterministic.

# Problem Solving

Let's extend that last code...

```
int32_t tx = threadIdx.x;  
int32_t ty = threadIdx.y;  
__shared__ float tile[TW][TW];  
tile[ty][tx] = myNumber;  
__syncthreads();  
otherNumber = tile[tx][ty];
```

Q: Do we need to add a call to `__syncthreads` here?

```
thirdNumber = tile[TW - tx - 1][ty];
```

A: No: no thread modifies `tile` after the first barrier.

# Problem Solving

And one final extension...

```
int32_t tx = threadIdx.x;  
int32_t ty = threadIdx.y;  
__shared__ float tile[TW][TW];  
tile[ty][tx] = myNumber;  
__syncthreads();  
otherNumber = tile[tx][ty];  
thirdNumber = tile[TW - tx - 1][ty];
```

Q: Do we need to add a call to `__syncthreads` here?

`tile[ty][tx] = myNumber * 2.0;`

A: Yes! Reads for `thirdNumber` must finish before other threads change `tile`.