

ECE408 / CS483 / CSE408
Summer 2025

Applied Parallel Programming

Lecture 6: More on Tiling



What Will You Learn Today?

- to handle boundary conditions in tiled algorithms

How to Handle Matrices of Other Sizes?

- Slide deck 5's tiled kernel
 - assumed integral number of tiles (thread blocks)
 - in all matrix dimensions.

How can we avoid this assumption?

- One answer: add padding, but not easy to reformat data, and adds transfer time.

Other ideas?

Let's Review Our Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, int Width)
{
1.    __shared__ float subTileM[TILE_WIDTH][TILE_WIDTH];
2.    __shared__ float subTileN[TILE_WIDTH][TILE_WIDTH];

3.    int bx = blockIdx.x;    int by = blockIdx.y;
4.    int tx = threadIdx.x;  int ty = threadIdx.y;


        // Identify the row and column of the P element to work on
5.    int Row = by * TILE_WIDTH + ty; // note: blockDim.x == TILE_WIDTH
6.    int Col = bx * TILE_WIDTH + tx; //          blockDim.y == TILE_WIDTH
7.    float Pvalue = 0;

        // Loop over the M and N tiles required to compute the P element
        // The code assumes that the Width is a multiple of TILE_WIDTH!
8.    for (int m = 0; m < Width/TILE_WIDTH; ++m) {
            // Collaborative loading of M and N tiles into shared memory
9.        subTileM[ty][tx] = M[Row*Width + m*TILE_WIDTH+tx];
10.       subTileN[ty][tx] = N[(m*TILE_WIDTH+ty)*Width+Col];
11.       __syncthreads();
12.       for (int k = 0; k < TILE_WIDTH; ++k)
13.           Pvalue += subTileM[ty][k] * subTileN[k][tx];
14.       __syncthreads();
15.   }

16.   P[Row*Width+Col] = Pvalue;
}
```

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign


Recall Second Tiles Loaded for Thread Block (0,0)



Thread Block (0,0) Computes on Shared Tiles (Iter 0)


Global Memory		
N _{0,0}	N _{0,1}	N _{0,2}
N _{1,0}	N _{1,1}	N _{1,2}
N _{2,0}	N _{2,1}	N _{2,2}

Each thread loads one value of N (from tile).



Threads use shared data in compute loop iteration 0.

Global Memory		
M _{0,0}	M _{0,1}	M _{0,2}
M _{1,0}	M _{1,1}	M _{1,2}
M _{2,0}	M _{2,1}	M _{2,2}




Each thread loads one value of M (from tile).

Thread Block (0,0) Computes on Shared Tiles (Iter 1)

Global Memory		
N _{0,0}	N _{0,1}	N _{0,2}
N _{1,0}	N _{1,1}	N _{1,2}
N _{2,0}	N _{2,1}	N _{2,2}

Neither are
tile values of N!

Global Memory		
M _{0,0}	M _{0,1}	M _{0,2}
M _{1,0}	M _{1,1}	M _{1,2}
M _{2,0}	M _{2,1}	M _{2,2}




Shared
Memory

Threads use shared
data in compute
loop iteration 1.

Tile values of
M are not
defined!

Global
Memory

Let's Look at the First Tile for Block(1,1) Next




Thread Block (1,1) Computes on Shared Tiles (Iter 0)

Global Memory		
N _{0,0}	N _{0,1}	N _{0,2}
N _{1,0}	N _{1,1}	N _{1,2}
N _{2,0}	N _{2,1}	N _{2,2}

RED and
PURPLE use
undefined N
value!

M _{0,0}	M _{0,1}	M _{0,2}
M _{1,0}	M _{1,1}	M _{1,2}
M _{2,0}	M _{2,1}	M _{2,2}

Global
Memory



P _{0,0}	P _{0,1}	P _{0,2}	
P _{1,0}	P _{1,1}	P _{1,2}	
P _{2,0}	P _{2,1}	P _{2,2}	

Global
Memory

Threads use shared
data in compute
loop iteration 0.



ORANGE and
PURPLE use
undefined M
value!

Thread Block (1,1) Computes on Shared Tiles (Iter 1)

Global Memory		
N _{0,0}	N _{0,1}	N _{0,2}
N _{1,0}	N _{1,1}	N _{1,2}
N _{2,0}	N _{2,1}	N _{2,2}

RED and
PURPLE use
undefined N
value!

Global Memory		
M _{0,0}	M _{0,1}	M _{0,2}
M _{1,0}	M _{1,1}	M _{1,2}
M _{2,0}	M _{2,1}	M _{2,2}



Threads use shared
data in compute
loop iteration 1.

ORANGE and
PURPLE use
undefined M
value!

Global
Memory

Major Cases in Toy Example

- Threads that calculate valid P elements but can step outside valid input
 - Second tile of Block(0,0), all threads when k is 1
- Threads that do not calculate valid P elements
 - Block(1,1), Thread(1,0), non-existent row
 - Block(1,1), Thread(0,1), non-existent column
 - Block(1,1), Thread(1,1), non-existent row and column

Solution: Write 0 for Missing Elements


- Test during tile load:
is **target within input matrix?**
 - **If yes**, proceed to **load**;
 - **otherwise**, just **write 0** to shared memory.
- The **benefit**?
 - **No specialization during tile use!**
 - Multiplying by 0 guarantees that unwanted terms do not contribute to the inner product.

Thread Block (0,0) Computes on Shared Tiles (Iter 1)

Global Memory		
N _{0,0}	N _{0,1}	N _{0,2}
N _{1,0}	N _{1,1}	N _{1,2}
N _{2,0}	N _{2,1}	N _{2,2}

Tile values of 0 have no effect on sum.

Global Memory		
M _{0,0}	M _{0,1}	M _{0,2}
M _{1,0}	M _{1,1}	M _{1,2}
M _{2,0}	M _{2,1}	M _{2,2}



Threads use shared data in compute loop iteration 1.

Tile values of 0 have no effect on sum.

Global
Memory

What About Threads Outside of P?

- If a **thread is not within P**,
 - All terms in sum are 0.
 - No harm in performing FLOPs.
 - No harm in writing to registers.
 - **Must not be allowed to write to global memory!**


So: **Threads outside of P calculate 0,
but store nothing.**

Thread Block (1,1) Computes on Shared Tiles (Iter 1)

Global Memory		
N _{0,0}	N _{0,1}	N _{0,2}
N _{1,0}	N _{1,1}	N _{1,2}
N _{2,0}	N _{2,1}	N _{2,2}

All but P_{2,2} computed as 0.

Global Memory		
M _{0,0}	M _{0,1}	M _{0,2}
M _{1,0}	M _{1,1}	M _{1,2}
M _{2,0}	M _{2,1}	M _{2,2}



Global Memory		
P _{0,0}	P _{0,1}	P _{0,2}
P _{1,0}	P _{1,1}	P _{1,2}
P _{2,0}	P _{2,1}	P _{2,2}

Threads use shared data in compute loop iteration 1.

All but P_{2,2} computed as 0.

Modifying the Tile Count

```
8.  for (int m = 0; m < Width/TILE_WIDTH; ++m) {
```

The bound for **m** implicitly assumes that **Width** is a multiple of **TILE_WIDTH**. We need to round up.

```
for (int m = 0; m < (Width - 1)/TILE_WIDTH + 1; ++m) {
```

For non-multiples of **TILE_WIDTH**:

- quotient is unchanged;
- add one to round up.

For multiples of **TILE_WIDTH**:

- quotient is now one smaller,
- but we add 1.

Modifying the Tile Loading Code

We had ...

```
// Collaborative loading of M and N tiles into shared memory  
9.    subTileM[ty][tx] = M[Row*Width + m*TILE_WIDTH+tx];  
10.   subTileN[ty][tx] = N[(m*TILE_WIDTH+ty)*Width+Col];
```

Note: the tests for M and N tiles are NOT the same.

```
if (Row < Width && m*TILE_WIDTH+tx < Width) {  
    // as before  
    subTileM[ty][tx] = M[Row*Width + m*TILE_WIDTH+tx];  
} else {  
    subTileM[ty][tx] = 0;  
}
```

And for Loading N...

We had ...

```
// Collaborative loading of M and N tiles into shared memory  
9.    subTileM[ty][tx] = M[Row*Width + m*TILE_WIDTH+tx];  
10.   subTileN[ty][tx] = N[(m*TILE_WIDTH+ty)*Width+Col];
```

Note: the tests for M and N tiles are NOT the same.

```
if (m*TILE_WIDTH+ty < Width && Col < Width) {  
    // as before  
    subTileN[ty][tx] = N[(m*TILE_WIDTH+ty)*Width+Col];  
} else {  
    subTileN[ty][tx] = 0;  
}
```

Modifying the Tile Use Code

We had ...

```
12. for (int k = 0; k < TILE_WIDTH; ++k)
13.     Pvalue += subTileM[ty][k] * subTileN[k][tx];
```

Note: **no changes are needed**, but we might save a little energy (fewer floating-point ops)?

```
if (Row < Width && Col < Width) {
    // as before
    for (int k = 0; k < TILE_WIDTH; ++k)
        Pvalue += subTileM[ty][k] * subTileN[k][tx];
}
```

Modifying the Write to P

We had ...

```
16. P[Row*Width+Col] = Pvalue;
```

We must test for threads outside of P:

```
if (Row < Width && Col < Width) {  
    // as before  
    P[Row*Width+Col] = Pvalue;  
}
```



Some Important Points

- For each thread, conditions are different for
 - Loading M element
 - Loading N element
 - Calculation/storing output elements
- Branch divergence
 - affects only blocks on boundaries, and
 - should be small for large matrices.
- What about rectangular matrices?



QUESTIONS?

READ CHAPTER 4!