



ECE408 / CS483 / CSE408
Summer 2025

Applied Parallel Programming

Lecture 12: Parallel Computation
Patterns – Parallel Scan (Prefix Sum)



What Will You Learn Today?

- parallel scan (prefix sum) algorithms based on reductions and reverse reductions
- the concept of double buffering
- tradeoffs between work efficiency and latency
- how to develop hierarchical algorithms (across multiple kernels)



Scan Includes all Partial Results

Reductions are a simplified form of scans.

In scan / parallel prefix,

- we need all of the partial sums
- (or whatever the operator might be).

(Inclusive) Scan (Prefix-Sum) Definition

Definition: *The scan operation takes a binary associative operator \oplus , and an array of n elements*

$$[x_0, x_1, \dots, x_{n-1}],$$

and returns the prefix-sum array

$$[x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-1})].$$

Example: If \oplus is addition, the scan operation
on the array [3 1 7 0 4 1 6 3],
returns [3 4 11 11 15 16 22 25].

Example: Sharing a Big Sandwich

You order a 100-inch sandwich to feed 10 people,
and you know how much each person wants in inches:

[3 5 2 7 28 4 3 0 8 1].

How do you cut the bread quickly?

How much of the sandwich is left over?

Method 1: sequentially!

Cut 3 inches, then cut 5 inches, then ...

Method 2: **calculate cutting offsets with prefix-sum**

[3, 8, 10, 17, 45, 49, 52, 52, 60, 61] (39 inches left)

Typical Applications of Scan

A simple and useful parallel building block.

Convert sequential recurrences

```
for (j = 1; j < n; j++)  
    out[j] = out[j-1] + f(j);
```

into parallel:

```
forall (j) { temp[j] = f(j); }  
scan (out, temp);
```

Typical Applications of Scan

- Useful for many parallel algorithms:
 - radix sort
 - quicksort
 - String comparison
 - Lexical analysis
 - Stream compaction
 - Polynomial evaluation
 - Solving recurrences
 - Tree operations
 - Histograms
 - Etc.



Other Applications

- Assigning camp slots
- Assigning farmer market space
- Allocating memory to parallel threads
- Allocating memory buffer to communication channels
- ...

An Inclusive Sequential Scan

Given a sequence $[x_0, x_1, x_2, \dots]$

Calculate output $[y_0, y_1, y_2, \dots]$

Such that

$$y_0 = x_0$$

$$y_1 = x_0 + x_1$$

$$y_2 = x_0 + x_1 + x_2$$

...

Using a recursive definition

$$y_i = y_{i-1} + x_i$$

A Sequential C Implementation

```
y[0] = x[0];  
for (i = 1; i < Max_i; i++)  
    y[i] = y[i-1] + x[i];
```

Computationally efficient:

N additions needed for N elements - O(N).

A Naïve Inclusive Parallel Scan

- Assign one thread to calculate each y element
- Have every thread to add up all x elements needed for the y element

$$y_0 = x_0$$

$$y_1 = x_0 + x_1$$

$$y_2 = x_0 + x_1 + x_2$$

“Parallel programming is easy as long
as you do not care about performance.”

Parallel Inclusive Scan using Reduction Trees

Calculate each output element as the reduction of all previous elements

- Some reduction partial sums will be shared among the calculation of output elements
- Based on hardware added design by Peter Kogge and Harold Stone at IBM in the 1970s – Kogge-Stone Trees
- Goal: low latency

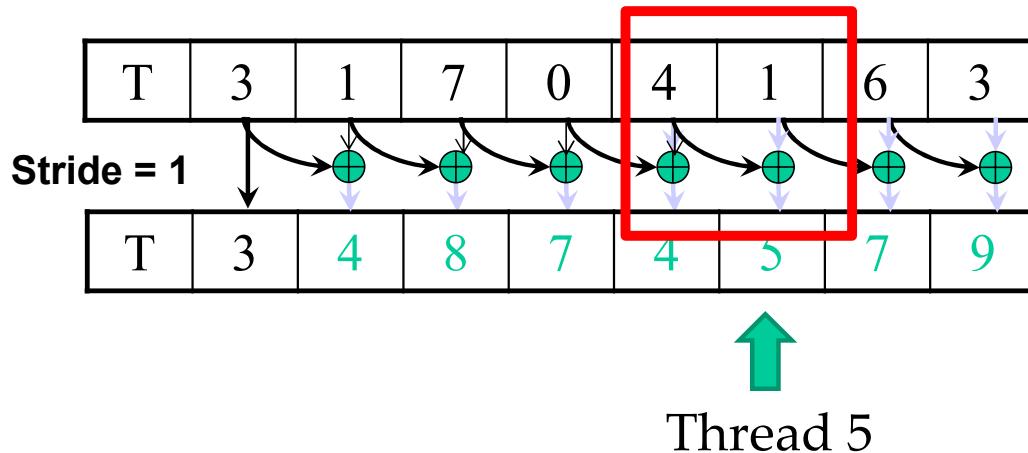
A Kogge-Stone Parallel Scan Algorithm

T	3	1	7	0	4	1	6	3
---	---	---	---	---	---	---	---	---

1. Load input from global memory into shared memory array T

Each thread loads one value from the input (global memory) array into shared memory array T.

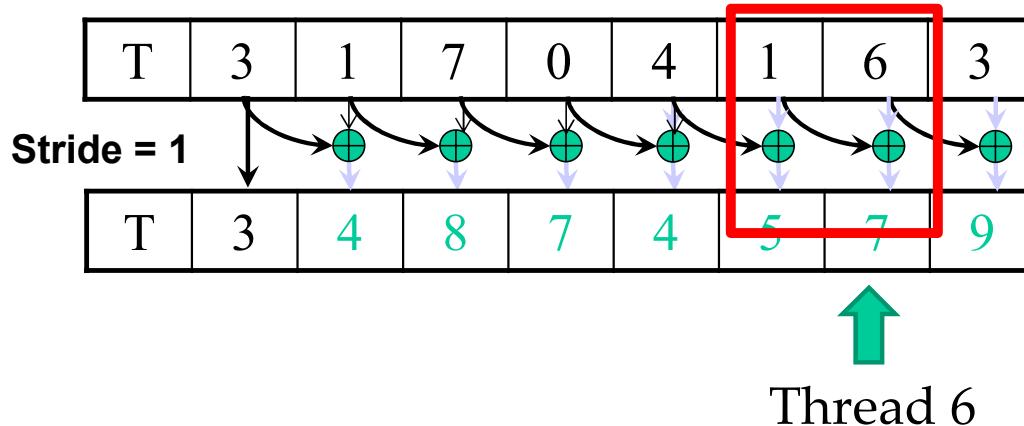
A Kogge-Stone Parallel Scan Algorithm



1. (previous slide)
2. Assuming n is a power of 2. Iterate $\log(n)$ times, stride from 1 to $n/2$. Threads *stride* to $n-1$ active: add pairs of elements that are *stride* elements apart.

- Iteration #1
Stride = 1
- Active threads: *stride* to $n-1$ ($n - \text{stride}$ active threads)
 - Thread j adds elements $T[j]$ and $T[j-\text{stride}]$ and writes result into element $T[j]$
 - Each iteration requires two syncthreads
 - make sure that input is in place
 - make sure that all input elements have been used

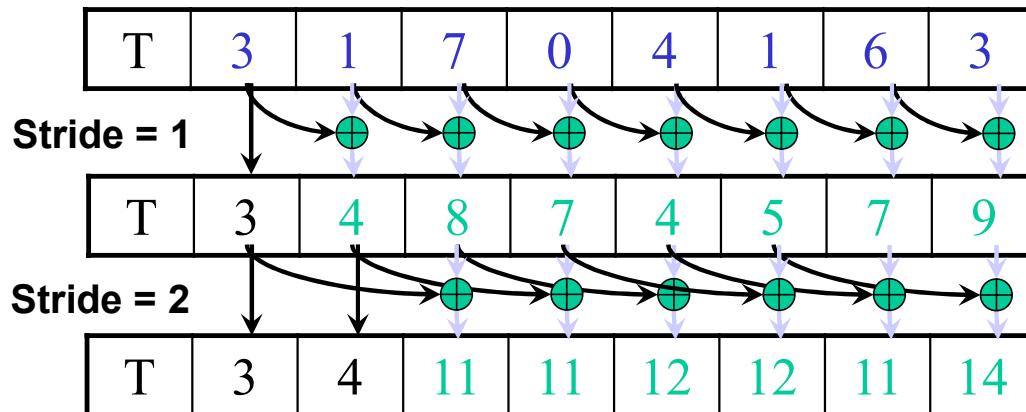
A Kogge-Stone Parallel Scan Algorithm



1. (previous slide)
2. Assuming n is a power of 2. Iterate $\log(n)$ times, stride from 1 to $n/2$. Threads *stride* to $n-1$ active: add pairs of elements that are *stride* elements apart.

- Iteration #1
Stride = 1
- Active threads: *stride* to $n-1$ ($n - \text{stride}$ active threads)
 - Thread j adds elements $T[j]$ and $T[j-\text{stride}]$ and writes result into element $T[j]$
 - Each iteration requires two syncthreads
 - `syncthreads();` // make sure that input is in place
 - `float temp = T[j] + T[j-stride];`
 - `syncthreads();` // make sure that previous output has been consumed
 - `T[j] = temp;`

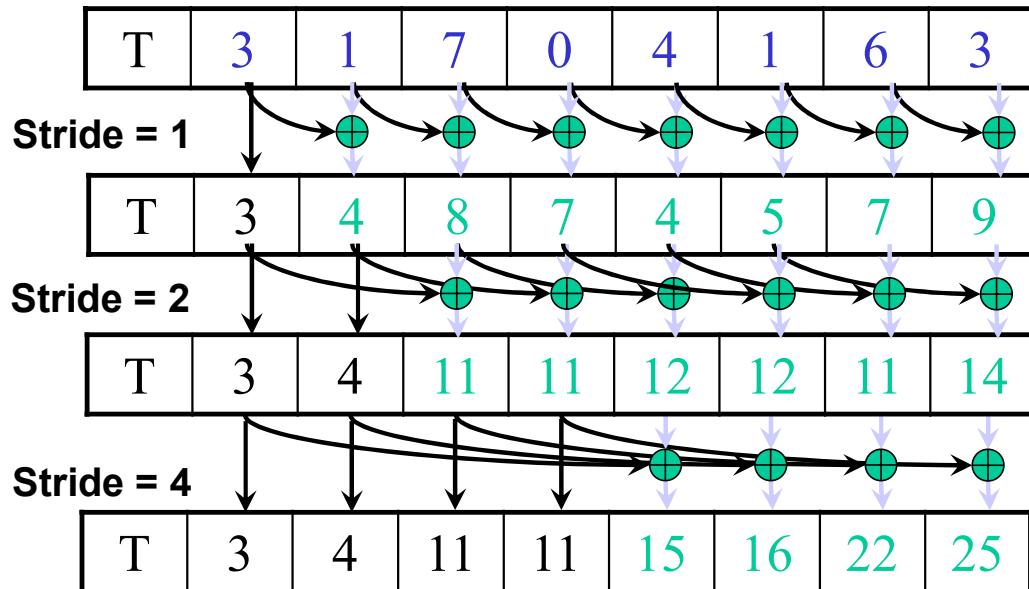
A Kogge-Stone Parallel Scan Algorithm



1. ...
2. Assuming n is a power of 2.
Iterate $\log(n)$ times, stride from 1 to $n/2$. Threads *stride* to $n-1$ *active*: add pairs of elements that are *stride* elements apart.

Iteration #2
Stride = 2

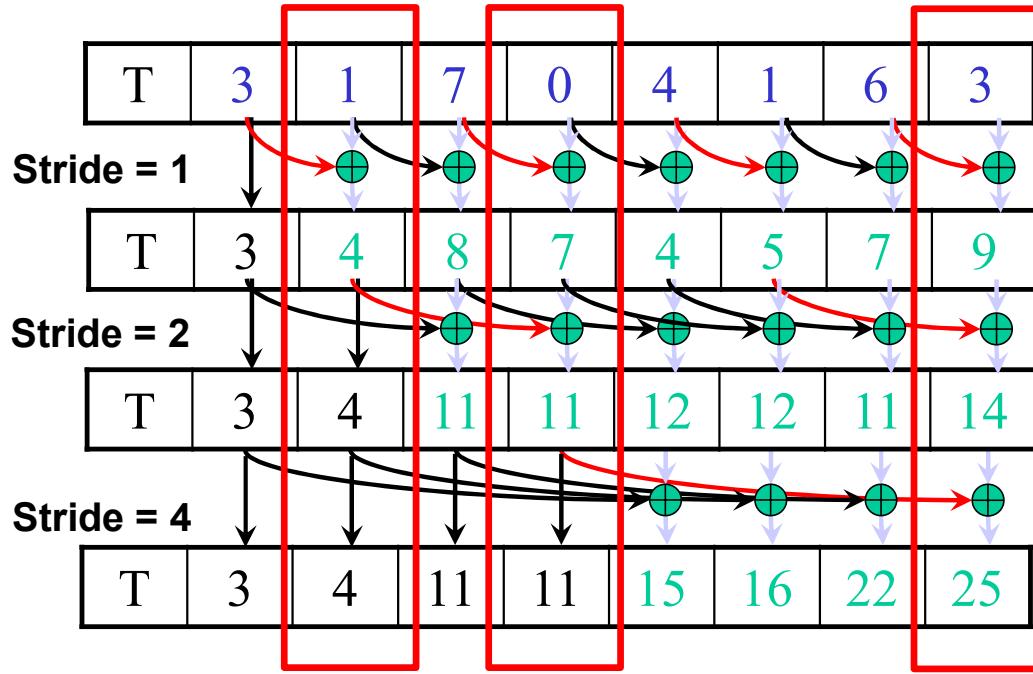
A Kogge-Stone Parallel Scan Algorithm



Iteration #3
Stride = 4

1. ...
2. ...
3. Write output from shared memory to device memory

Sharing Computation in Kogge-Stone



Iteration #3
Stride = 4

(Incorrect) Implementation with Single Barrier

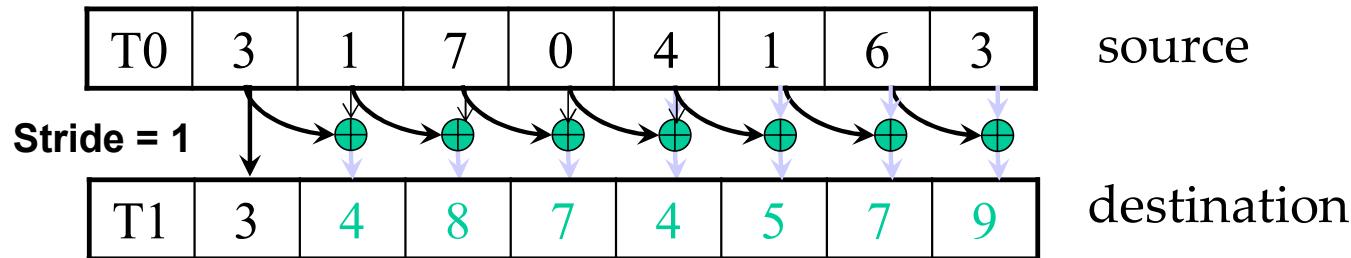
```
__global__
void Kogge_Stone_scan_kernel(float *X, float *Y, int InputSize)
{
    __shared__ float T[SECTION_SIZE];
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < InputSize) T[threadIdx.x] = X[i];

    for (unsigned int stride = 1; stride < blockDim.x; stride *= 2) {
        __syncthreads();
        if (threadIdx.x >= stride) // This code has a data race condition
            T[threadIdx.x] += T[threadIdx.x-stride];
    }
    Y[i] = T[threadIdx.x];
}
```

Double Buffering

- Use two copies of data T0 and T1
- Start by using T0 as input and T1 as output
- Switch input/output roles after each iteration
 - Iteration 0: T0 as input and T1 as output
 - Iteration 1: T1 as input and T0 and output
 - Iteration 2: T0 as input and T1 as output
- This is typically implemented with two pointers, *source* and *destination* that swap their contents from one iteration to the next
- This eliminates the need for the second `__syncthreads()` call

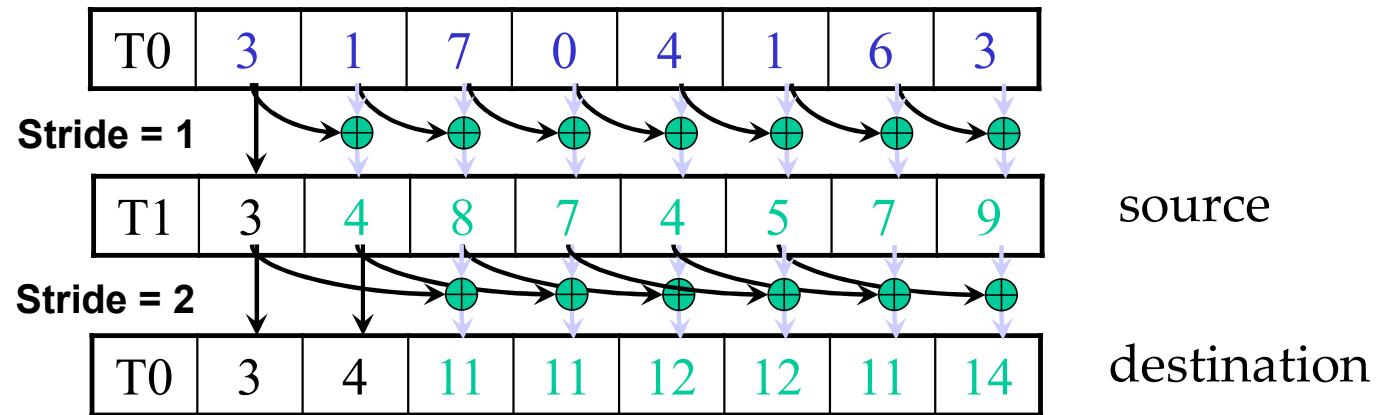
A Double-Buffered Kogge-Stone Parallel Scan Algorithm



Iteration #1
Stride = 1

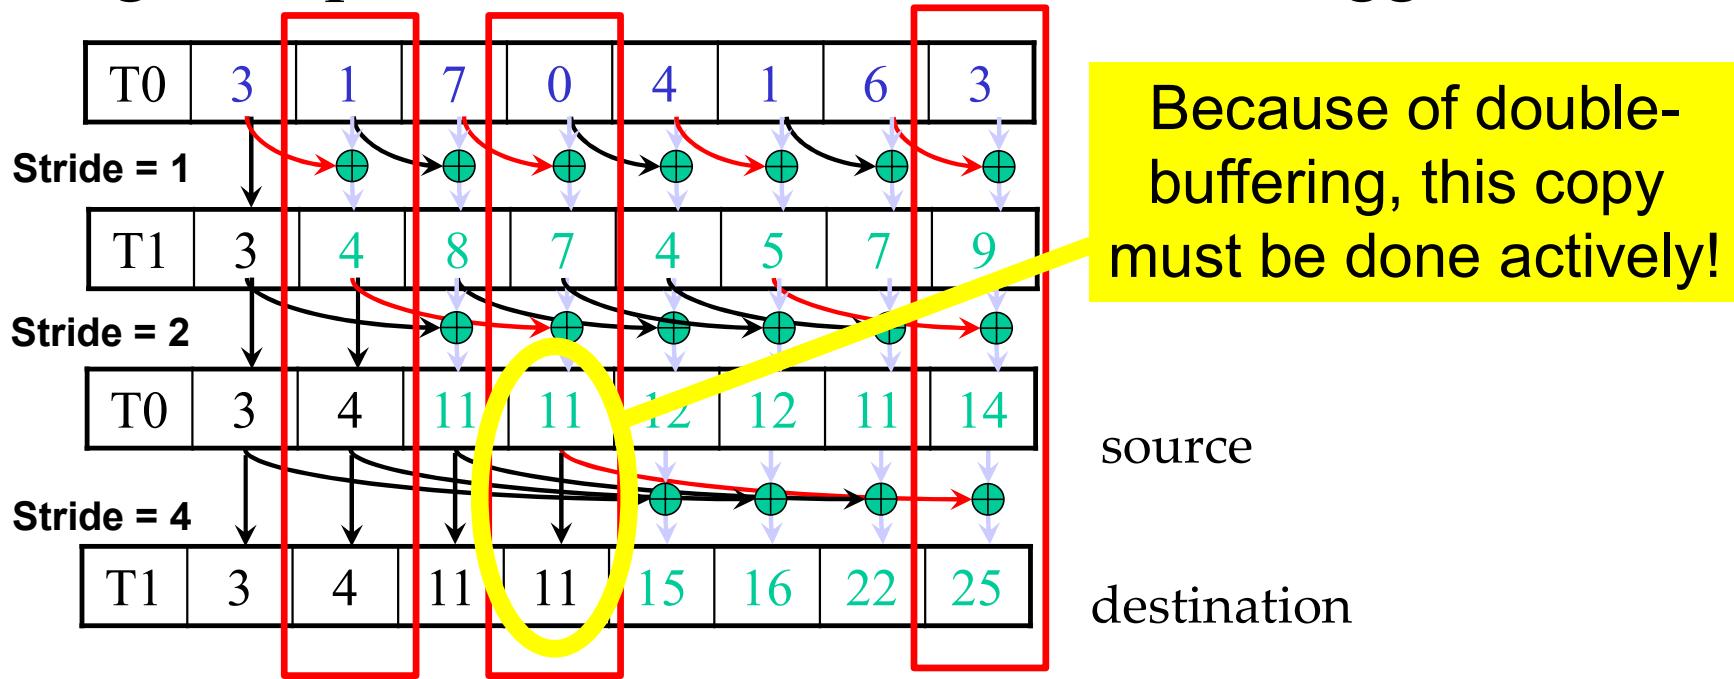
- source = &T0[0]; destination = &T1[0];
- Each iteration requires only one syncthreads()
 - syncthreads(); // make sure that input is in place
 - float destination[j] = source[j] + source[j-stride];
 - temp = destination; destination = source; source = temp;
- After the loop, write source (swapped) contents to global memory

A Double-Buffered Kogge-Stone Parallel Scan Algorithm



Iteration #2
Stride = 2

Sharing Computation in Double-Buffed Kogge-Stone



Iteration #3
Stride = 4

- Each iteration requires only one syncthreads()
 - `syncthreads(); // make sure that input is in place`
 - `float destination[j] = source[j] + source[j-stride];`
 - `temp = destination; destination = source; source = temp;`
- After the loop, write source (swapped) contents to global memory

Work Efficiency Analysis

- A Kogge-Stone scan kernel executes $\log(n)$ parallel iterations
 - The steps do $(n-1), (n-2), (n-4), \dots, (n - n/2)$ add operations each
 - Total # of add operations: $n * \log(n) - (n-1) \rightarrow O(n * \log(n))$ work
- This scan algorithm is not very work efficient
 - Sequential scan algorithm does n adds
 - A factor of $\log(n)$ hurts: 20x for 1,000,000 elements!
 - Typically used within each block, where $n \leq 1,024$
- A parallel algorithm can be slow when execution resources are saturated due to low work efficiency.

Improving Efficiency

A common parallel algorithm pattern: *Balanced Trees*

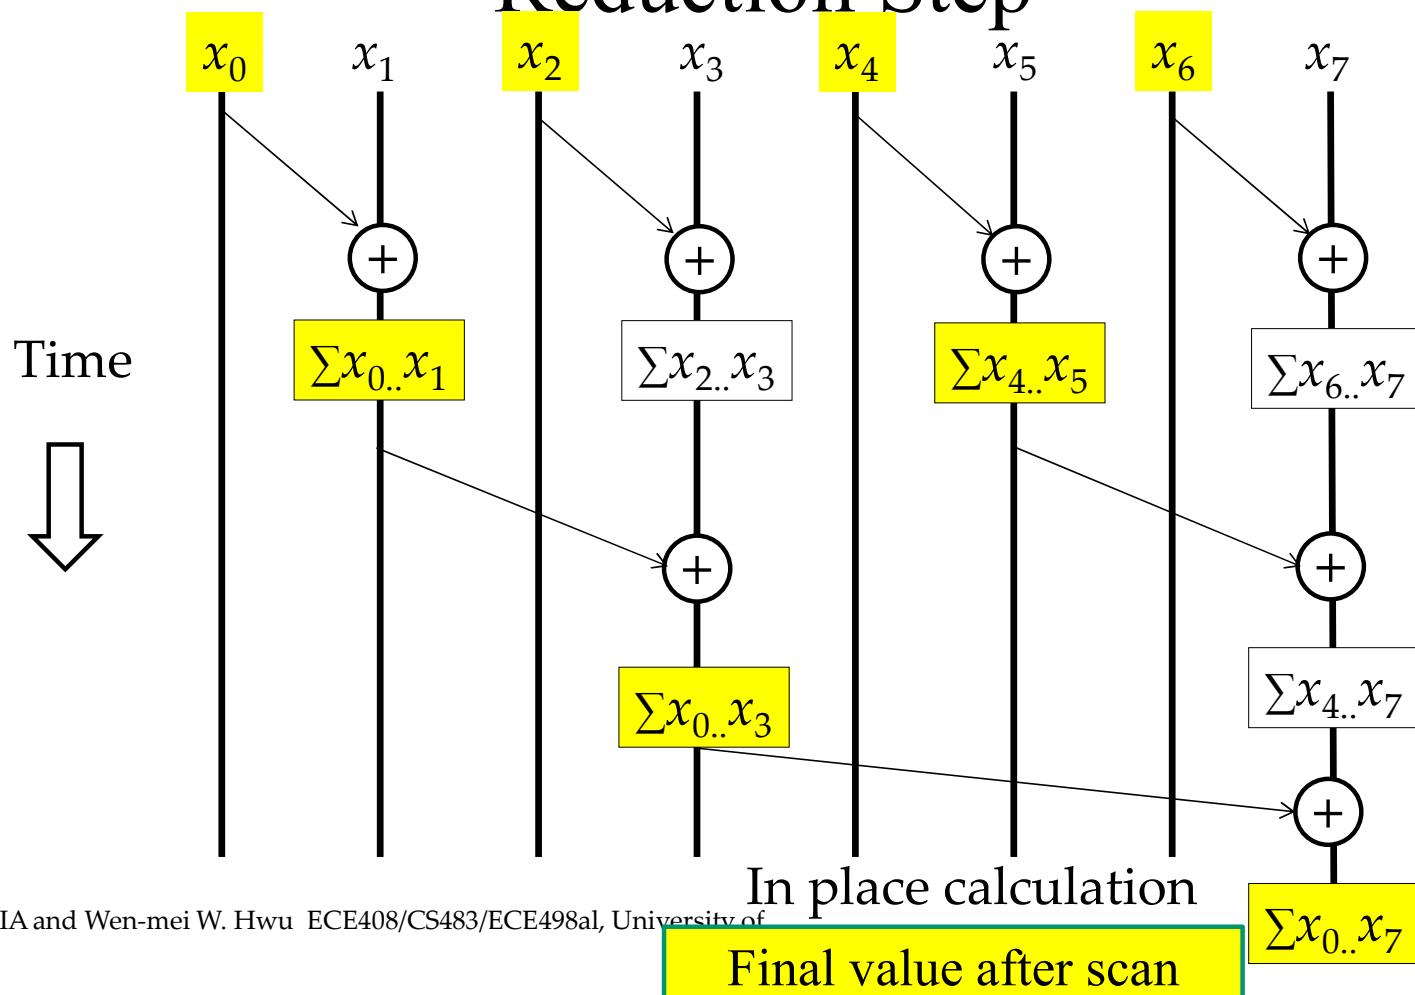
- Build balanced binary tree on input data
- Sweep it to and from the root
- Tree is not an actual data structure, but a concept to determine what each thread does at each step

For scan:

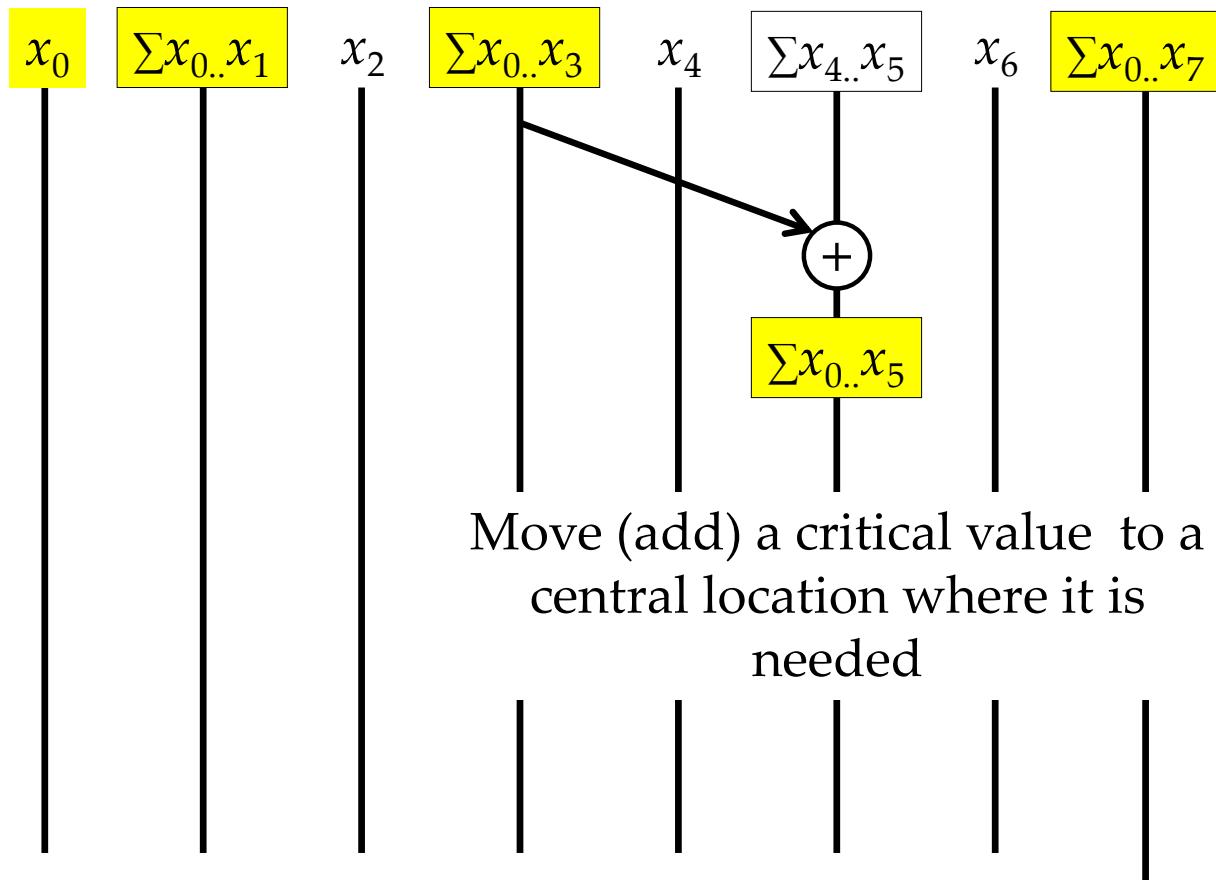
- Traverse down from leaves to root building partial sums at internal nodes in the tree
- Root holds sum of all leaves
- Traverse back up the tree building the scan from the partial sums

Brent-Kung Parallel Scan

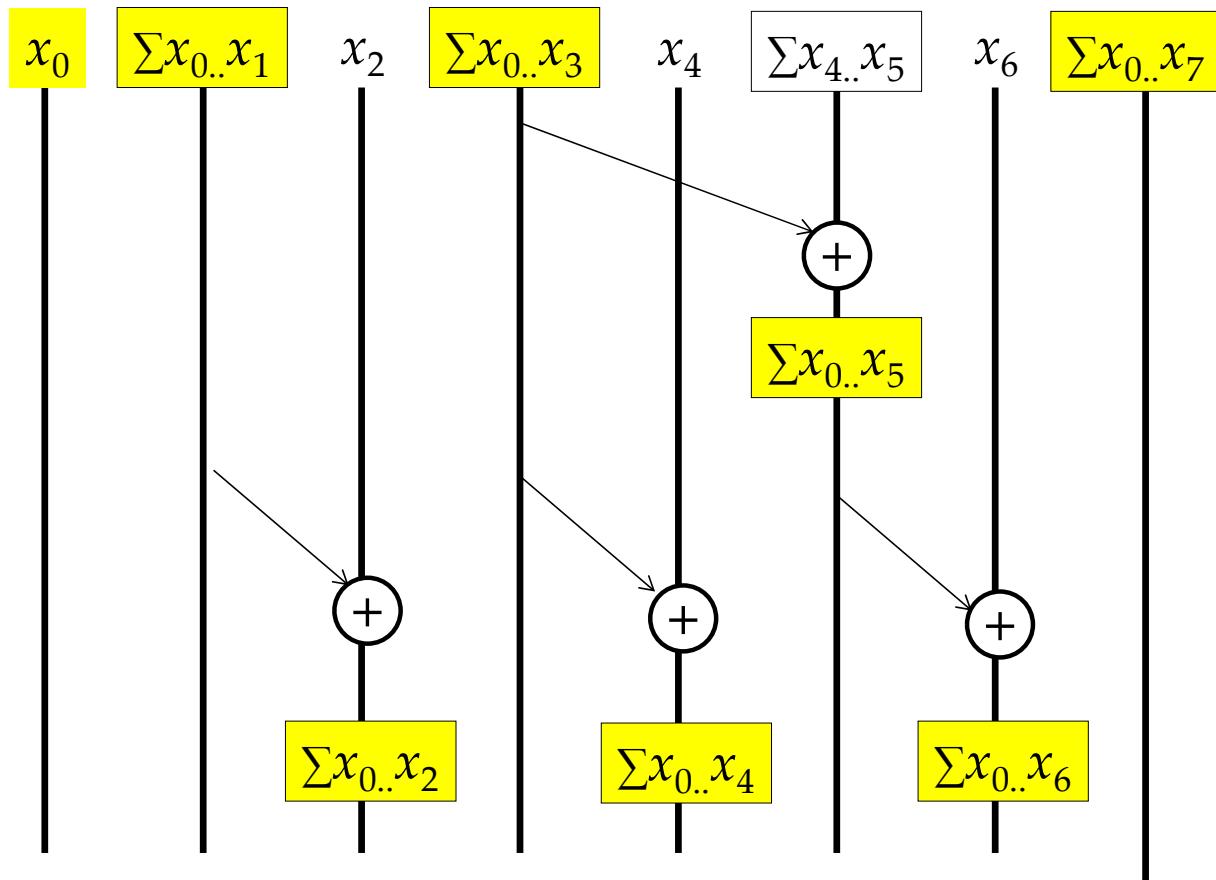
- Reduction Step



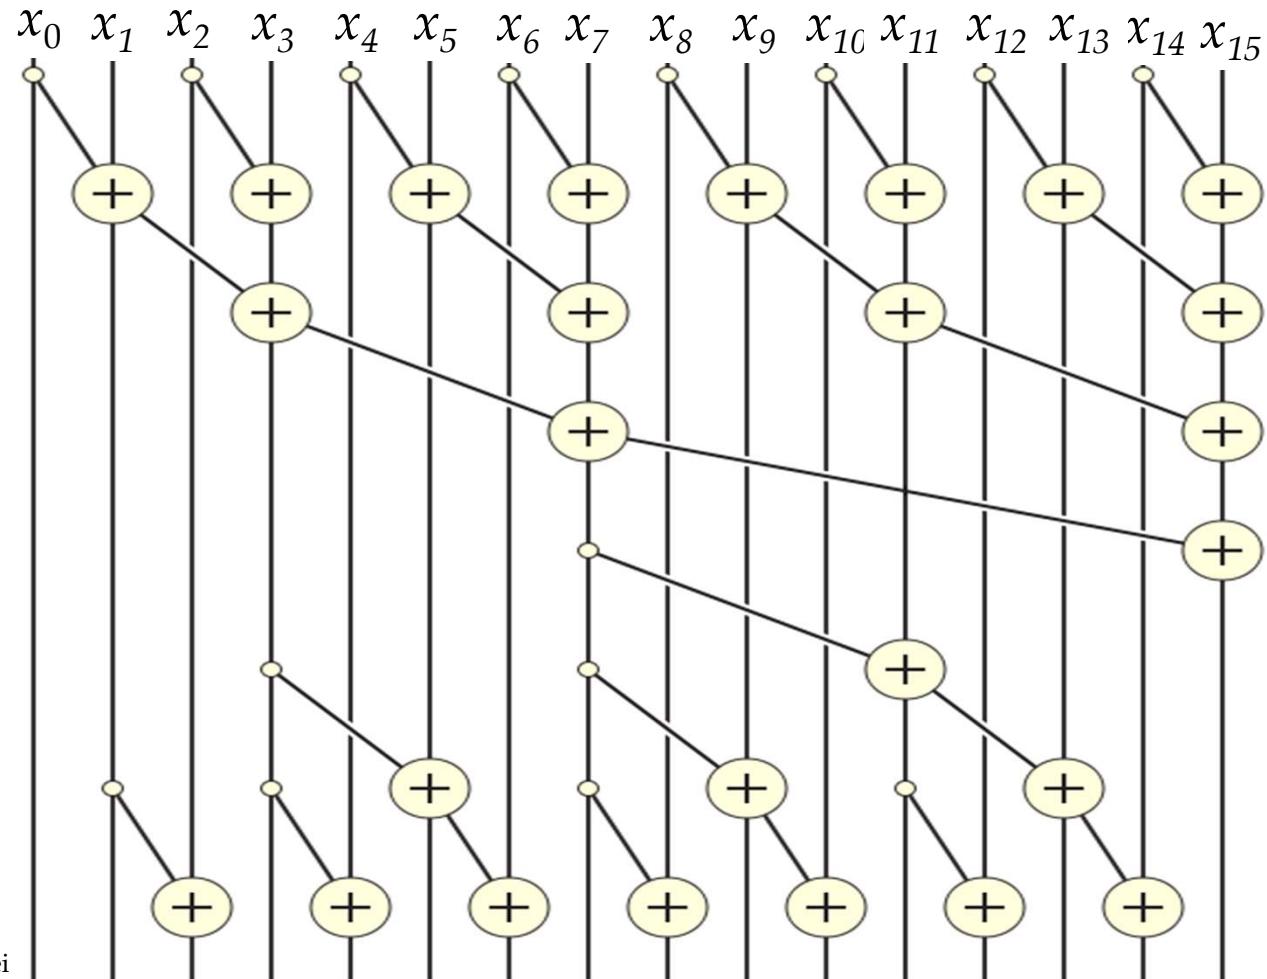
Inclusive Post-Scan Step



Inclusive Post Scan Step



Putting it Together (Data View)



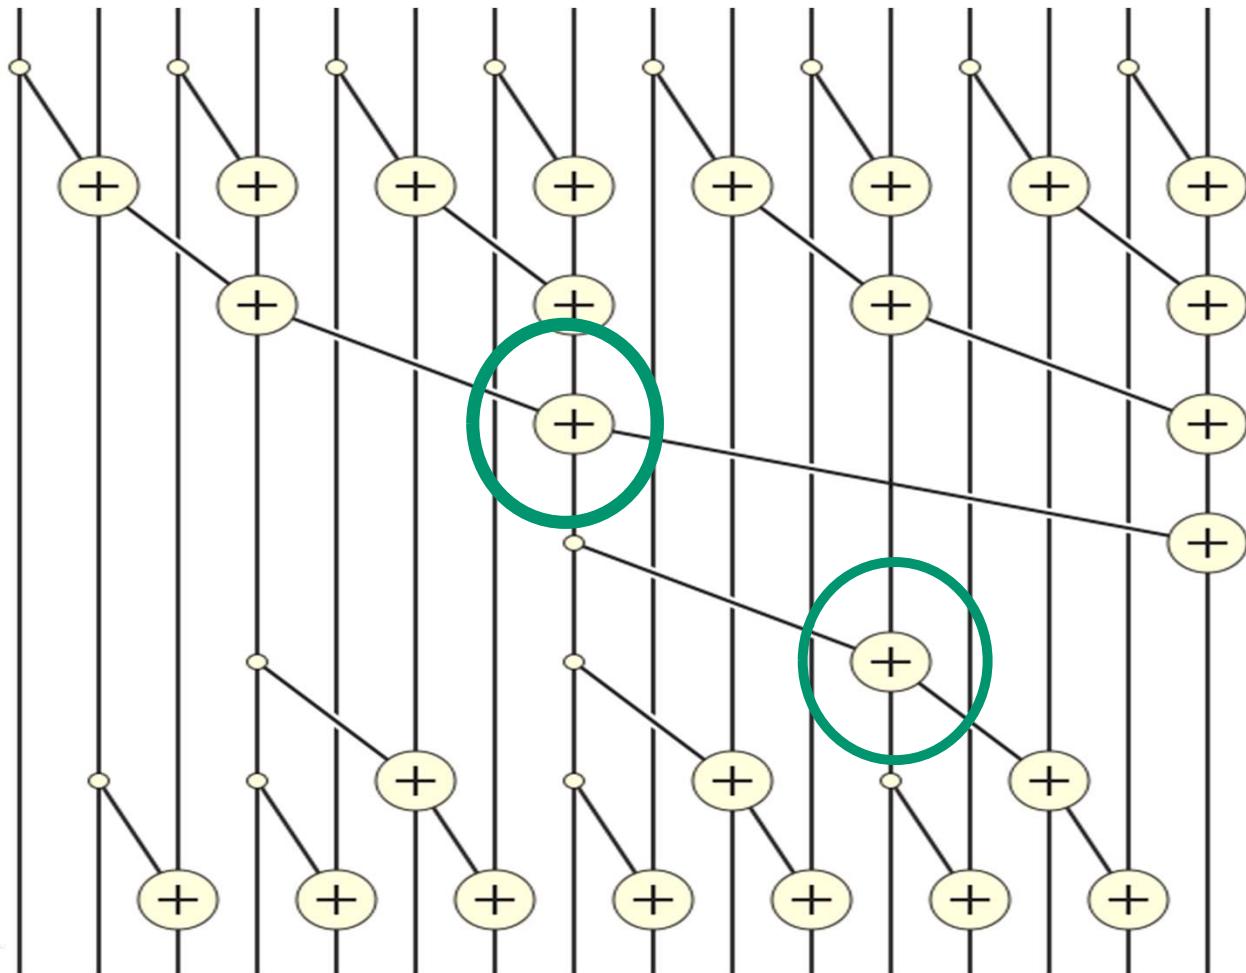
Reduction Step Kernel Code

```
// float T[2*BLOCK_SIZE] is in shared memory
// for previous slide, BLOCK_SIZE is 8

int stride = 1;
while(stride < 2*BLOCK_SIZE) {
    __syncthreads();
    int index = (threadIdx.x+1)*stride*2 - 1;
    if(index < 2*BLOCK_SIZE && (index-stride) >= 0)
        T[index] += T[index-stride];
    stride = stride*2;
}

// In our example,
// threadIdx.x+1      = 1, 2, 3, 4, 5, 6, 7, 8
// stride = 1, index = 1, 3, 5, 7, 9, 11, 13, 15
```

Putting it Together



Post Scan Step (Distribution Tree)

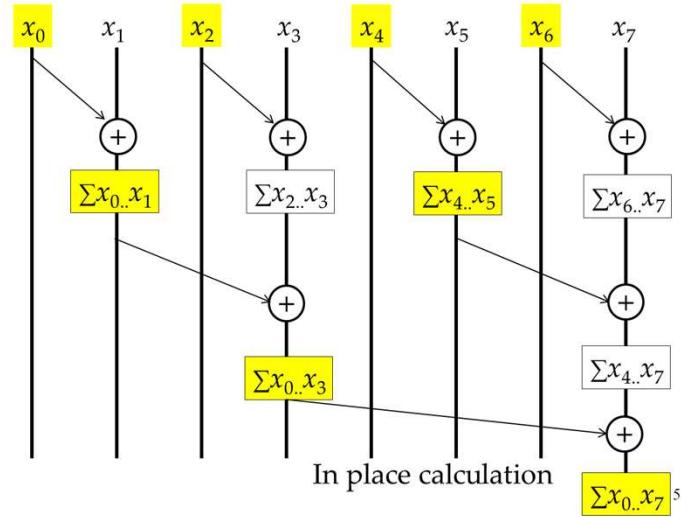
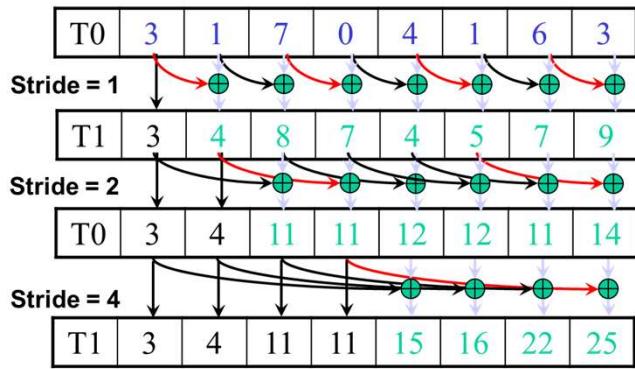
```
int stride = BLOCK_SIZE/2;
while(stride > 0) {
    __syncthreads();
    int index = (threadIdx.x+1)*stride*2 - 1;
    if ((index+stride) < 2*BLOCK_SIZE)
        T[index+stride] += T[index];
    stride = stride / 2;
}
```

```
// In our example,
// BLOCK_SIZE=8 stride=4, 2, 1
// for first iteration, active thread = 0 index = 7, stride = 11
```

Work Analysis

- The parallel Scan executes $2*\log(n)$ parallel iterations
 - $\log(n)$ in reduction and $\log(n)$ in post scan
 - The iterations do $n/2, n/4,..1, (2-1), \dots, (n/4-1), (n/2-1)$ useful adds
 - In our example, $n = 16$, the number of useful adds is $16/2 + 16/4 + 16/8 + 16/16 + (16/8-1) + (16/4-1) + (16/2-1)$
 - Total adds: $(n-1) + (n-2) - (\log(n) - 1) = 2*(n-1) - \log(n) \rightarrow O(n)$ work
- The total number of adds is no more than twice of that done in the efficient sequential algorithm
 - The benefit of parallelism can easily overcome the $2\times$ work when there is sufficient hardware

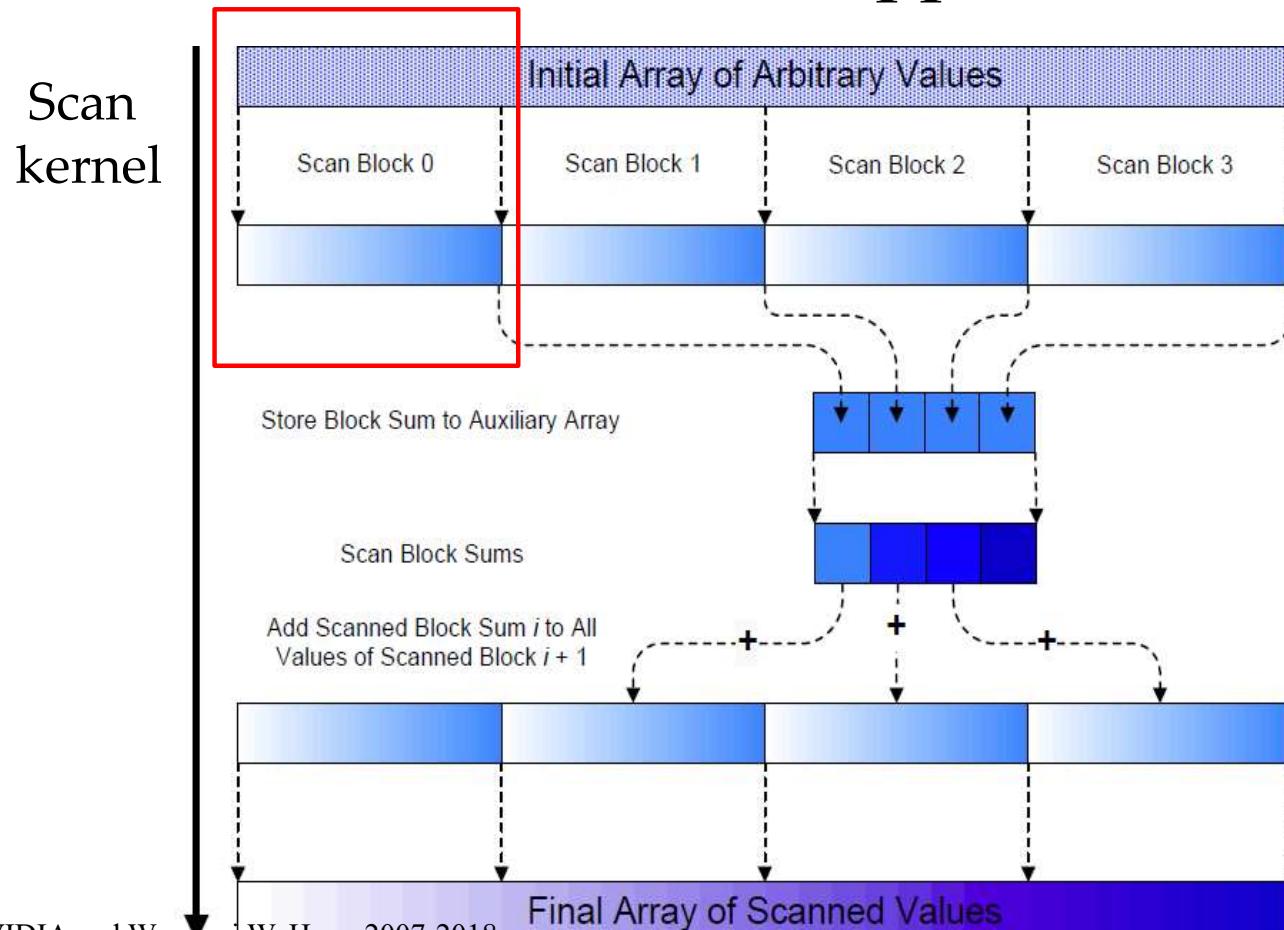
Kogge-Stone vs. Brent-Kung



- Brent-Kung uses half the number of threads compared to Kogge-Stone
 - Each thread should load two elements into the shared memory
- Brent-Kung takes twice the number of steps compared to Kogge-Stone
 - Kogge-Stone is more popular for parallel scan with blocks in GPUs

Overall Flow of Complete Scan

A Hierarchical Approach



Using Global Memory Contents in CUDA

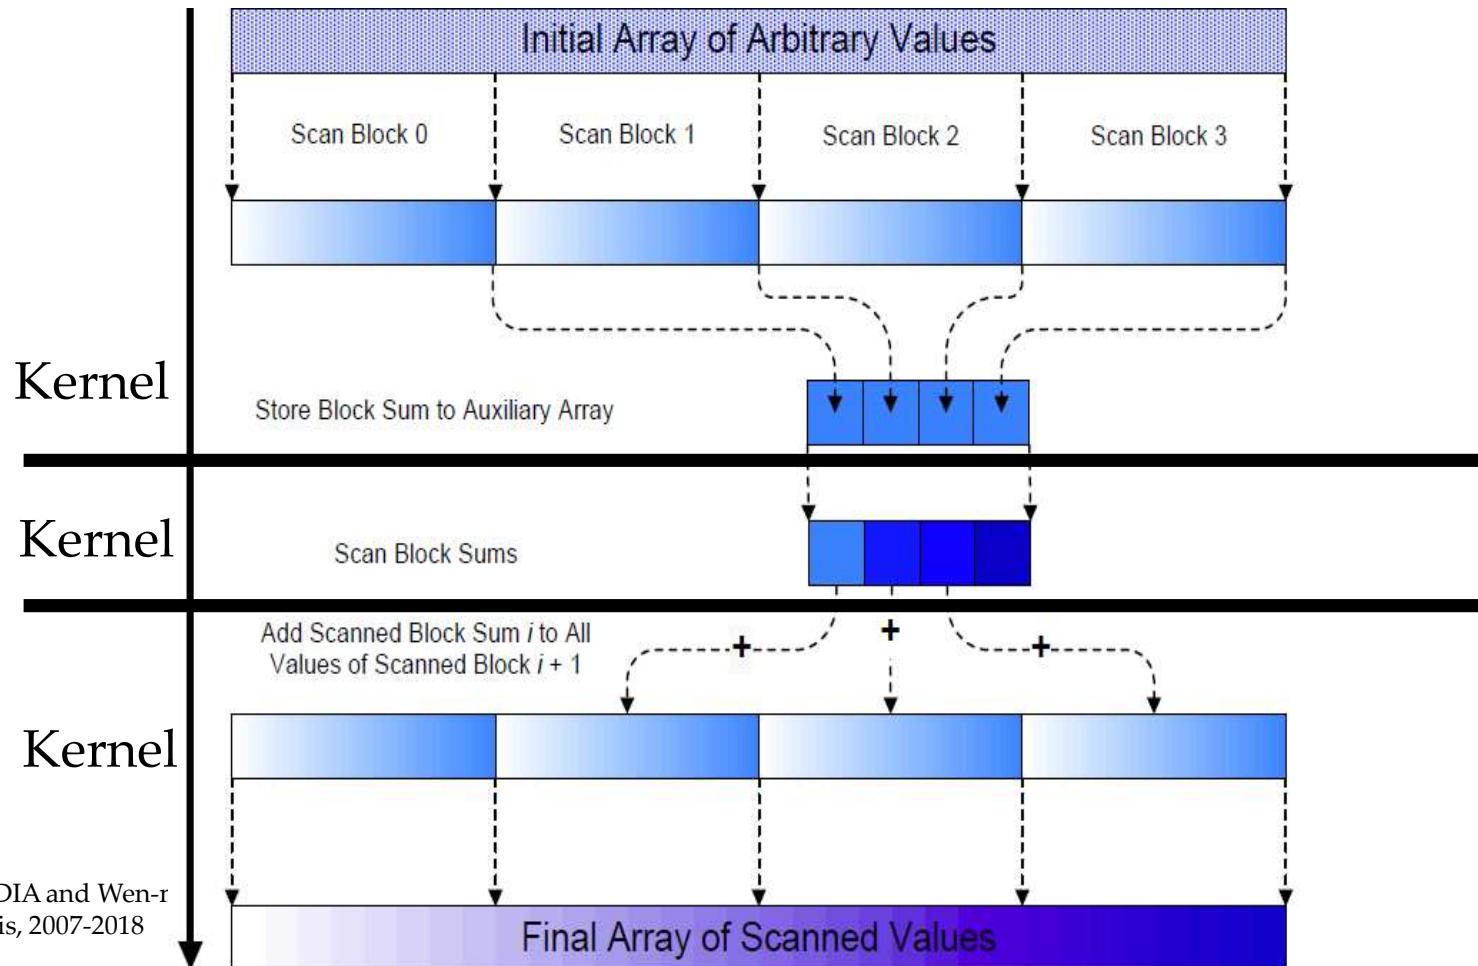
- Data in registers and shared memory of one thread block are not visible to other blocks
- To make data visible, the data has to be written into global memory
- However, any data written to the global memory are not visible until a memory fence. This is typically done by terminating the kernel execution
- Launch another kernel to continue the execution. The global memory writes done by the terminated kernels are visible to all thread blocks.

Scan of Arbitrary Length Input

- Build on the scan kernel that handles up to $2 * \text{blockDim.x}$ elements from Brent-Kung
 - For Kogge-Stone, have each section of blockDim.x elements assigned to a block
- Have each block write the sum of its section into a Sum array using its blockIdx.x as index
- Run parallel scan on the Sum array
 - May need to break down Sum into multiple sections if it is too big for a block
- Add the scanned Sum array values to the elements of corresponding sections

Overall Flow of Complete Scan

A Hierarchical Approach



(Exclusive) Scan Definition

Definition: *The exclusive scan operation takes a binary associative operator \oplus , and an array of n elements*

$$[x_0, x_1, \dots, x_{n-1}]$$

and returns the array

$$[0, x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-2})].$$

Example: If \oplus is addition, then the exclusive scan operation on

[3 1 7 0 4 1 6 3],
would return [0 3 4 11 11 15 16 22].

Why Exclusive Scan

- To find the beginning address of allocated buffers
- Inclusive and Exclusive scans can be easily derived from each other; it is a matter of convenience

[3 1 7 0 4 1 6 3]

Exclusive [0 3 4 11 11 15 16 22]

Inclusive [3 4 11 11 15 16 22 25]

A simple exclusive scan kernel

- Adapt an inclusive, Kogge-Stone scan kernel
 - Block 0:
 - Thread 0 loads 0 into (shared) $T[0]$
 - Other threads load (global) $X[threadIdx.x-1]$ into $T[threadIdx.x]$
 - All other blocks:
 - All threads load $X[blockIdx.x * blockDim.x + threadIdx.x - 1]$ into $T[threadIdx.x]$
- Similar adaption for Brent-Kung kernel but pay attention that each thread loads two elements
 - Only one zero should be loaded
 - All elements should be shifted by only one position



QUESTIONS?

READ CHAPTER 8!

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483,
University of Illinois, Urbana-Champaign

Problem Solving

- Q: Suppose we have a kernel function that performs partial sum reduction. The block dimension is (64,1,1). During the **second** iteration of the for loop, is there any warp that has a control divergence?

```
__shared__ int partialSum[2*BLOCK_SIZE];
// omitted code to fill partialSum from global memory
for (int stride = blockDim.x; stride >= 1; stride /= 2)  {
    __syncthreads();
    if (threadIdx.x < stride) partialSum[t] += partialSum[t + stride];
}
```

- A: No: stride 32 means warp 0 is fully active, while warp 1 is fully inactive

Problem Solving

- Q: Consider a kernel that performs Brent-Kung scan algorithm and assume that there are 1024 elements in a section, and a warp size is 32. In which iteration will there be at least one warp that has a control divergence? The stride is 1 for the first iteration.
- A: In the 6th iteration: strides are 512, 256, 128, 64, 32, 16, ... (and active threads have $\text{threadIdx.x} < \text{stride}$)

Problem Solving

- Q: Suppose we use Brent-Kung in a hierarchical approach to perform parallel scan on a 1D input array of 2^{42} elements. We use 1024 threads per block in all our Brent-Kung kernels and our GPU supports at most 2048 blocks per grid. Each block processes 2048 elements. What is the best approximation of the number of floating-point add operations performed per thread block in the reduction and post scan steps in the Brent-Kung kernel (summed together)?
- A: 2048×2