



ECE408 / CS483 / CSE408
Summer 2025

Applied Parallel Programming

Lecture 15: Parallel Sparse Methods
(Part 2)

What Will You Learn Today?


to regularize irregular data by

- limiting variations with clamping,
- sorting, and
- transposition

to write

- a high-performance SpMV kernel
- based on JDS transposed format

Review: Sparse Matrix-Vector Multiplication (SpMV)



Review: Compressed Sparse Row (CSR) Format


CSR Representation

		Row 0	Row 2	Row 3
Nonzero values	data[7]	{ 3, 1, 2, 4, 1, 1, 1 }		
Column indices	col_index[7]	{ 0, 2, 1, 2, 3, 0, 3 }		
Row Pointers	row_ptr[5]	{ 0, 2, 2, 5, 7 }		

Dense representation

Row 0	3	0	1	0	Thread 0
Row 1	0	0	0	0	Thread 1
Row 2	0	2	4	1	Thread 2
Row 3	1	0	0	1	Thread 3

Review: Regularizing SpMV with ELL(PACK) Format



- Pad all rows to the same length
 - Inefficient if a few rows are much longer than others
- Transpose (Column Major) for DRAM efficiency
- Both data and col_index padded/transposed


Review: Coordinate (COO) format

- Explicitly list the column and row indices for every non-zero element

		Row 0	Row 2	Row 3
Nonzero values	data[7]	{ 3, 1,	2, 4, 1,	1, 1 }
Column indices	col_index[7]	{ 0, 2,	1, 2, 3,	0, 3 }
Row indices	row_index[7]	{ 0, 0,	2, 2, 2,	3, 3 }


Review: Hybrid Format (ELL + COO)

- ELL handles *typical* entries
- COO handles *exceptional* entries
 - Implemented with segmented reduction




Often implemented in sequential host code in practice

Review: Reduced Padding with Hybrid Format




CSR Run-time




Block performance is determined by longest row

JDS (Jagged Diagonal Sparse) Kernel Design for Load Balancing



Sort rows into descending order according to number of non-zero.
Keep track of the original row numbers so that the output vector can be generated correctly.

Sorting Rows According to Length (Regularization)




CSR to JDS Conversion

		Row 0	Row 2	Row 3
Nonzero values	data[7]	{ 3, 1,	2, 4, 1,	1, 1 }
Column indices	col_index[7]	{ 0, 2,	1, 2, 3,	0, 3 }
Row Pointers	row_ptr[5]	{ 0, 2, 2,	5,	7 }
		Row 2	Row 0	Row 3
Nonzero values	data[7]	{ 2, 4, 1,	3, 1,	1 1 }
Column indices	col_index[7]	{ 1, 2, 3,	0, 2,	0, 3 }
JDS Row Pointers	jds_row_ptr[5]	{ 0, 3,	5,	7, 7 }
JDS Row Indices	jds_row_perm[4]	{ 2,	0,	3, 1 }

13

JDS Summary

Nonzero values `data[7]` { 2, 4, 1, 3, 1, 1, 1 }
Column indices `Jds_col_index[7]` { 1, 2, 3, 0, 2, 0, 3 }
JDS row indices `Jds_row_perm[4]` { 2, 0, 3, 1 }
JDS Row Ptrs `Jds_row_ptr[5]` { 0, 3, 5, 7, 7 }



A Parallel SpMV/JDS Kernel

```
1. __global__ void SpMV_JDS(int num_rows, float *data,
   int *col_index, int *jds_row_ptr,int *jds_row_perm,
   float *x, float *y) {
2.     int row = blockIdx.x * blockDim.x + threadIdx.x;
3.     if (row < num_rows) {
4.         float dot = 0;
5.         int row_start = jds_row_ptr[row];
6.         int row_end = jds_row_ptr[row+1];
7.         for (int elem = row_start; elem < row_end; elem++) {
8.             dot += data[elem] * x[col_index[elem]];
9.         }
10.        y[jds_row_perm[row]] = dot;
11.    }
12. }
```

Nonzero values data[7]

Row 2	Row 0	Row 3
{ 2, 4, 1, }	{ 3, 1, }	{ 1, 1 }

Column indices col_index[7]

{ 1, 2, 3, 0, 2, 0, 3 }

JDS Row Pointers jds_row_ptr[5]

{ 0, 3, 5, 7, 7 }

© JDS Row Indices jds_row_perm[4]

{ 2, 0, 3, 1 }

JDS vs. CSR - Control Divergence

- Threads still execute different number of iterations in the JDS kernel for-loop
 - However, neighboring threads tend to execute similar number of iterations because of sorting.
 - Better thread utilization, less control divergence

Nonzero values `data[7]` { 2, 4, 1, 3, 1, 1, 1 }

Column indices `col_index[7]` { 1, 2, 3, 0, 2, 0, 3 }

JDS row indices `Jds_row_perm[4]` { 2, 0, 3, 1 }

JDS Row Ptrs `Jds_row_ptr[5]` { 0, 3, 5, 7, 7 }

`data`


2	4	1	3	1	1	1
---	---	---	---	---	---	---

`col_index`

1	2	3	0	2	0	3
---	---	---	---	---	---	---


JDS vs. CSR Memory Divergence

- Adjacent threads still access non-adjacent memory locations




17


JDS with Transposition




Transposition for Memory Coalescing




JDS Format with Transposed Layout




JDS with Transposition Memory Coalescing



JDS with Transposition Memory Coalescing



JDS with Transposition Memory Coalescing



A Parallel SpMV/JDS_T Kernel


```
1. __global__ void SpMV_JDS_T(int num_rows, float *data,
   int *col_index, int *jds_t_col_ptr, int *jds_row_perm,
   float *x, float *y) {
2.     int row = blockIdx.x * blockDim.x + threadIdx.x;
3.     if (row < num_rows) {
4.         float dot = 0;
5.         unsigned int sec = 0;
6.         while (jds_t_col_ptr[sec+1]-jds_t_col_ptr[sec] > row) {
7.             dot += data[jds_t_col_ptr[sec]+row] *
8.                 x[col_index[jds_t_col_ptr[sec]+row]];
9.             sec++;
10.        }
11.        y[jds_row_perm[row]] = dot;
12.    }
13. }
```

Column indices col_index[7]	{ 1, 0, 3, 2, 2, 3 3 }
JDS_T Column Pointers jds_t_col_ptr[5]	{ 0, 3, 6, 7, 7 }
JDS Row Indices jds_row_perm[4]	{ 2, 0, 3, 1 }


Lab 8 Variable Names

JDS_T Length of Cols matRows[4] {3, 2, 2, 0 }

		Sec 0	Sec 1	Sec 2
Nonzero values	matData[7]	{ 2, 3, 1,	4, 1, 1	1 }
Column indices	matCols[7]	{ 1, 0, 0,	2, 2, 3	3 }
JDS_T Column Pointers	matColStart[4]	{0, 3, 6, 7 }		
JDS Row Indices	matRowPerm[4]	{2, 0, 3, 1 }		




Roughly Random...




Roughly Random...

Probably best with ELL.

- Padding will be uniformly distributed
- Sparse representation will be uniform




High variance in rows...




High variance in rows

Probably best with ELL/COO

- Benefit of ELL for most cases
- Outliers are captured with COO




Very sparse...




Very sparse

Probably best with COO

- Not a lot of data, compute is sparse



Roughly triangular...





Roughly triangular...

Probably best with JDS

- Takes advantage of sparsity structure

Banded Matrix...







Banded Matrix...

Probably best with ELL

- Small amount of variance in rows



Other formats


- Diagonal (DIA): for strictly banded/diagonal matrices
- Packet (PKT): create diagonal submatrices by reordering rows/cols
- Dictionary of Keys (DOK): map of (row/col) to data
- Compressed Sparse Column (CSC): when to use over CSR?
- Blocked CSR: useful for block sparse matrices
- Hybrids of these...

Sparse Matrices as Foundation for Advanced Algorithm Techniques

- Graphs are often represented as sparse adjacency matrices
 - Used extensively in social network analytics, natural language processing, etc.
 - Sparse Matrix-Matrix multiplication (SpMM) is a fundamental operator in GNNs, which performs a multiplication between a sparse matrix and a dense matrix.
- Binning techniques often use sparse matrices for data compaction
 - Used extensively in ray tracing, particle-based fluid dynamics methods, and games
- These will be covered in ECE508/CS508

Sparse Matrices as Foundation for Advanced Algorithm Techniques

- Graphs are often represented as sparse adjacency matrices
 - Used extensively in social network analytics, natural language processing, etc.
 - Sparse Matrix-Matrix multiplication (SpMM) is a fundamental operator in GNNs, which performs a multiplication between a sparse matrix and a dense matrix.
- Binning techniques often use sparse matrices for data compaction
 - Used extensively in ray tracing, particle-based fluid dynamics methods, and games
- These will be covered in ECE508/CS508



QUESTIONS?

READ CHAPTER 10!

Problem Solving

- Q: Consider the following sparse Matrix:
- For each of the following **data** layouts in memory, select the option that best matches all the sparse matrix formats that can store the data in memory as depicted.

- A:
 - 1) CSR, COO
 - 2) ???
 - 3) JDS, COO
 - 4) COO
 - 5) JDS-Transposed, COO

$$\begin{bmatrix} 1 & 0 & 4 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 \\ 0 & 7 & 0 & 9 & 3 \\ 0 & 0 & 0 & 0 & 0 \\ 6 & 5 & 0 & 0 & 8 \end{bmatrix}$$

Layout 1:

1	4	2	7	9	3	6	5	8
---	---	---	---	---	---	---	---	---

Layout 2:

1	2	7	6	4	0	9	5	0	0	3	8
---	---	---	---	---	---	---	---	---	---	---	---

Layout 3:

7	9	3	6	5	8	1	4	2
---	---	---	---	---	---	---	---	---

Layout 4:

9	7	1	2	4	3	5	8	6
---	---	---	---	---	---	---	---	---

Layout 5:

7	6	1	2	9	5	4	3	8
---	---	---	---	---	---	---	---	---