



ECE408 / CS483 / CSE408
Summer 2025

Applied Parallel Programming

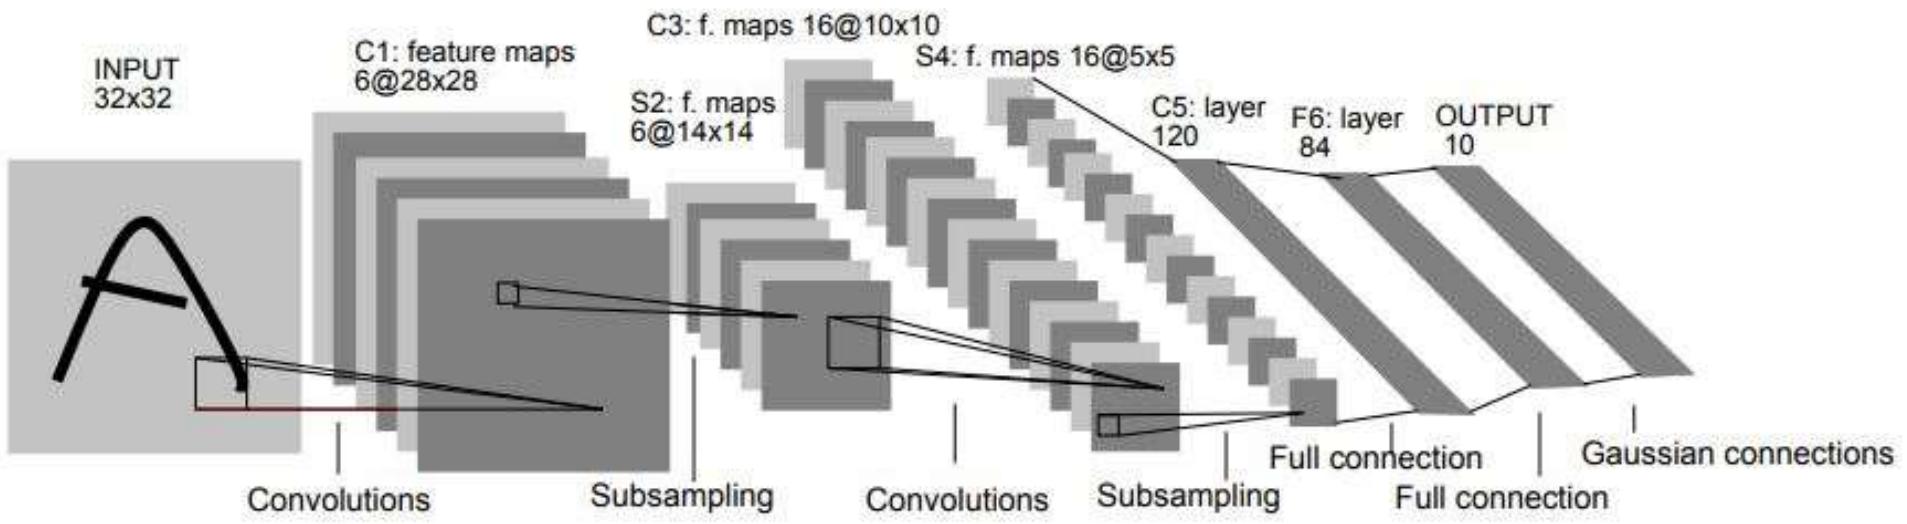
Lecture 18: Computation in Deep
Neural Networks



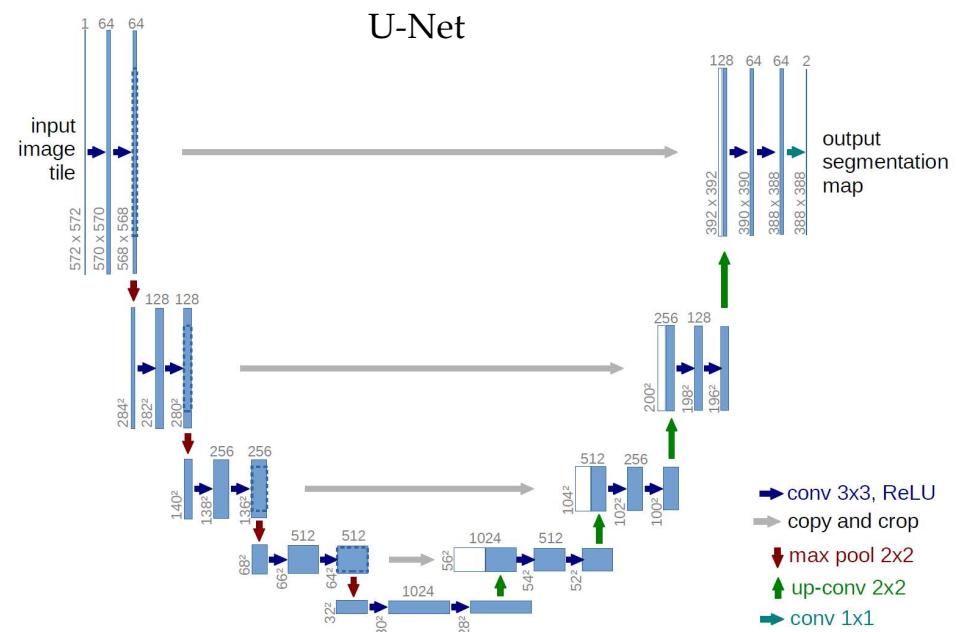
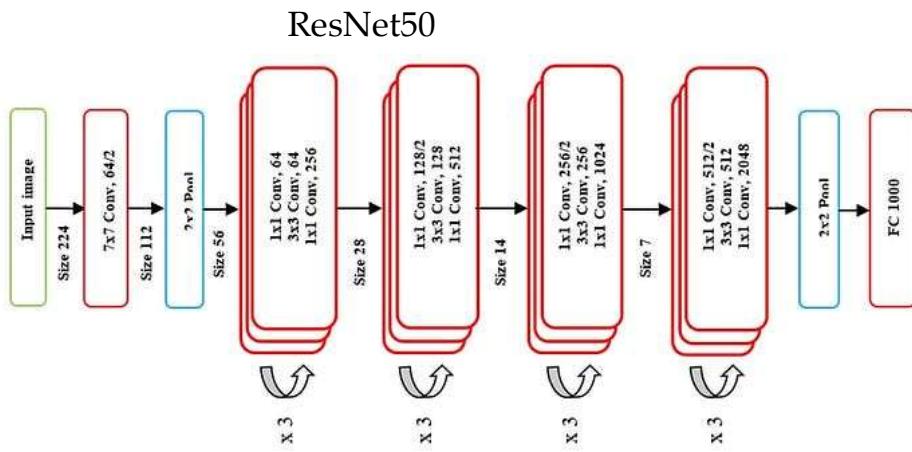
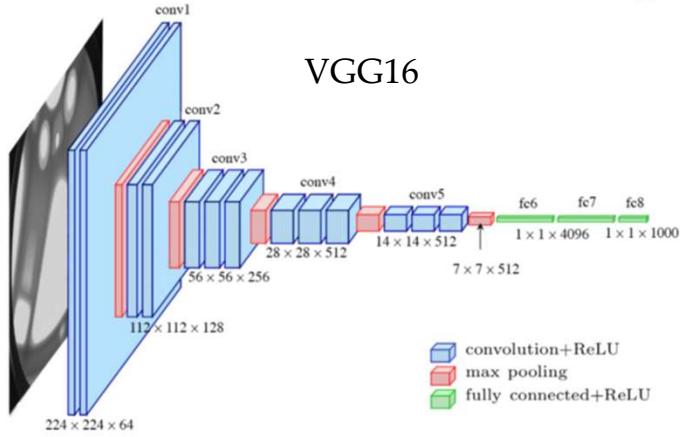
What Will You Learn Today?

to implement the different types of layers
in a Convolutional Neural Network (CNN)

LeNet-5:CNN for hand-written digit recognition



Many Types of CNNs



Anatomy of a Convolution Layer

Input features

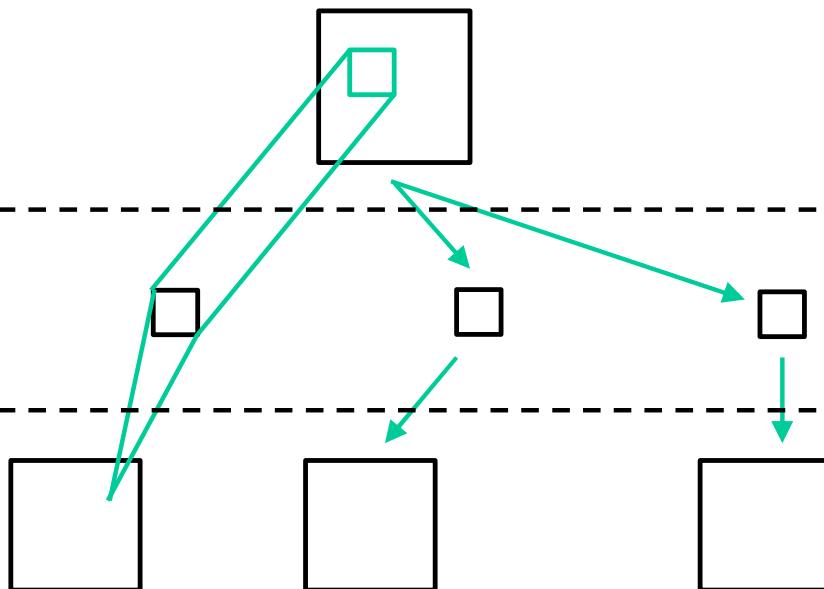
- A inputs each $N_1 \times N_2$

Convolution Layer

- B convolution kernels each $K_1 \times K_2$

Output Features (total of B)

- $A \times B$ outputs each $(N_1 - K_1 + 1) \times (N_2 - K_2 + 1)$



Notion of a Channel in Input Layer

Some Set of Input Features are Related

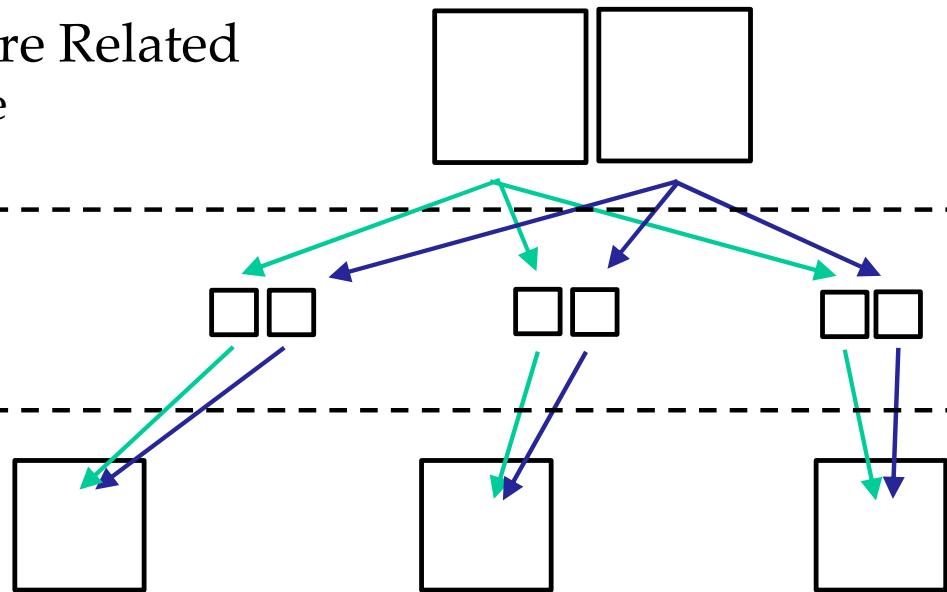
- For example: Red, Green, Blue

Convolution Layer

- Different kernels per channel

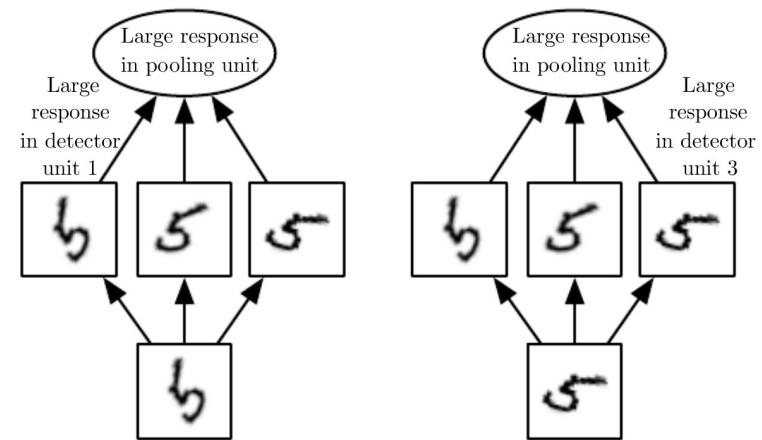
Output Features

- Channels combine per output feature

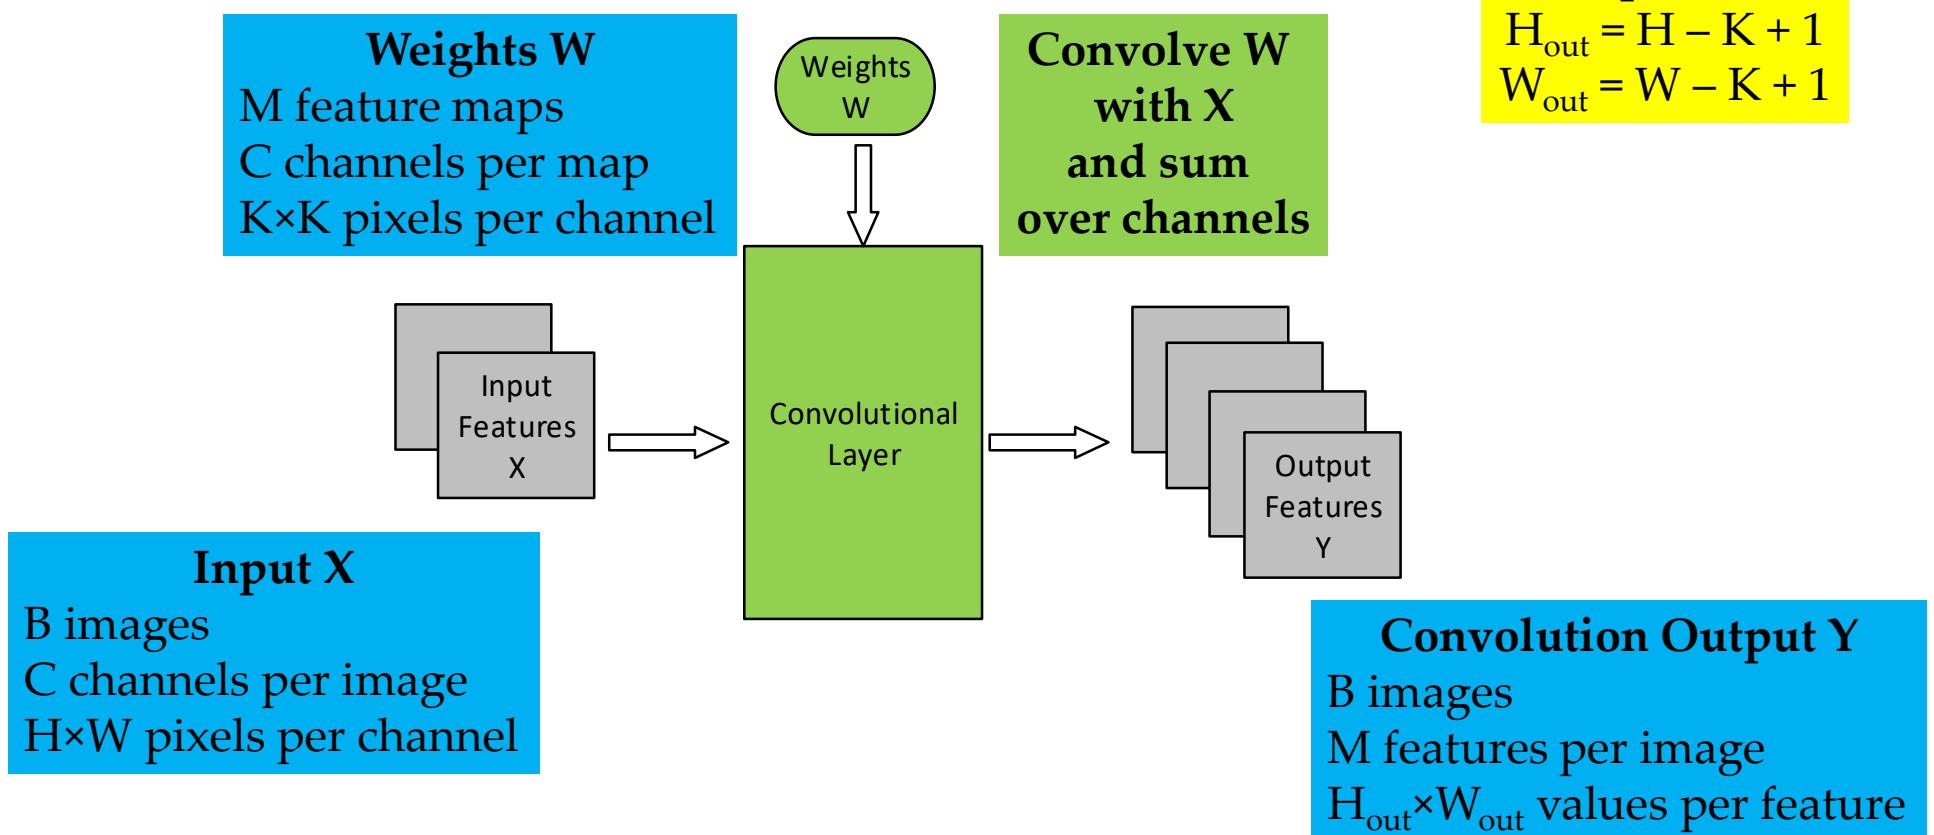


2-D Pooling (Subsampling)

- A subsampling layer
 - Sometimes with bias and non-linearity built in
- Common types
 - max, average, L^2 norm, weighted average
- Helps make representation invariant to size scaling and small translations in the input



Forward Propagation



Outputs Must Use Full Mask/Kernel

| X | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
|---|---|---|---|---|---|---|---|
| 2 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
| 3 | 2 | 3 | 4 | 5 | 6 | 7 | 8 |
| 4 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
| 5 | 4 | 5 | 6 | 7 | 8 | 5 | 6 |
| 6 | 5 | 6 | 7 | 8 | 5 | 6 | 7 |
| 7 | 6 | 7 | 8 | 9 | 0 | 1 | 2 |
| 8 | 7 | 8 | 9 | 0 | 1 | 2 | 3 |

| W | 1 | 2 | 3 | 2 | 1 |
|---|---|---|---|---|---|
| 2 | 1 | 2 | 3 | 2 | 1 |
| 3 | 2 | 3 | 4 | 3 | 2 |
| 4 | 3 | 4 | 5 | 4 | 3 |
| 5 | 2 | 3 | 4 | 3 | 2 |
| 6 | 1 | 2 | 3 | 2 | 1 |

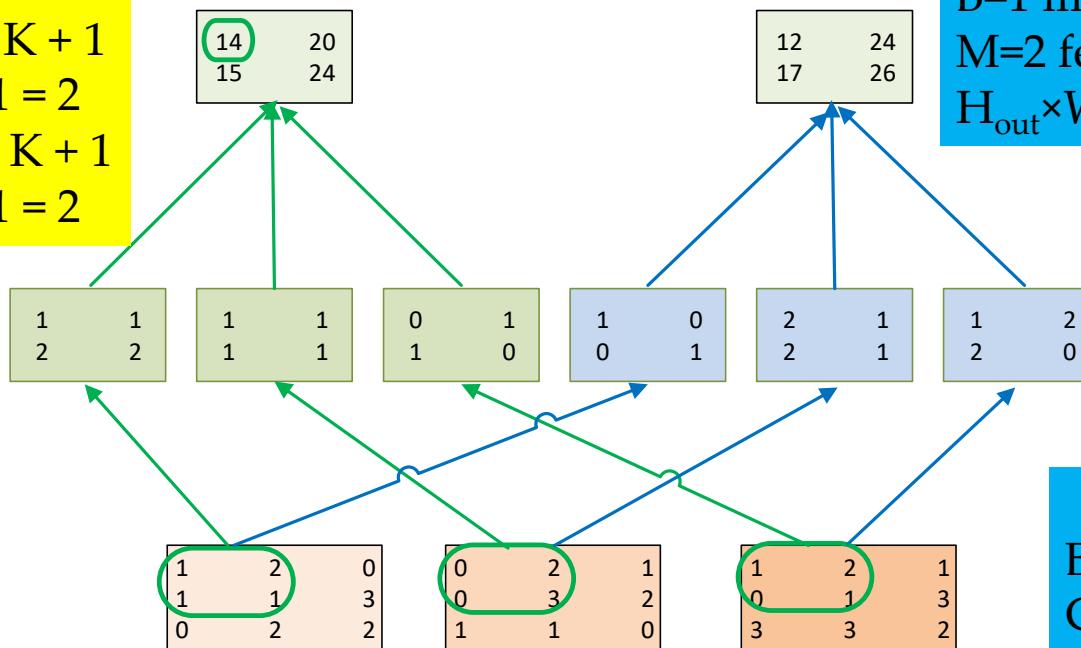
Compute only
this part of Y.

| Y | | | | | | | | | |
|---|--|--|---|---|---|--|--|--|--|
| | | | | | | | | | |
| | | | 3 | 2 | 1 | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |

| | | | | | |
|-----|-----|-----|-----|-----|-----|
| 1 | 4 | 9 | 16 | 25 | 36 |
| 4 | 9 | 16 | 25 | 36 | 49 |
| 9 | 16 | 25 | 36 | 49 | 64 |
| 16 | 25 | 36 | 49 | 64 | 81 |
| 25 | 36 | 49 | 64 | 81 | 100 |
| 36 | 49 | 64 | 81 | 100 | 121 |
| 49 | 64 | 81 | 100 | 121 | 144 |
| 64 | 81 | 100 | 121 | 144 | 169 |
| 81 | 100 | 121 | 144 | 169 | 196 |
| 100 | 121 | 144 | 169 | 196 | 225 |

Example of the Forward Path of a Convolution Layer

Output Size
 $H_{out} = H - K + 1$
 $= 3 - 2 + 1 = 2$
 $W_{out} = W - K + 1$
 $= 3 - 2 + 1 = 2$



Weights W
 $M=2$ feature maps
 $C=3$ channels per map
 $K \times K=2 \times 2$ pixels per channel

Input X
 $B=1$ image
 $C=3$ channels
 $H \times W=3 \times 3$ pixels per channel

Sequential Code: Forward Convolutional Layer

```
void convLayer_forward(int B, int M, int C, int H, int W, int K, float* X, float* W, float* Y) {  
    int H_out = H - K + 1;                                // calculate H_out, W_out  
    int W_out = W - K + 1;  
  
    for (int b = 0; b < B; ++b)                          // for each image  
        for(int m = 0; m < M; m++)                      // for each output feature map  
            for(int h = 0; h < H_out; h++)                // for each output value (two loops)  
                for(int w = 0; w < W_out; w++) {  
                    Y[b, m, h, w] = 0.0f;                  // initialize sum to 0  
                    for(int c = 0; c < C; c++)              // sum over all input channels  
                        for(int p = 0; p < K; p++)          // KxK filter  
                            for(int q = 0; q < K; q++)  
                                Y[b, m, h, w] += X[b, c, h + p, w + q] * W[m, c, p, q];  
                }  
}
```

A Small Convolution Layer Example

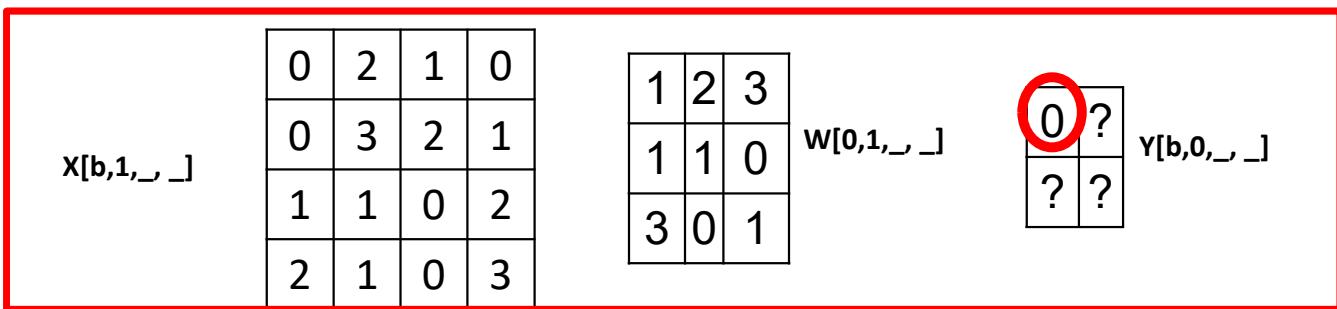
$X[b, 1, \dots]$ input channel
 $w[0,1,\dots]$
 $y[b, 0, \dots]$ output map

Image b in mini batch
 $x[b,0,\dots]$

| | | | |
|---|---|---|---|
| 1 | 2 | 0 | 1 |
| 1 | 1 | 3 | 2 |
| 0 | 2 | 2 | 0 |
| 2 | 1 | 0 | 3 |

| | | |
|---|---|---|
| 1 | 1 | 1 |
| 2 | 2 | 3 |
| 2 | 1 | 0 |

$w[0,0,\dots]$



$x[b,2,\dots]$

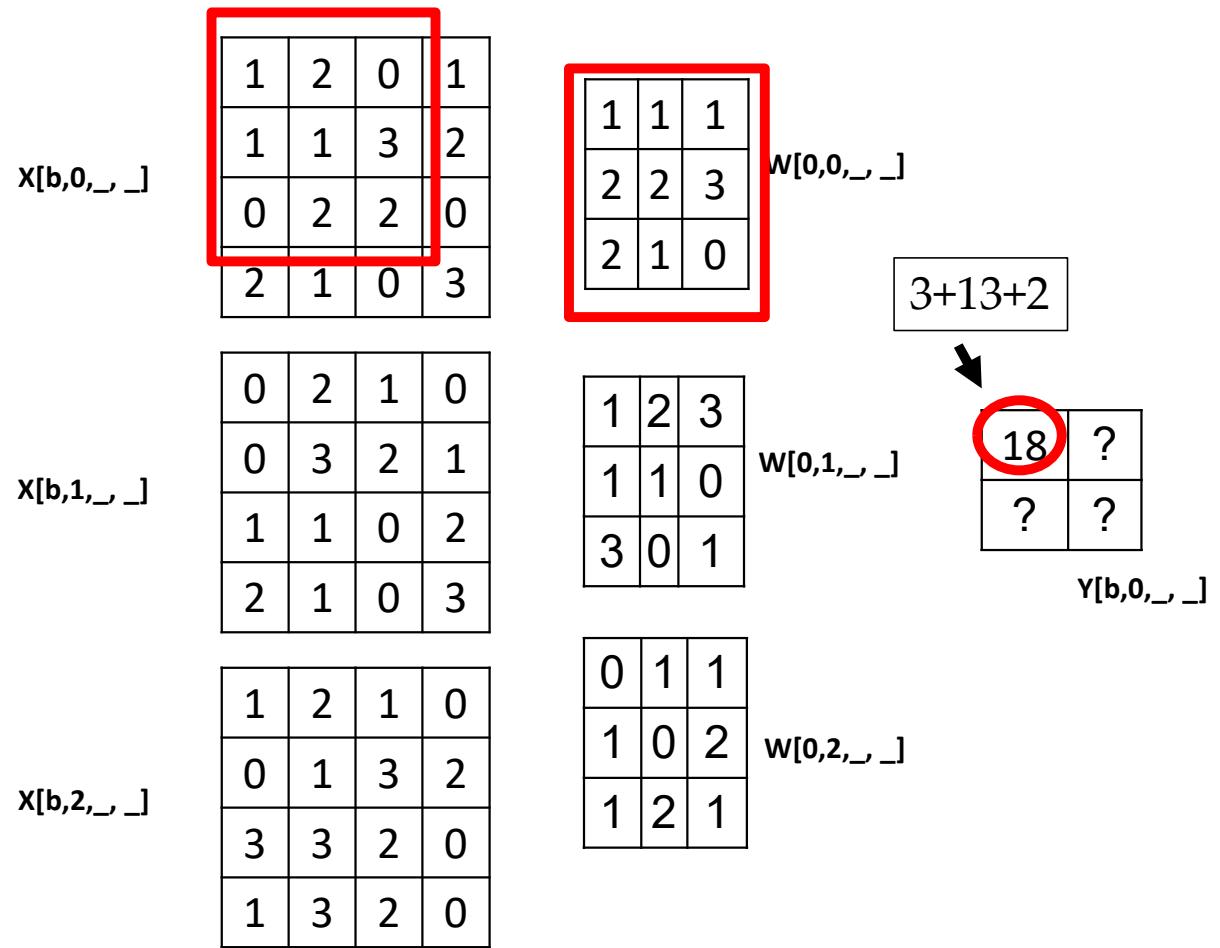
| | | | |
|---|---|---|---|
| 1 | 2 | 1 | 0 |
| 0 | 1 | 3 | 2 |
| 3 | 3 | 2 | 0 |
| 1 | 3 | 2 | 0 |

| | | |
|---|---|---|
| 0 | 1 | 1 |
| 1 | 0 | 2 |
| 1 | 2 | 1 |

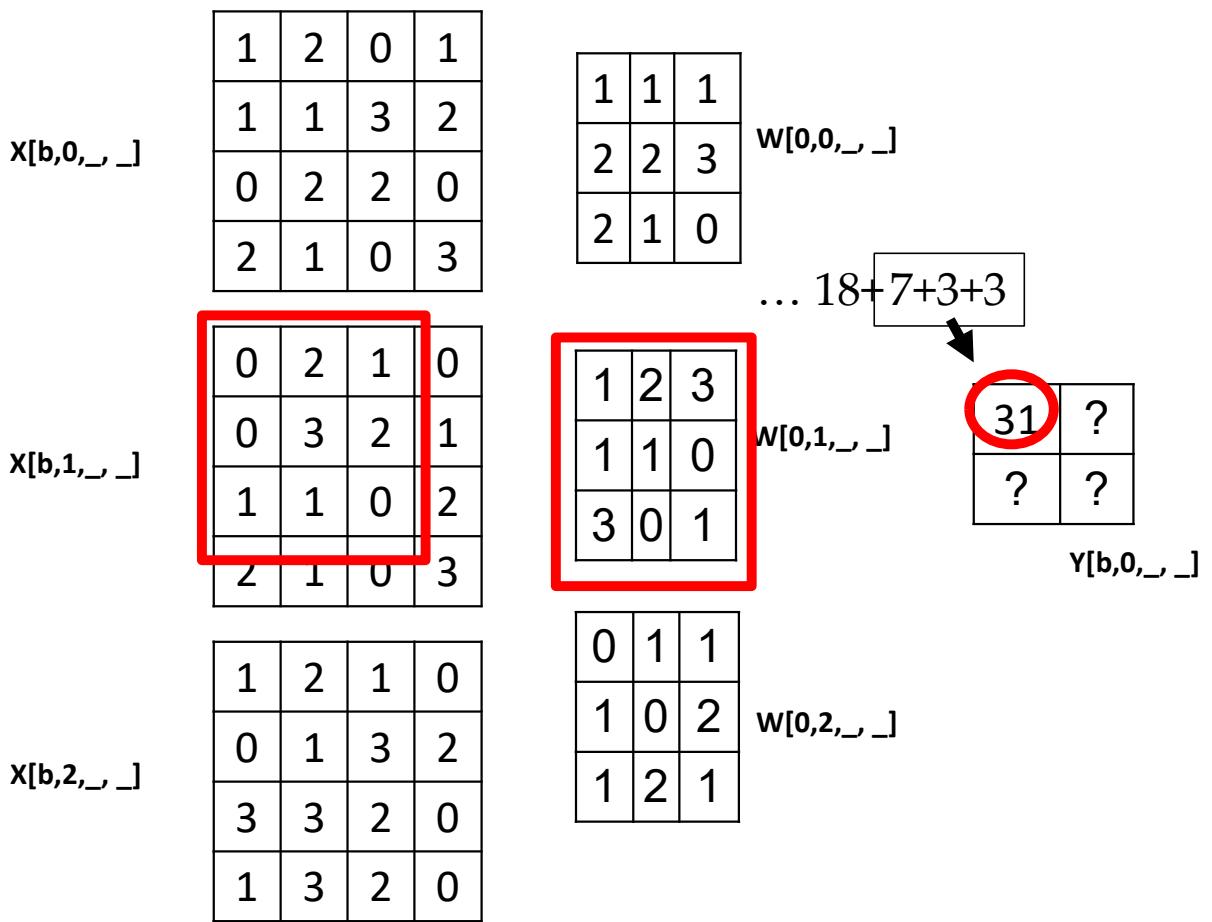
$w[0,2,\dots]$

A Small Convolution Layer Example

$c = 0$

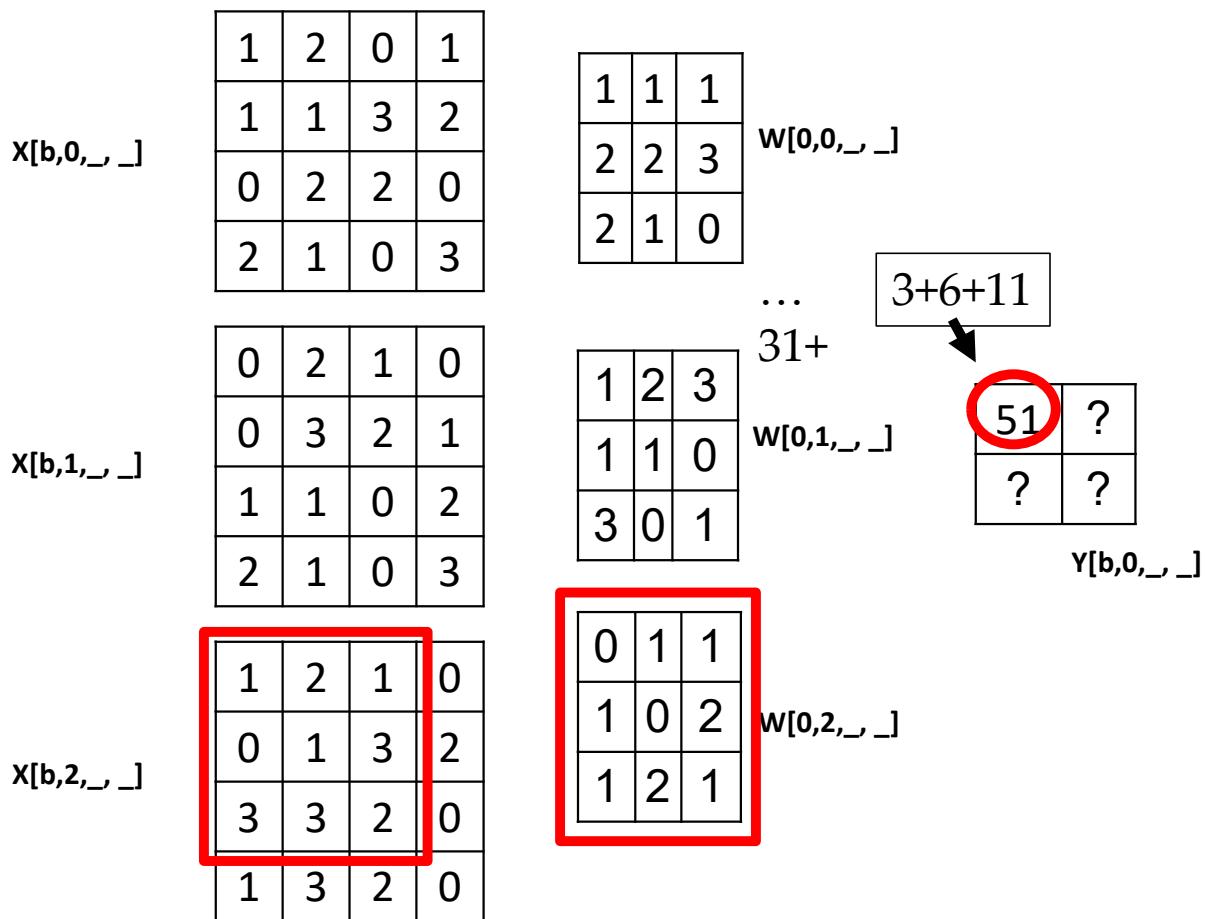


A Small Convolution Layer Example $c = 1$



A Small Convolution Layer Example

$c = 2$



Parallelism in a Convolution Layer

Output feature maps can be calculated in parallel

- Usually a small number, not sufficient to fully utilize a GPU

All **output feature map pixels** can be calculated in parallel

- All rows can be done in parallel
- All pixels in each row can be done in parallel
- Large number but diminishes as we go into deeper layers

All **input feature maps** can be processed in parallel,
but need atomic operation or tree reduction (we'll learn later)

Different layers may demand different strategies.

Subsampling (Pooling) by Scale N

Convolution Output Y
B images
M features per image
 $H_{out} \times W_{out}$ values per feature

Average over $N \times N$ blocks,
then calculate sigmoid

Output Size
 $H_{S(N)} = \text{floor}(H_{out} / N)$
 $W_{S(N)} = \text{floor}(W_{out} / N)$

Subsampling/Pooling Output S
B images
M features per image
 $H_{S(N)} \times W_{S(N)}$ values per feature

Sequential Code: Forward Pooling Layer

```
void poolingLayer_forward(int B, int M, int H_out, int W_out, int N, float* Y, float* S)
{
    for (int b = 0; b < B; ++b)                      // for each image
        for (int m = 0; m < M; ++m)                  // for each output feature map
            for (int x = 0; x < H_out/N; ++x)        // for each output value (two loops)
                for (int y = 0; y < W_out/N; ++y) {
                    float acc = 0.0f                     // initialize sum to 0
                    for (int p = 0; p < N; ++p)          // loop over NxN block of Y (two loops)
                        for (int q = 0; q < N; ++q)
                            acc += Y[b, m, N*x + p, N*y + q];
                    acc /= N * N;                      // calculate average over block
                    S[b, m, x, y] = sigmoid(acc + bias[m]) // bias, non-linearity
                }
}
```

Kernel Implementation of Subsampling Layer

- Straightforward mapping from grid to subsampled output feature map pixels
- in GPU kernel,
 - need to manipulate index mapping
 - for accessing the output feature map pixels
 - of the previous convolution layer.
- Often merged into the previous convolution layer to save memory bandwidth

Design of a Basic Kernel

- Each block computes
 - a tile of output pixels for one feature
 - `TILE_WIDTH` pixels in each dimension
- Grid's X dimension maps to M output feature maps
- Grid's Y dimension maps to the tiles in the output feature maps (linearized order).
- (Grid's Z dimension is used for images in batch, which we omit from slides.)

tiles covering an output feature map, marked with linearized indices

| | | | | |
|----|----|----|----|----|
| 0 | 1 | 2 | 3 | 4 |
| 5 | 6 | 7 | 8 | 9 |
| 10 | 11 | 12 | 13 | 14 |
| 15 | 16 | 17 | 18 | 19 |

A Small Example

Assume

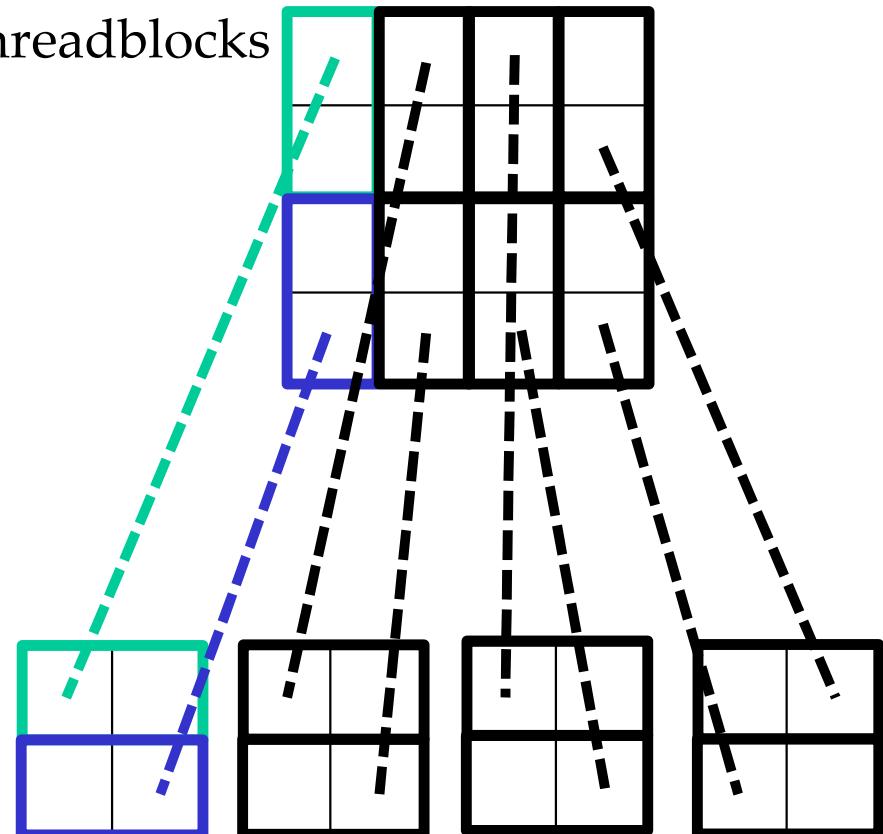
- **M = 4** (4 output feature maps),
- thus 4 blocks in the X dimension, and
- **W_out = H_out = 8** (8x8 output features).

If **TILE_WIDTH = 4**,

we also need 4 blocks in the Y dimension:

- for each output feature,
- top two blocks in each column calculates the top row of tiles, and
- bottom two calculate the bottom row.

CUDA Grid and Threadblocks



Output Feature Maps and Tiles 22

Host Code for a Basic Kernel: CUDA Grid

Consider an output feature map:

- width is **W_out**, and
- height is **H_out**.
- Assume these are multiples of **TILE_WIDTH**.

| | | | | |
|----|----|----|----|----|
| 0 | 1 | 2 | 3 | 4 |
| 5 | 6 | 7 | 8 | 9 |
| 10 | 11 | 12 | 13 | 14 |
| 15 | 16 | 17 | 18 | 19 |

Let **X_grid** be the number of blocks needed in X dim (5 above).

Let **Y_grid** be the number of blocks needed in Y dim (4 above).

Host Code for a Basic Kernel: CUDA Grid

(Assuming W_out and H_out are multiples of TILE_WIDTH.)

```
#define TILE_WIDTH 16          // We will use 4 for small examples.  
W_grid = W_out/TILE_WIDTH;   // number of horizontal tiles per output map  
H_grid = H_out/TILE_WIDTH;   // number of vertical tiles per output map  
Y = H_grid * W_grid;  
  
dim3 blockDim(TILE_WIDTH, TILE_WIDTH, 1); // output tile for untiled code  
dim3 gridDim(M, Y, 1);  
  
ConvLayerForward_Kernel<<< gridDim, blockDim >>>(...);
```

Sequential Code: Forward Convolutional Layer

```
void convLayer_forward(int B, int M, int C, int H, int W, int K, float* X, float* W, float* Y) {  
    int H_out = H - K + 1;                      // calculate H_out, W_out  
    int W_out = W - K + 1;  
  
    for (int b = 0; b < B; ++b)                  // for each image  
        for(int m = 0; m < M; m++)              // for each output feature map  
            for(int h = 0; h < H_out; h++)        // for each output value (two loops)  
                for(int w = 0; w < W_out; w++) {  
                    Y[b, m, h, w] = 0.0f;           // initialize sum to 0  
                    for(int c = 0; c < C; c++)      // sum over all input channels  
                        for(int p = 0; p < K; p++)    // KxK filter  
                            for(int q = 0; q < K; q++)  
                                Y[b, m, h, w] += X[b, c, h + p, w + q] * W[m, c, p, q];  
                }  
}
```

Computed by the grid

Computed by a thread

Partial Kernel Code for a Convolution Layer

```
__global__ void ConvLayerForward_Basic_Kernel
    (int C, int W_grid, int K, float* X, float* W, float* Y)
{
    int m = blockIdx.x;
    int h = (blockIdx.y / W_grid) * TILE_WIDTH + threadIdx.y;
    int w = (blockIdx.y % W_grid) * TILE_WIDTH + threadIdx.x;
    float acc = 0.0f;
    for (int c = 0; c < C; c++) {           // sum over all input channels
        for (int p = 0; p < K; p++) {       // loop over KxK filter
            for (int q = 0; q < K; q++)
                acc += X[c, h + p, w + q] * W[m, c, p, q];
    }
    Y[m, h, w] = acc;
}
```

Some Observations

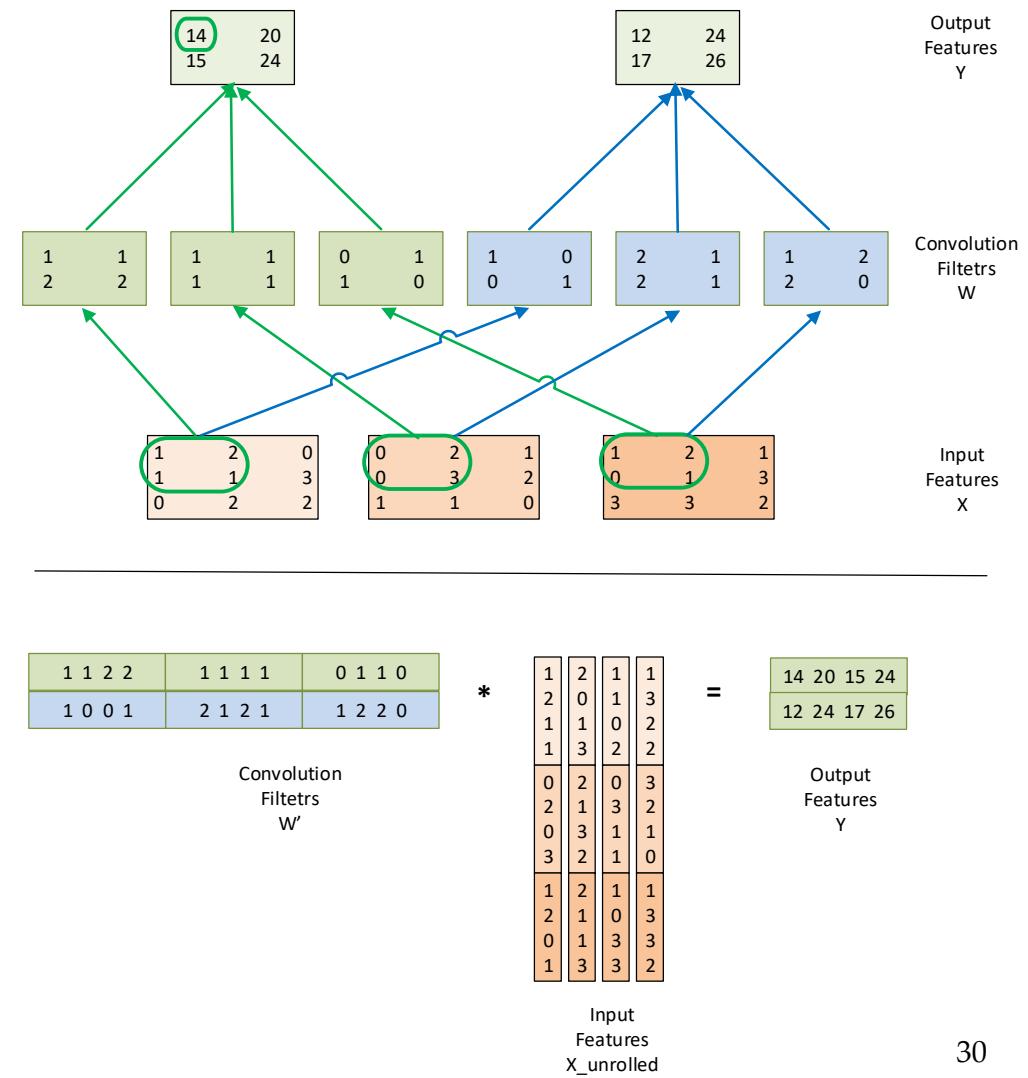
Enough parallelism

- if the total number of pixels
- across all output feature maps is large
- (often the case for CNN layers)

Each input tile

- loaded M times (number of output features), so
- **not efficient in global memory bandwidth,**
- but block scheduling in X dimension should give cache benefits.

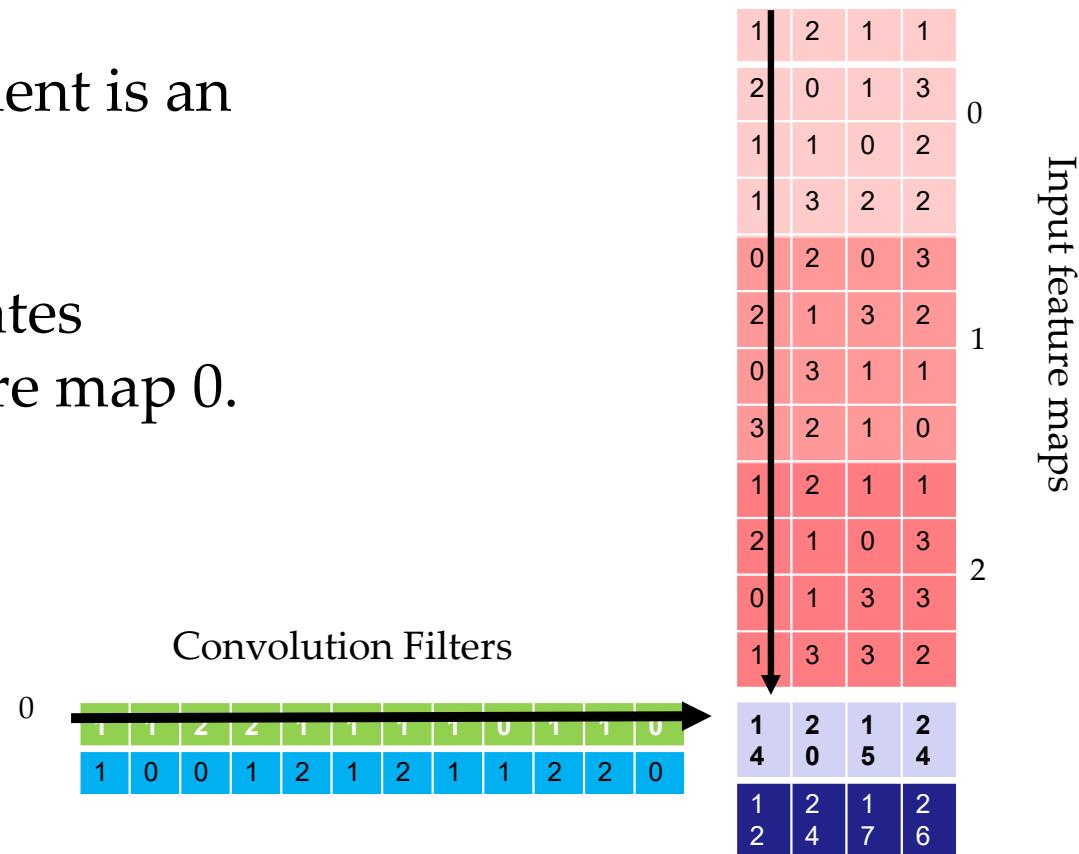
Implementing a Convolution Layer with Matrix Multiplication



Simple Matrix Multiplication

Each product matrix element is an output feature map pixel.

This inner product generates element 0 of output feature map 0.

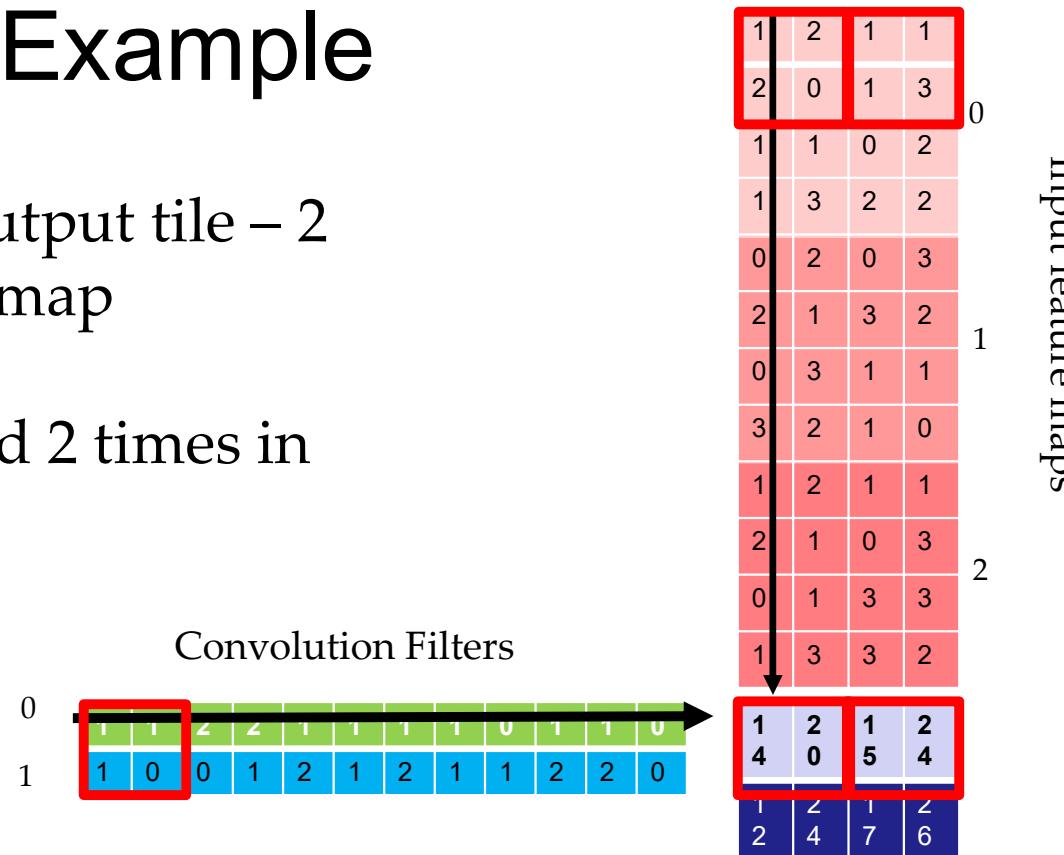


Tiled Matrix Multiplication

2x2 Example

Each block calculates one output tile – 2 elements from each output map

Each input element is reused 2 times in the shared memory

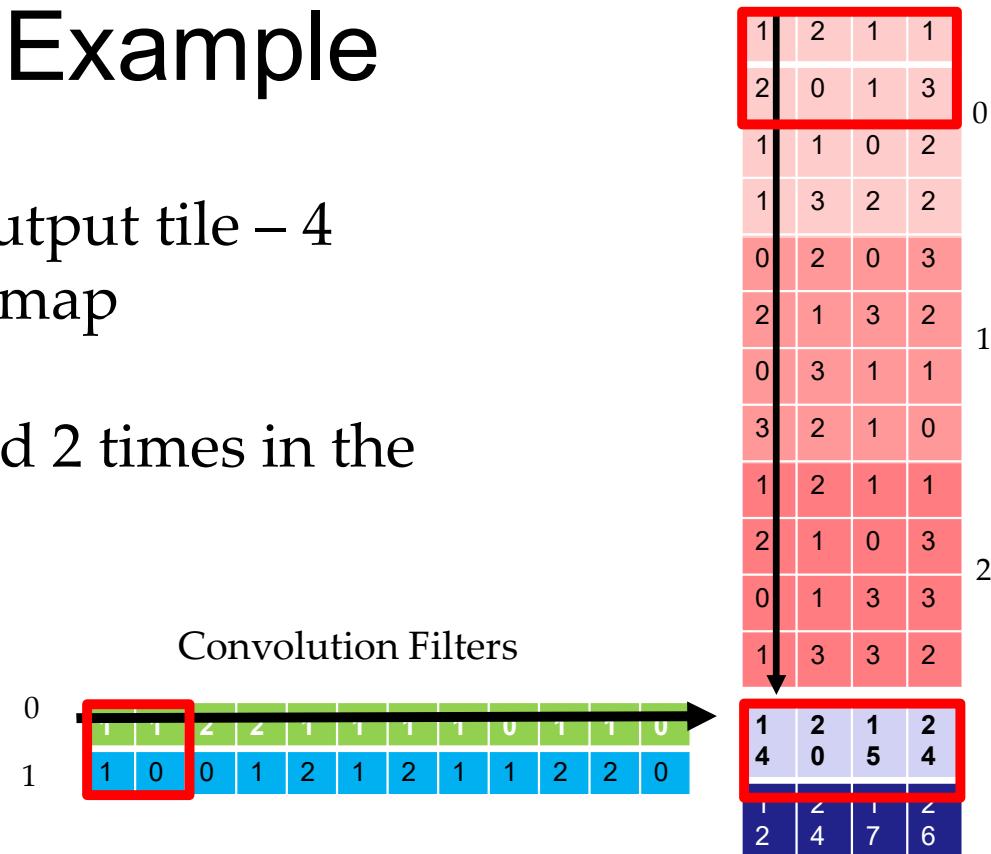


Tiled Matrix Multiplication

2x4 Example

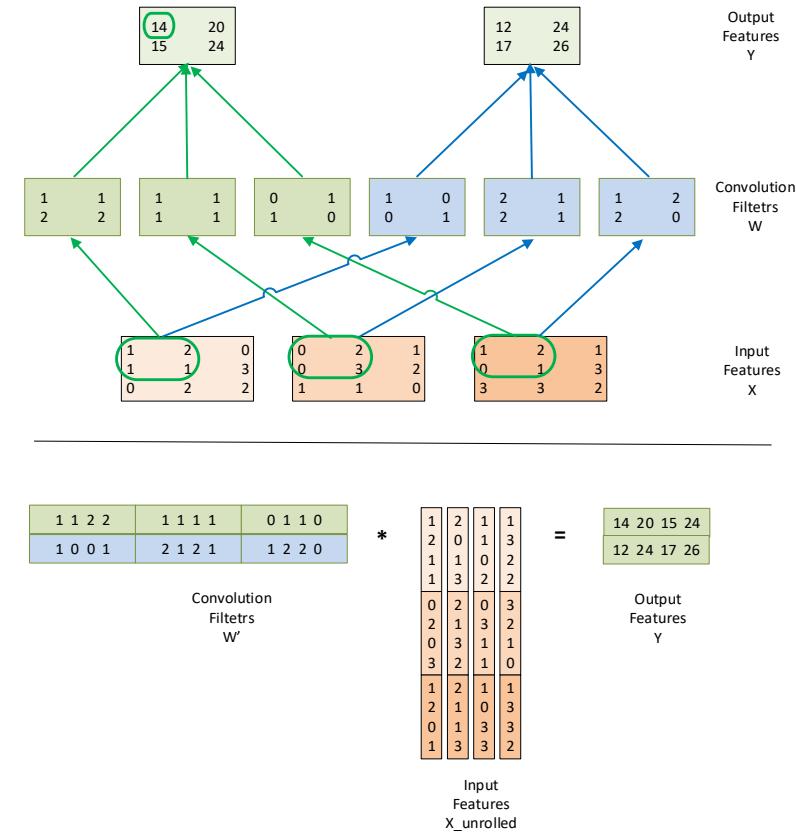
Each block calculates one output tile – 4 elements from each output map

Each input element is reused 2 times in the shared memory



Efficiency Analysis: Total Input Replication

- Replicated input features are shared among output maps
 - There are $H_{out} * W_{out}$ output feature map elements
 - Each requires K^2 elements from the input feature maps
 - So, the total number of input element after replication is $H_{out} * W_{out} * K^2$ times for each input feature map
 - The total number of elements in each original input feature map is $(H_{out}+K-1) * (W_{out}+K-1)$



Analysis of a Small Example

$$H_{out} = 2$$

$$W_{out} = 2$$

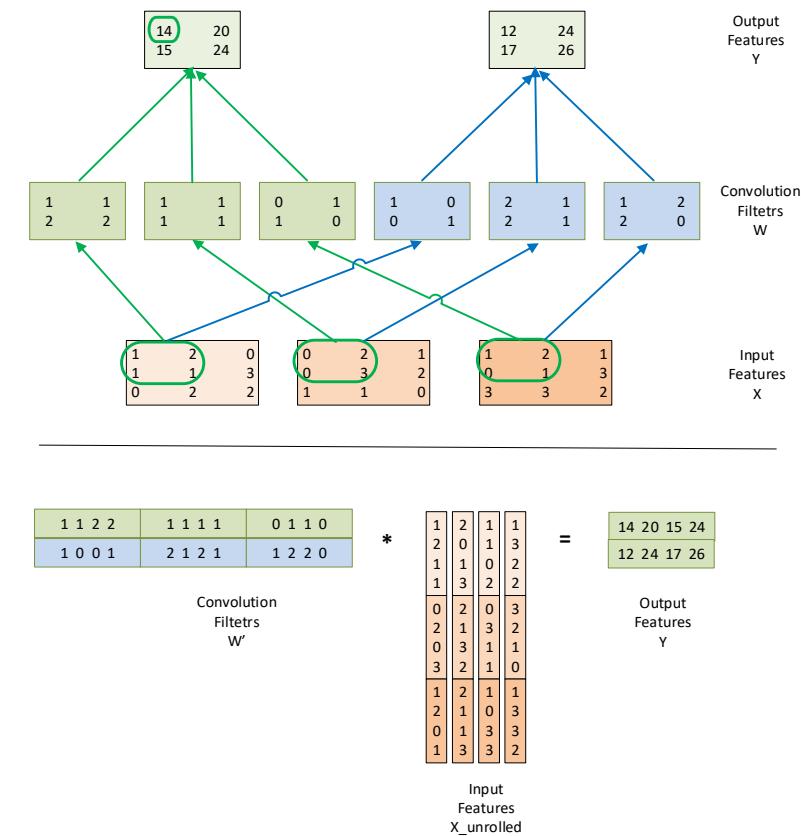
$$K = 2$$

There are 3 input maps (channels)

The total number of input elements in the replicated (“unrolled”) input matrix is $3*2*2*2*2$

The replicating factor is

$$(3*2*2*2*2)/(3*3*3) = 1.78$$

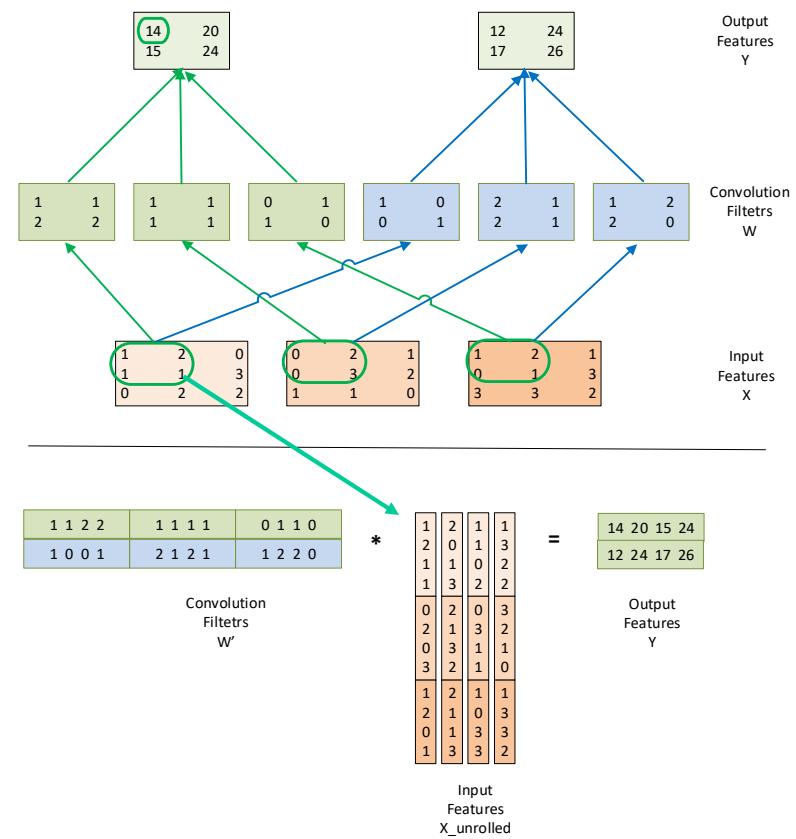


Memory Access Efficiency of Original Convolution Algorithm

- Assume that we use tiled 2D convolution
- For input elements
 - Each output tile has TILE_WIDTH^2 elements
 - Each input tile has $(\text{TILE_WIDTH}+K-1)^2$
 - The total number of input feature map element accesses was $\text{TILE_WIDTH}^2 \times K^2$
 - The reduction factor of the tiled algorithm is $K^2 \times \text{TILE_WIDTH}^2 / ((\text{TILE_WIDTH}+K-1)^2)$
- The convolution filter weight elements are reused within each output tile

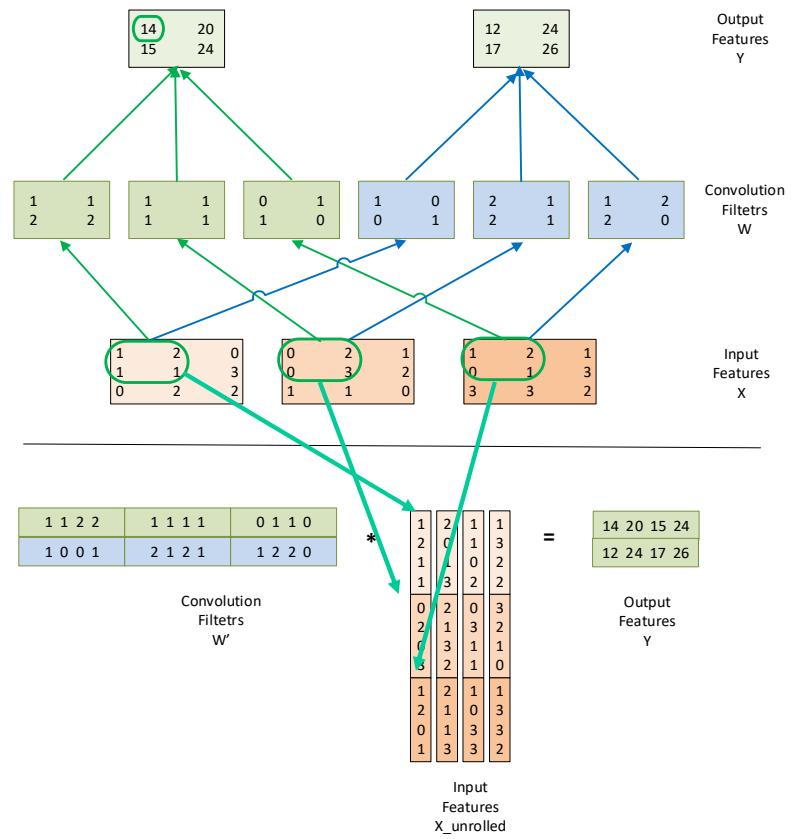
Properties of the Unrolled Matrix

- Each unrolled column corresponds to an output feature map element
- For an output feature element (h, w) , the index for the unrolled column is $h * W_{out} + w$ (linearized index of the output feature map element)



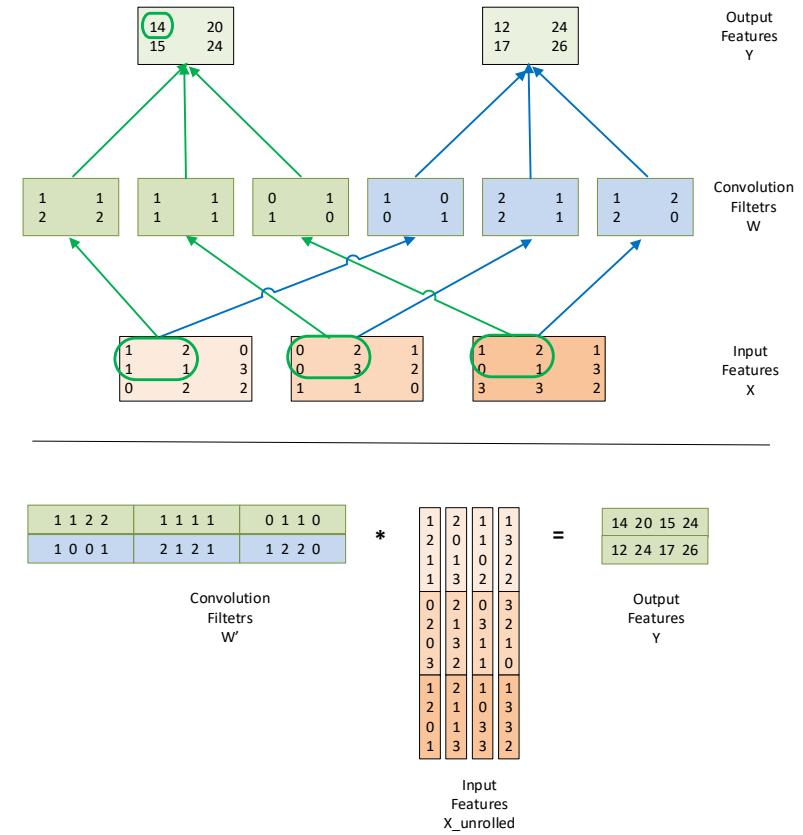
Properties of the Unrolled Matrix (cont.)

- Each section of the unrolled column corresponds to an input feature map
- Each section of the unrolled column has k^k elements (convolution mask size)
- For an input feature map c , the vertical index of its section in the unrolled column is $c \cdot k^k$ (linearized index of the output feature map element)



To Find the Input Elements

- For output element (h, w) , the base index for the upper left corner of the input feature map c is (c, h, w)
- The input element index for multiplication with the convolution mask element (p, q) is $(c, h+p, w+q)$



Input to Unrolled Matrix Mapping

Output element (h, w)

Mask element (p, q)

Input feature map c

```
// calculate the horizontal matrix index
```

```
int w_unroll = h * w_out + w;
```

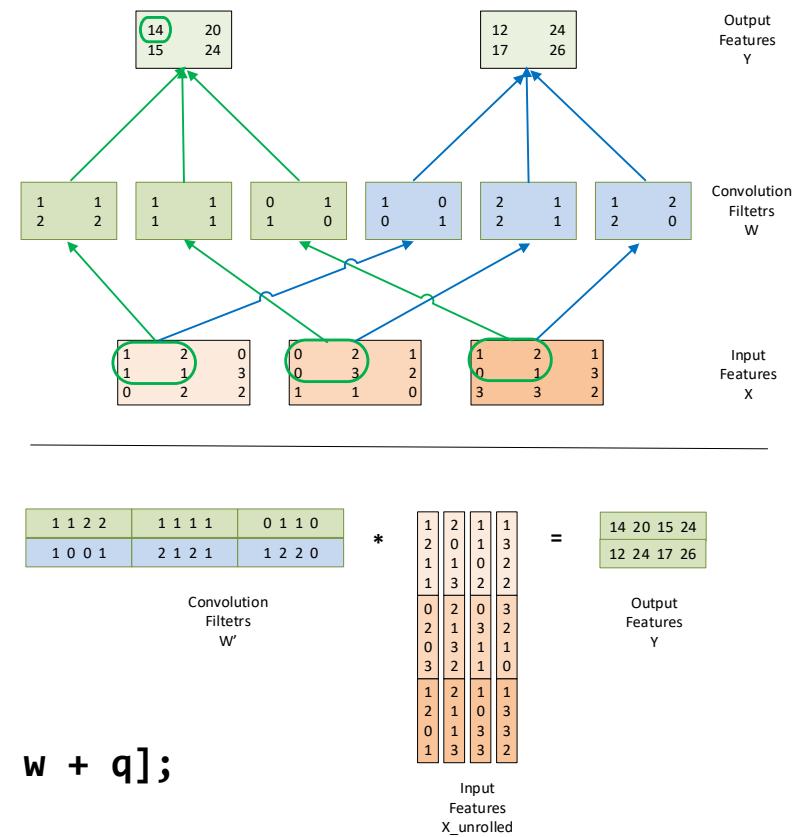
```
// find the beginning of the unrolled
```

```
int w_base = c * (K*K);
```

```
// calculate the vertical matrix index
```

```
int h_unroll = w_base + p * K + q;
```

```
X_unroll[b, h_unroll, w_unroll] = X[b, c, h + p, w + q];
```



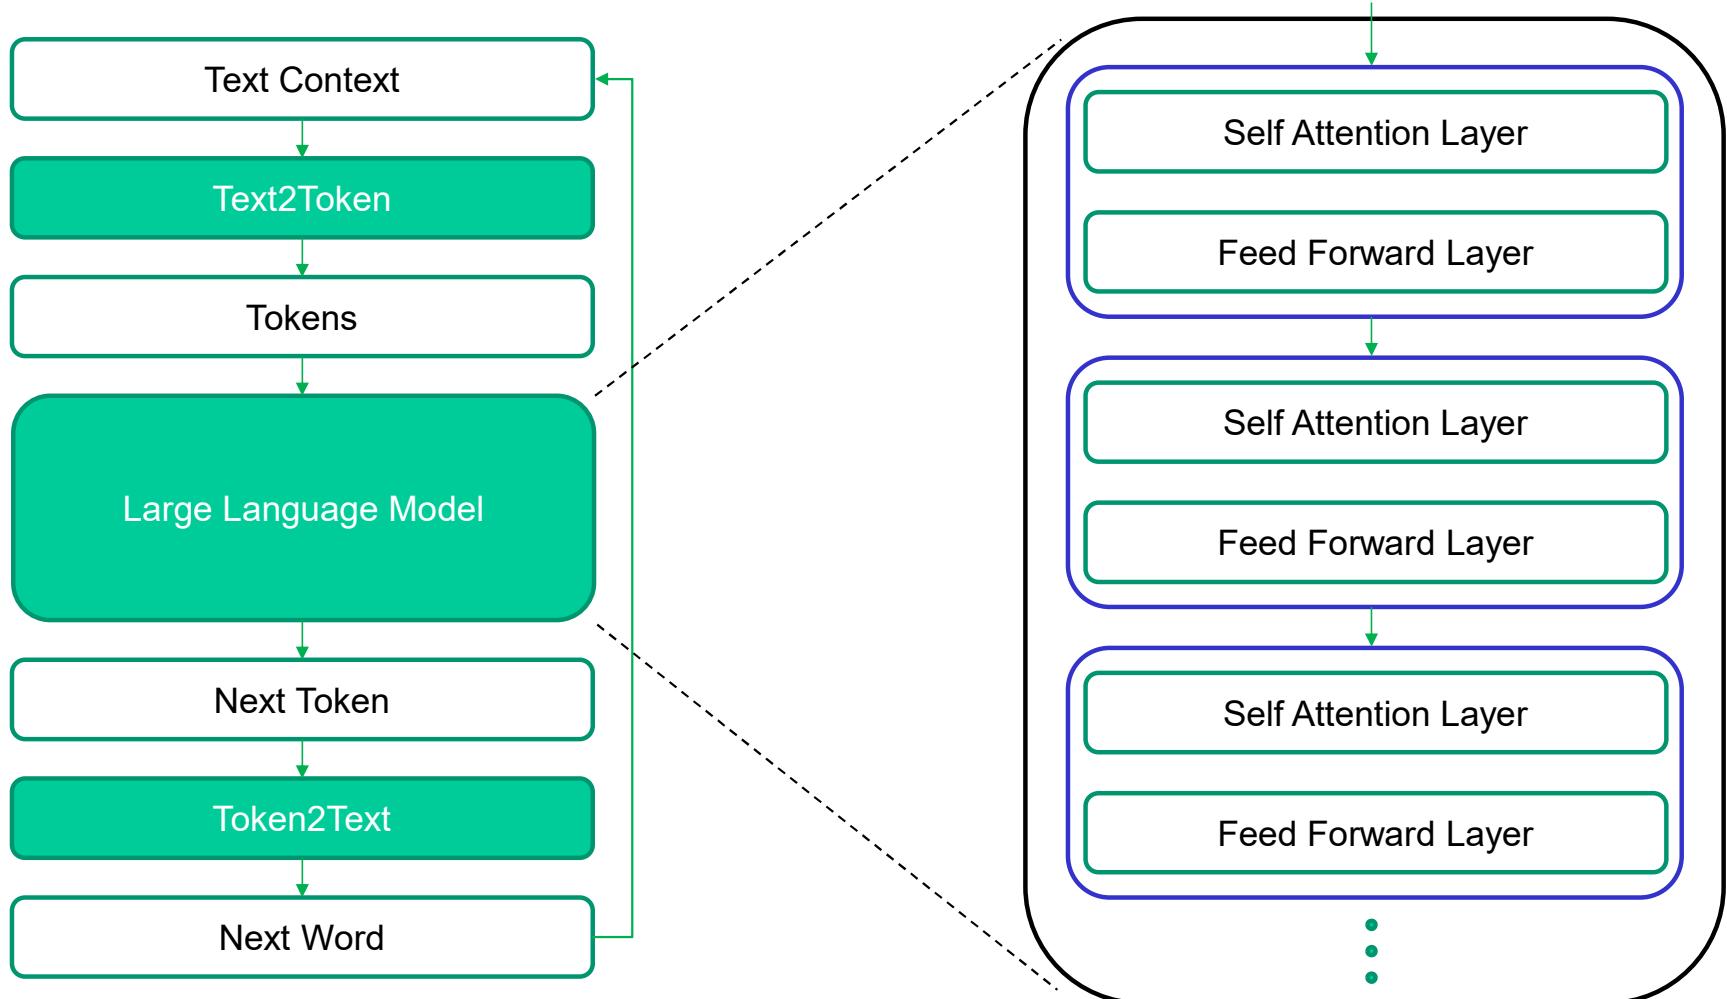
Function to generate “unrolled” X

```
void unroll(int B, int C, int H, int W, int K, float* X, float* X_unroll)
{
    int H_out = H - K + 1;                                // calculate H_out, W_out
    int W_out = W - K + 1;
    for (int b = 0; b < B; ++b)                          // for each image
        for (int c = 0; c < C; ++c) {                     // for each input channel
            int w_base = c * (K*K);                      // per-channel offset for smallest X_unroll index
            for (int p = 0; p < K; ++p) {                  // for each element of KxK filter (two loops)
                for (int q = 0; q < K; ++q) {
                    for (int h = 0; h < H_out; ++h)        // for each thread (each output value, two loops)
                        for (int w = 0; w < W_out; ++w) {
                            int h_unroll = w_base + p * K + q; // data needed by one thread
                            int w_unroll = h * W_out + w;      // smallest index--across threads (output values)
                            X_unroll[b, h_unroll, w_unroll] = X[b, c, h + p, w + q]; // copy input pixels
                        }
                }
            }
        }
}
```

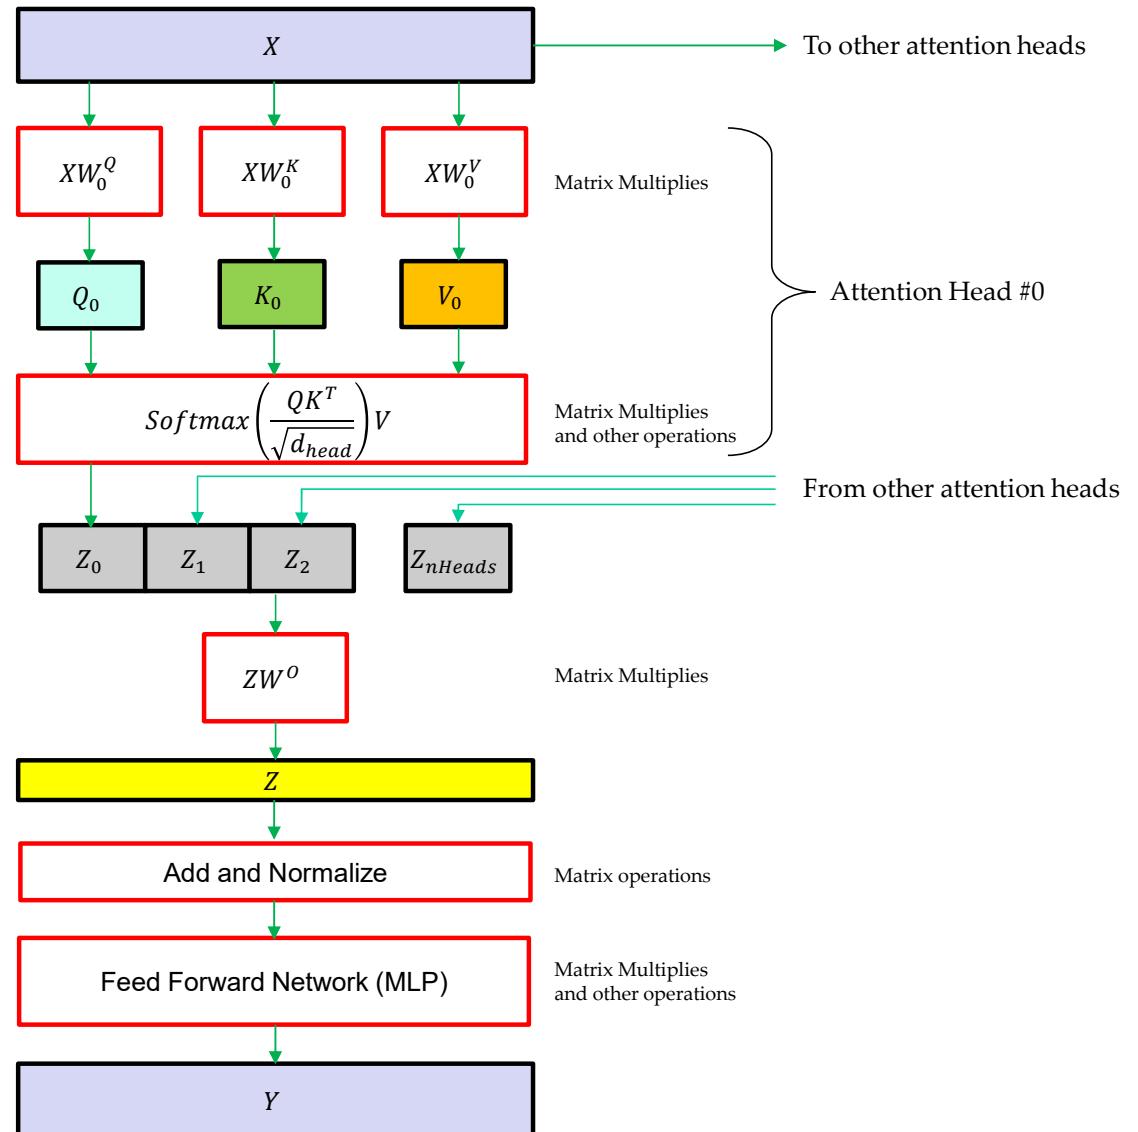
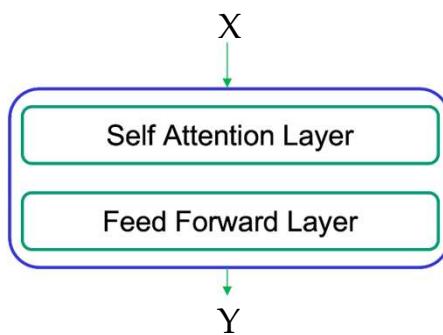
Implementation Strategies for a Convolution Layer

- **Baseline**
 - Tiled 2D convolution implementation, use constant memory for convolution masks
- **Matrix-Multiplication Baseline**
 - Input feature map unrolling kernel, constant memory for convolution masks as an optimization
 - Tiled matrix multiplication kernel
- **Matrix-Multiplication with built-in unrolling**
 - Perform unrolling only when loading a tile for matrix multiplication
 - The unrolled matrix is only conceptual
 - When loading a tile element of the conceptual unrolled matrix into the shared memory, use the properties in the lecture to load from the input feature map
- **More advanced Matrix-Multiplication**
 - Use joint register-shared memory tiling

Transformer-based Language Models



Single Layer Computational Flow



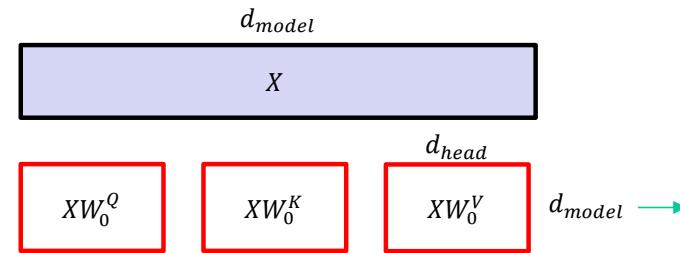
GPT-3, as an example

Language Models are Few-Shot Learners, Brown et al., OpenAI, July 2020

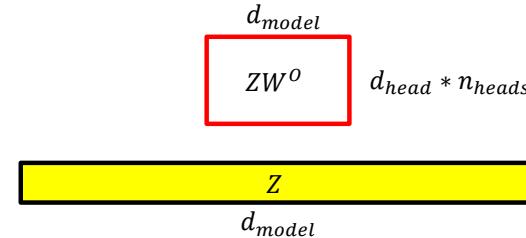
| Model Name | n_{params} | n_{layers} | d_{model} | n_{heads} | d_{head} | Batch Size | Learning Rate |
|-----------------------|---------------------|---------------------|--------------------|--------------------|-------------------|------------|----------------------|
| GPT-3 Small | 125M | 12 | 768 | 12 | 64 | 0.5M | 6.0×10^{-4} |
| GPT-3 Medium | 350M | 24 | 1024 | 16 | 64 | 0.5M | 3.0×10^{-4} |
| GPT-3 Large | 760M | 24 | 1536 | 16 | 96 | 0.5M | 2.5×10^{-4} |
| GPT-3 XL | 1.3B | 24 | 2048 | 24 | 128 | 1M | 2.0×10^{-4} |
| GPT-3 2.7B | 2.7B | 32 | 2560 | 32 | 80 | 1M | 1.6×10^{-4} |
| GPT-3 6.7B | 6.7B | 32 | 4096 | 32 | 128 | 2M | 1.2×10^{-4} |
| GPT-3 13B | 13.0B | 40 | 5140 | 40 | 128 | 2M | 1.0×10^{-4} |
| GPT-3 175B or "GPT-3" | 175.0B | 96 | 12288 | 96 | 128 | 3.2M | 0.6×10^{-4} |

Table 2.1: Sizes, architectures, and learning hyper-parameters (batch size in tokens and learning rate) of the models which we trained. All models were trained for a total of 300 billion tokens.

GPT-3 has 96 Layers, 55B parameters just from Self Attention. 220 Gflop per X vector, per output token.



In GPT-3, each W is 12288x128 matrix, 1.5M parameters, 3MB (@ 16 bit floats), or 6 MFlop per X vector



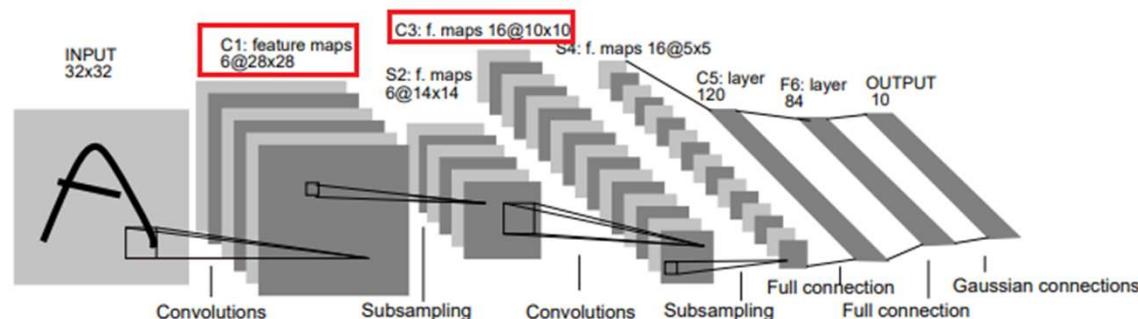
W^O is 12288x12288 matrix, 150M parameters, 300MB (@ 16 bit floats), or 600 MFlop per X vector

Project Overview

- Optimize the forward pass of the convolutional layers in a modified LeNet-5 CNN using CUDA. (CNN implemented using Mini-DNN, a C++ framework)
- The network will be classifying Fashion MNIST dataset
- Some network parameters to be aware of
 - Input Size: 86x86 pixels, batch of 1k-10k images
 - Input Channels: 1
 - Convolutional kernel size: 7x7
 - Number of kernels: Variable (your code should support this)



<https://github.com/zalandoresearch/fashion-mnist>



<http://yann.lecun.com/exdb/publis/pdf/lecun-01a.pdf>



QUESTIONS?

READ CHAPTER 16!