



ECE408 / CS483 / CSE408
Summer 2025

Applied Parallel Programming

Lecture 21:
Accelerating Matrix Operations



What Will You Learn Today?

recent optimizations to support dense matrix multiplication

- tensor units (new hardware)
- programming abstractions
- overall benefit to speed

Tiled Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, int Width)
{
1. __shared__ float subTileM[TILE_WIDTH][TILE_WIDTH];
2. __shared__ float subTileN[TILE_WIDTH][TILE_WIDTH];


3. int bx = blockIdx.x; int by = blockIdx.y;
4. int tx = threadIdx.x; int ty = threadIdx.y;

        // Identify the row and column of the P element to work on
5. int Row = by * TILE_WIDTH + ty;
6. int Col = bx * TILE_WIDTH + tx;
7. float Pvalue = 0;

        // Loop over the M and N tiles required to compute the P element
        // The code assumes that the Width is a multiple of TILE_WIDTH!
8. for (int q = 0; q < Width/TILE_WIDTH; ++q) {
            // Collaborative loading of M and N tiles into shared memory
9.     subTileM[ty][tx] = M[Row*Width + (q*TILE_WIDTH+tx)];
10.    subTileN[ty][tx] = N[(q*TILE_WIDTH+ty)*Width+Col];
11.    __syncthreads();
12.    for (int k = 0; k < TILE_WIDTH; ++k)
13.        Pvalue += subTileM[ty][k] * subTileN[k][tx];
14.    __syncthreads();
15. }
16. P[Row*Width+Col] = Pvalue;
}
```

At the Operation Level, per Thread


```
12.     for (int k = 0; k < TILE_WIDTH; ++k)
13.         Pvalue += subTileM[ty][k] * subTileN[k][tx];
```



```
ld rx <- subTileM[ty][k]           // Load into register rx
ld ry <- subTileN[k][tx]           // Load into register ry
mad rd <- rd + rx * ry           // Multiply Add Instruction
```

Each thread is able to complete the dot product in $O(TILE_WIDTH)$ cycles.


At the Operation Level, per Warp



With enough thread blocks, we have a throughput of $2 * \text{width}$ of SM FP Ops per cycle

Let's look at a 4x4 tile example

```
for (int k = 0; k < TILE_WIDTH; ++k)
    Pvalue += subTileM[ty][k] * subTileN[k][tx];
```



Shared Memory Loads into Register File by Loop Iteration

	t0	t1	t2	t3	t4	t5
k = 0	M _{0,0} , N _{0,0}	M _{0,0} , N _{0,1}	M _{0,0} , N _{0,2}	M _{0,0} , N _{0,3}	M _{1,0} , N _{0,0}	M _{1,0} , N _{0,1}
k = 1	M _{0,1} , N _{1,0}	M _{0,1} , N _{1,1}	M _{0,1} , N _{1,2}	M _{0,1} , N _{1,3}	M _{1,1} , N _{1,0}	M _{1,1} , N _{1,1}
k = 2	M _{0,2} , N _{2,0}	M _{0,2} , N _{2,1}	M _{0,2} , N _{2,2}	M _{0,2} , N _{2,3}	M _{1,2} , N _{2,0}	M _{1,2} , N _{2,1}

More Efficient Pattern


Entire M tile

M _{0,0}	M _{0,1}	M _{0,2}	...	M _{3,3}
N _{0,0}	N _{0,1}	N _{0,2}	...	N _{3,3}

Entire N tile

With two loads from shared memory across a warp (only 16 threads needed), we can load the entire 4x4 subtiles of M and N

Optimized Hardware





Notes on Optimized HW

- If we optimize for 16-bit types (and smaller)
 - We need $16b * 32$ threads = 512 bits per register
 - 4 16-bit multipliers per lane, but 32-bit adders
- Max Throughput is now increased by 4x

Nvidia Tensor Cores

$$D = \left(\begin{array}{cccc} A_{0,0} & A_{0,1} & A_{0,2} & A_{0,3} \\ A_{1,0} & A_{1,1} & A_{1,2} & A_{1,3} \\ A_{2,0} & A_{2,1} & A_{2,2} & A_{2,3} \\ A_{3,0} & A_{3,1} & A_{3,2} & A_{3,3} \end{array} \right) \left(\begin{array}{cccc} B_{0,0} & B_{0,1} & B_{0,2} & B_{0,3} \\ B_{1,0} & B_{1,1} & B_{1,2} & B_{1,3} \\ B_{2,0} & B_{2,1} & B_{2,2} & B_{2,3} \\ B_{3,0} & B_{3,1} & B_{3,2} & B_{3,3} \end{array} \right) + \left(\begin{array}{cccc} C_{0,0} & C_{0,1} & C_{0,2} & C_{0,3} \\ C_{1,0} & C_{1,1} & C_{1,2} & C_{1,3} \\ C_{2,0} & C_{2,1} & C_{2,2} & C_{2,3} \\ C_{3,0} & C_{3,1} & C_{3,2} & C_{3,3} \end{array} \right)$$

FP16 or FP32 FP16 FP16 or FP32

Basic Tensor Operator: Matrix Multiply + Accumulate



CUDA TENSOR CORE PROGRAMMING

WMMA Matrix Multiply and Accumulate Operation


Warp-level operation to perform matrix multiply and accumulate

```
wmma::mma_sync(Dmat, Amat, Bmat, Cmat);
```

$$D = \left(\begin{matrix} \text{teal} & \text{teal} & \text{teal} \\ \text{teal} & \text{teal} & \text{teal} \\ \text{teal} & \text{teal} & \text{teal} \\ \text{teal} & \text{teal} & \text{teal} \end{matrix} \right) \left(\begin{matrix} \text{purple} & \text{purple} & \text{purple} \\ \text{purple} & \text{purple} & \text{purple} \\ \text{purple} & \text{purple} & \text{purple} \\ \text{purple} & \text{purple} & \text{purple} \end{matrix} \right) + \left(\begin{matrix} \text{green} & \text{green} & \text{green} \\ \text{green} & \text{green} & \text{green} \\ \text{green} & \text{green} & \text{green} \\ \text{green} & \text{green} & \text{green} \end{matrix} \right)$$


TENSOR SYNCHRONIZATION

Full Warp 16x16 Matrix Math



CUDA TENSOR CORE PROGRAMMING

WMMA datatypes




Per-Thread fragments to hold components of matrices for use with Tensor Cores

```
wmma::fragment<matrix_a, ...> Amat;
```

CUDA TENSOR CORE PROGRAMMING

WMMA load and store operations


Warp-level operation to fetch components of matrices into fragments



CUDA TENSOR CORE PROGRAMMING

WMMA load and store operations

Warp-level operation to fetch components of matrices into fragments



TENSOR CORE EXAMPLE

Create Fragments

Initialize Fragments

Perform MatMul

Store Results

```
__device__ void tensor_op_16_16_16(
    float *d, half *a, half *b, float *c)
{
    wmma::fragment<matrix_a, ...> Amat;
    wmma::fragment<matrix_b, ...> Bmat;
    wmma::fragment<matrix_c, ...> Cmat;

    wmma::load_matrix_sync(Amat, a, 16);
    wmma::load_matrix_sync(Bmat, b, 16);
    wmma::fill_fragment(Cmat, 0.0f);


    wmma::mma_sync(Cmat, Amat, Bmat, Cmat);

    wmma::store_matrix_sync(d, Cmat, 16,
                           wmma::row_major);
}
```

CUDA C++
Warp-Level Matrix Operations

Gains from



- Number Representation
 - FP32, FP16, Int8
 - (TF32, BF16)
 - ~16x
- Complex Instructions
 - DP4, HMMA, IMMA
 - ~12.5x
- Process
 - 28nm, 16nm, 7nm, 5nm
 - ~2.5x
- Sparsity
 - ~2x
- Model efficiency has also improved – overall gain > 1000x



Bill Dally, Chief Scientist, Nvidia



Google TPU (v1 in 2016)



Intel Advance Matrix Extensions (AMX)

- Introduced in 2020
- Introduced into x86 processors
- Usable directly from CPU code

```
// Load tile configuration
init_tile_config (&tile_data);

// Init src matrix buffers with data
init_buffer (src1, 2);
print_buffer(src1, rows, colsb);

init_buffer (src2, 2);
print_buffer(src2, rows, colsb);

// Init dst matrix buffers with data
init_buffer32 (res, 0);

// Load tile rows from memory
_tile_loadd (2, src1, STRIDE);
_tile_loadd (3, src2, STRIDE);
_tile_loadd (1, res, STRIDE);

// Compute dot-product of bytes in tiles with a
// source/destination accumulator
_tile_dpbsd (1, 2, 3);

// Store the tile data to memory
_tile_stored (1, res, STRIDE);
```



QUESTIONS?

**READ NVIDIA TECHNICAL BLOG &
CUDA PG: WARP MATRIX FUNCTIONS**